

15-319 / 15-619

Cloud Computing

Recitation 4

February 6th, 2017

Administrative Issues

- Make use of office hours
 - We will have to make sure that you have tried yourself before you ask
- Monitor AWS (and Azure / GCP) expenses regularly
- Always do the cost calculation before launching services
- Terminate your instances when not in use
- Stopping instances still has an EBS cost (\$0.1/GB-Month)
- Make sure spot instances are tagged right after launch

Important Notice

- **DON'T EVER EXPOSE YOUR AWS CREDENTIALS!**
 - Github
 - Bitbucket
 - Anywhere public...
- **DON'T EVER EXPOSE YOUR GCP CREDENTIALS!**
- **DON'T EVER EXPOSE YOUR Azure CREDENTIALS!**
 - ApplicationId, ApplicationKey
 - StorageAccountKey, EndpointUrl

Reflection

- Last week's reflection
 - Project 1.2, Quiz 2
- Theme - **Big data analytics**
 - P1.1: Sequential Analysis of **100s MB** of wikipedia data
 - P1.2: Parallel Analysis of **35GB** compressed / **128GB** decompressed wikipedia data
- Power of parallel analysis
 - Amount of work done remains the same
 - Span is reduced significantly

Reflection

- You should have learned
 - How to process big data sets with MapReduce
 - How MapReduce works
 - How to write a Mapper and a Reducer
 - Performance/cost tradeoff
 - How to debug MapReduce
 - How to save overall cost by testing using small data sets
- Don't forget about MapReduce just yet!
 - Will be relevant in the Team Project and Project 4

This Week

- Quiz 3 (OLI Modules 5 & 6)
 - Due on Friday,
Feb 9th, 2018,
11:59PM ET
- Project 2.1
 - Due on Sunday,
Feb 11th, 2018,
11:59PM ET

OLI Module 5 - Cloud Management

Cloud Software stack - enables provisioning, monitoring and metering of virtual user “resources” on top of the Cloud Service Provider’s (CSP) infrastructure.

- Cloud middleware
- Provisioning
- Metering
- Orchestration and automation
- Case Study: Openstack - Open-source cloud stack implementation

OLI Module 6 - Cloud Software Deployment Considerations

- Programming the cloud
- Deploying applications on the cloud
- Build fault-tolerant cloud services
- Load balancing
- Scaling resources
- Dealing with tail latency
- Economics for cloud applications

Project 2 Overview

Scaling and Elasticity with
VMs, Containers & Functions

- **2.1 Scaling Virtual Machines**
 - Horizontal scaling in / out using AWS APIs
 - Load balancing, failure detection, and cost management on AWS
 - Infrastructure as Code (Terraform)
- **2.2 Scaling with Containers**
 - Building your own container-based services. Frontend and Backend services.
 - Docker containers
 - Manage multiple Kubernetes Cluster
 - Multi Cloud deployments
- **2.3 Functions as a Service**
 - Develop event driven cloud functions
 - Deploy multiple functions to build a video processing pipeline

Project 2.1 Learning Objectives

- **Design** solutions and invoke cloud APIs to programmatically provision and deprovision cloud resources based on the current load.
- **Explore** and compare the usability and performance of APIs used in AWS.
- **Configure** and deploy an Elastic Load Balancer along with an Auto Scaling Group on AWS.
- **Develop** solutions that manage cloud resources with the ability to deal with resource failure.
- **Account** for cost as a constraint when provisioning cloud resources and **analyze** the performance tradeoffs due to budget restrictions.
- **Experience** using cloud orchestration and automation tools such as Terraform.

Quality of Service (QoS)

Quantitatively Measure QoS

- **Performance: Throughput, Latency**
(Very helpful in Projects 2 & Team Project)
- **Availability:** the probability that a system is operational at a given time *(Projects P2.1 and P2.2)*
- **Reliability:** the probability that a system will produce a correct output up to a given time *(Project P2.1 and P2.2)*

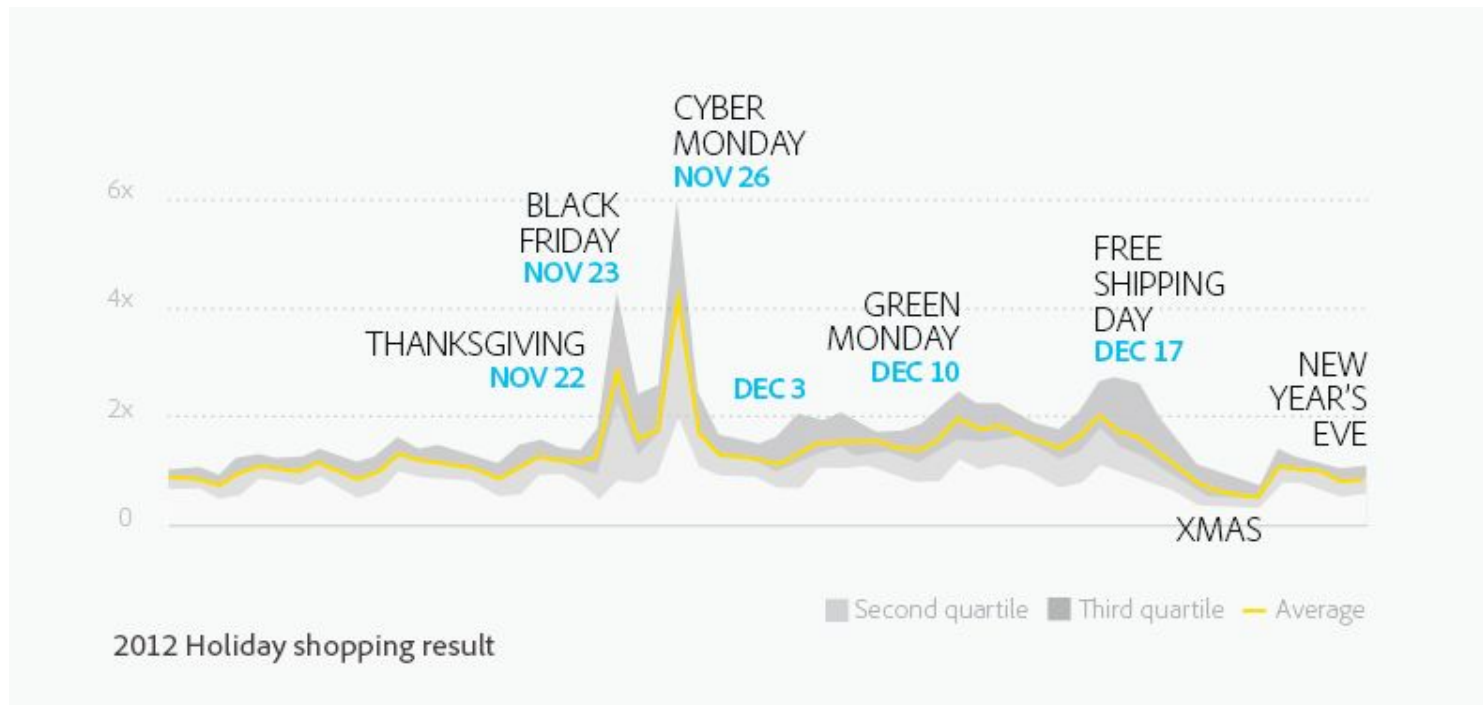
QoS Matters:

- Amazon found every **100ms** of latency cost them **1%** in sales.



Reality, human patterns...


- Daily
- Weekly
- Monthly
- Yearly
- ...




Cloud Comes to the Rescue!

Scaling!

P0: Vertical Scaling




Resources in Cloud Infrastructure

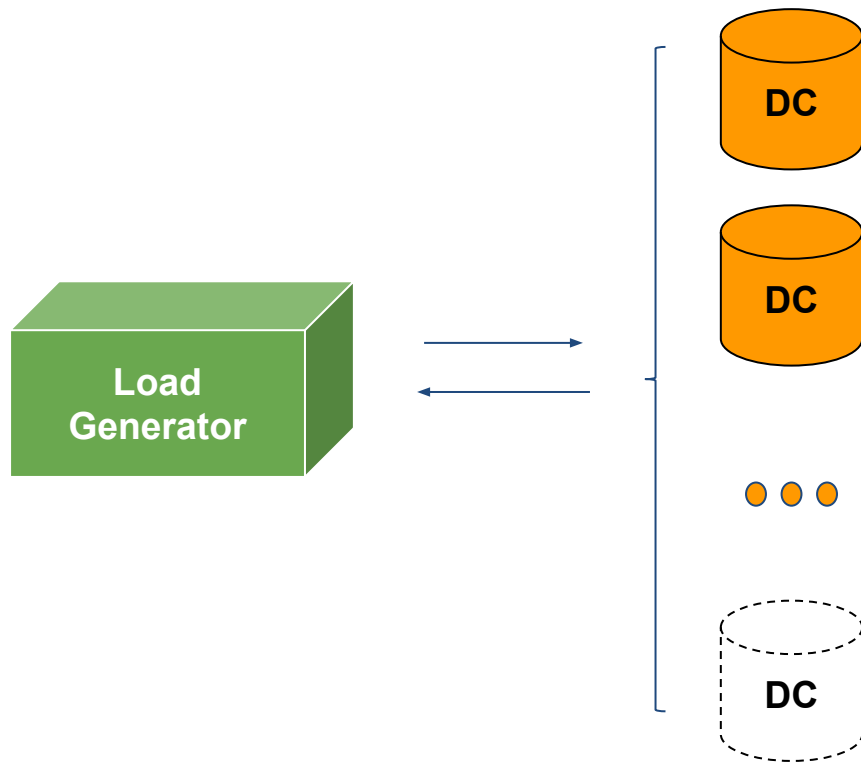


P0: Vertical Scaling Limitation


- However, one instance will always have limited resources.
- Reboot/Downtime.




Horizontal Scaling



How do we distribute load?



Instance Failure?



CPU utilization, memory utilization...




Available capacity

What You Need


- Make sure that workload is even on each server
- Do not assign load to servers that are down
- Increase/Remove servers according to changing load

How does a cloud service help solve these problems?




Load balancer

- “Evenly” distribute the load
 - Simplest distribution strategy
 - Round Robin
 - Health Check
-
- What if the Load Balancer becomes the bottleneck?
 - Elastic Load Balancer
 - Could scale up based on load
 - Elastic, but it takes time
 - Through the warm-up process



Load Balancer

Reality...



sapient.com


Scaling

Manual Scaling:

- Expensive on manpower
- Low utilization or over provisioning
- Manual control
- Lose customers

Autoscaling:

- Automatically adjust the size based on demand
- Flexible capacity and scaling sets
- Save cost



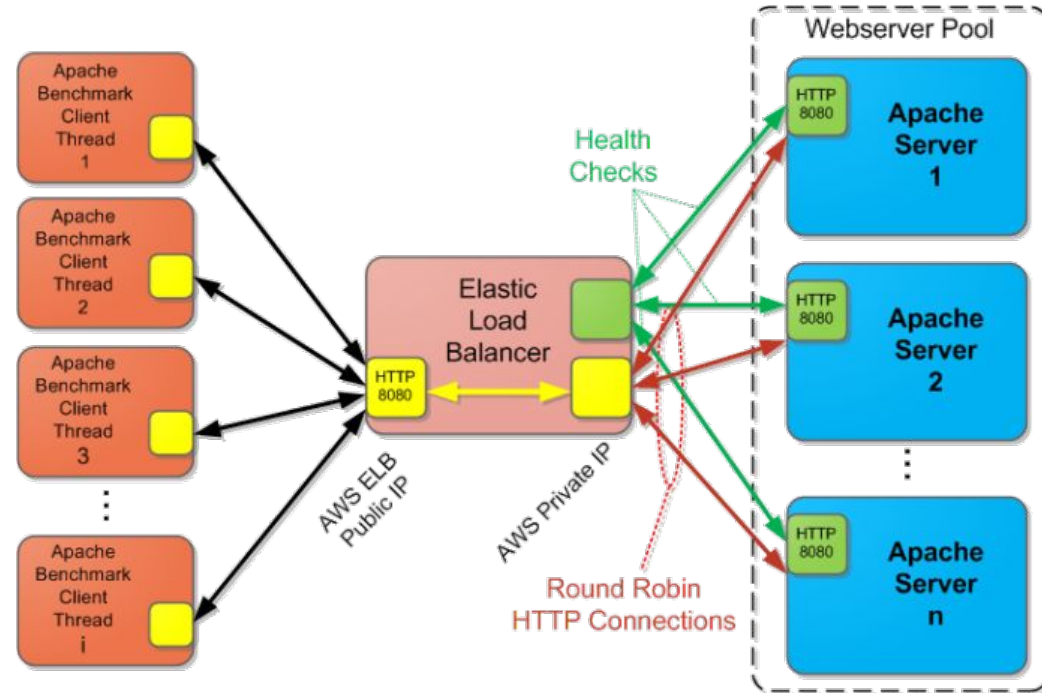
AWS Autoscaling

Auto Scaling on AWS


Using the AWS APIs:

- CloudWatch
- ELB
- Auto Scaling Group
- Auto Scaling Policy
- EC2

You can build a load balanced auto-scaled web service.

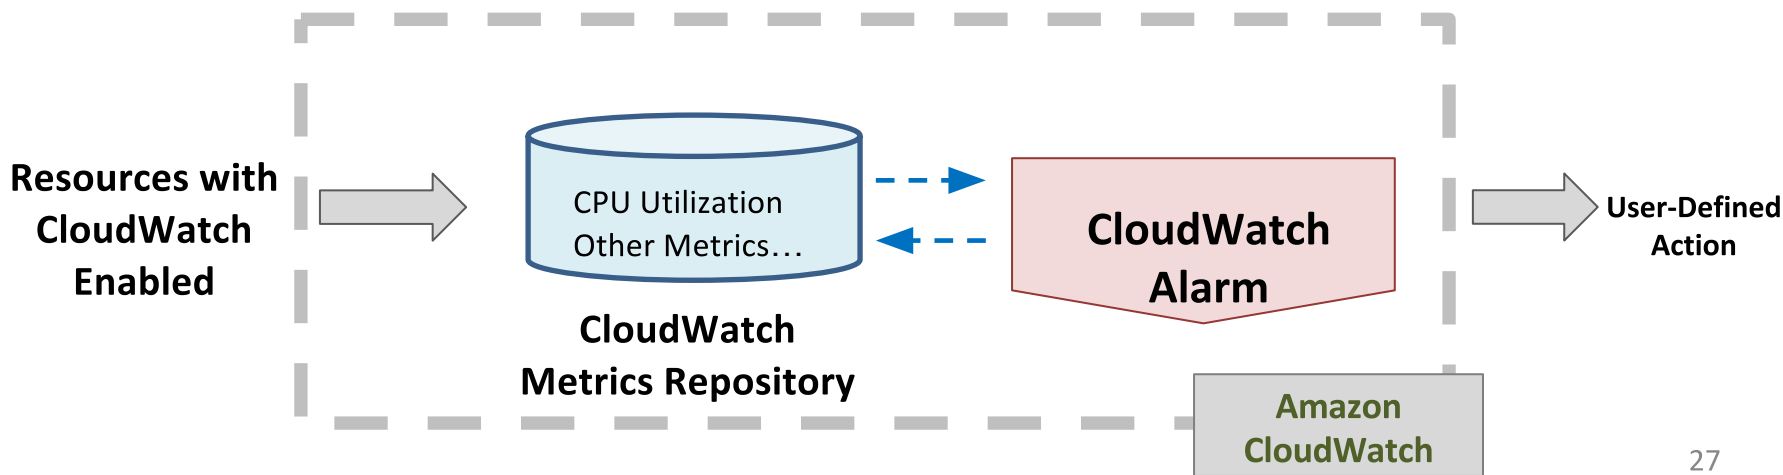


Amazon Auto Scaling Group



Amazon's CloudWatch Alarm

- Monitor CloudWatch metrics for some specified alarm conditions
- Take automated action when the condition is met



Terraform Configuration

- **Providers**

- A provider is responsible for understanding API interactions and exposing resources.

- **Resources**

- The resource block defines a resource that exists within the infrastructure.

- **AWS Provider**

- <https://www.terraform.io/docs/providers/aws/index.html>

```
$ cat main.tf
provider "aws" {
  region  = "us-east-1"
}

resource "aws_instance" "cmucc" {
  ami          = "ami-2757f631"
  instance_type = "t2.micro"

  tags {
    Project = "2.1"
  }

  key_name = "my-ssh-key"
}
```

Terraform CLI

- **init**
 - Initializes a working directory containing Terraform configuration files.
- **plan**
 - Creates an execution plan.
- **apply**
 - Apply the changes required to reach the desired state.
- **destroy**
 - Destroy the Terraform-managed infrastructure.

```
$ terraform plan
```

```
...
```

Terraform will perform the following actions:

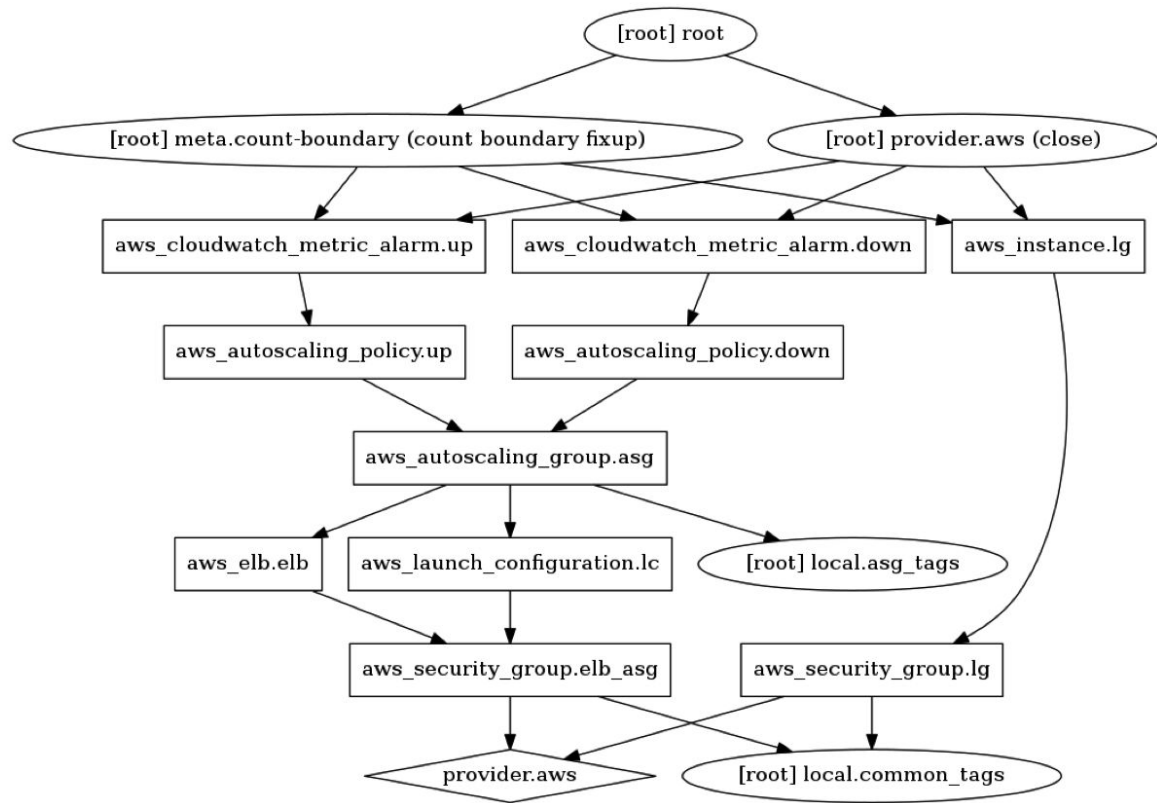
```
+ aws_instance.cmucc
  id:                  <computed>
  ami:                  "ami-2757f631"
  ...
```

Plan: 1 to add, 0 to change, 0 to destroy.

```
$ terraform apply
```

```
...
```


Declarative Infrastructures with Terraform



Do not expect this to exactly match your solution (e.g. the load balancer is supposed to be an Application Load Balancer)!


Project 2.1 Scaling on AWS

- Step 1
 - AWS Horizontal Scaling
- Step 2
 - AWS Auto Scaling
- Step 3
 - AWS Auto Scaling with Terraform



Project 2.1 Scaling on AWS


- Step 1 - AWS Horizontal Scaling
- Implement Horizontal Scaling in AWS.
- Write a program that launches the data center instances and ensures that the target total RPS is reached.
- Your program should be fully automated: launch LG->submit password-> Launch DC-> start test-> check log -> add more DC...



Project 2.1 Scaling on AWS


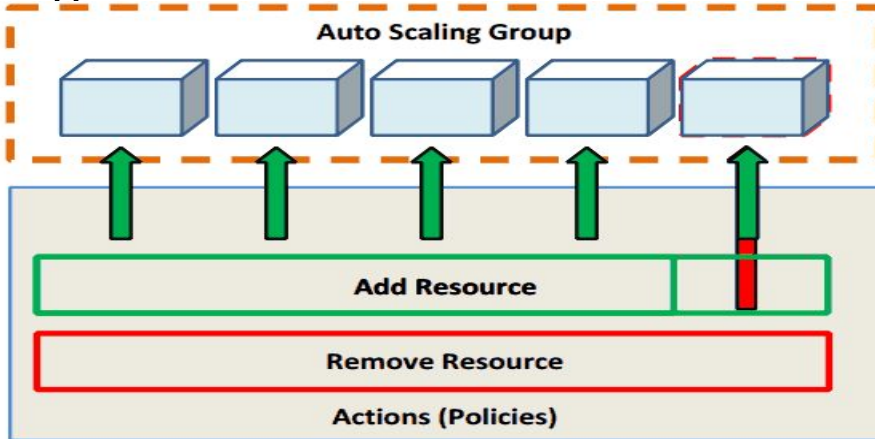


- Step 2 - AWS Auto Scaling



P2.1 - Step 2 - Your Tasks

- Programmatically create an Application Load Balancer (ALB) and an Auto Scaling Policy. Attach the policy to Auto-Scaling Group (ASG) and link ASG to ALB.
- Test by submitting a URL request and observe logs, ALB, and CloudWatch.
- Decide on the Scale-Out and Scale-In policies
- Mitigate failure



Hints for Project 2.1 AWS Autoscaling



Step 2 - AWS Auto Scaling

- **Autoscaling Test could be very EXPENSIVE!**
 - on-demand and charged by the hour
- **Determine if there is a less expensive means to test your solution**
- **Creating and deleting security groups can be tricky**
- **CloudWatch and monitoring in ELB is helpful**
- **Explore ways to check if your instance is ready**
- **Understanding the API documents could take time**
- **Finish parts 1-3 first, the experience will help**

Project 2.1 Scaling on AWS

- Step 3 - AWS Auto Scaling with Terraform
- Read the Infrastructure as Code primer to learn about infrastructure automation
- Update the started Terraform configuration to launch the AWS resource required for autoscaling (ALB, Alarms, ASG...)
- Remove the code to launch these resources (as they are in the TF config)
- Makes sure that **terraform plan** generates the expected resource

Project 2.1 Code Submission

AWS:

- Submit the horizontal scaling task on AWS's load generator (LG) instance
- Submit the autoscaling task to the AWS load generator (LG) instance
- Submit the Terraform autoscaling task to the AWS load generator (LG) instance
 - You can use the test id from the AWS autoscaling task when creating the ZIP file

Penalties for 2.1

Violation	Penalty of the project grade
Spending more than \$20 for this project phase on AWS	-10%
Spending more than \$35 for this project phase on AWS	-100%
Failing to tag all your resources in either parts (EC2 instances, ELB, ASG) for this project with the tag: key=Project, value=2.1 (AWS only)	-10%
Submitting your AWS/Andrew credentials in your code for grading	-100%
Using instances other than t2.micro (testing only) or m3.medium for Horizontal scaling on AWS	-100%
Using instances other than t2.micro (testing only), m3.medium for Autoscaling on AWS	-100%
Submitting executables (.jar, .pyc, etc.) instead of human-readable code (.py,.java, .sh, etc.)	-100%

Penalties for 2.1 cont.

Violation	Penalty of the project grade
Attempting to hack/tamper the autograder in any way	-100%
Cheating, plagiarism or unauthorized assistance (please refer to the university policy on academic integrity and our syllabus)	-100%

AWS Cloud APIs



- AWS CLI ([link](#))
- AWS Java SDK ([link](#))
- AWS Python SDK ([link](#))

This Week's Deadlines

- **Quiz 3 (OLI Modules 5 & 6)**
 - Due on **Friday**,
Feb 9th, 2018,
11:59PM ET
- **Project 2.1**
 - Due on **Sunday**,
Feb 11th, 2018,
11:59PM ET

Questions?