

2013

# Implementing a dynamic scaling of web applications in a virtualized cloud computing environment

Mohammed Aljebreen

Follow this and additional works at: <http://scholarworks.rit.edu/theses>

---

## Recommended Citation

Aljebreen, Mohammed, "Implementing a dynamic scaling of web applications in a virtualized cloud computing environment" (2013). Thesis. Rochester Institute of Technology. Accessed from

**Implementing a Dynamic Scaling of Web Applications  
in a Virtualized Cloud Computing Environment****By**

Mohammed Aljebreen

**Committee Chair:**

Professor Sharon Mason

**Committee Members:**Professor Larry Hill  
Professor Jim Leone

Thesis submitted in partial fulfillment of the requirements for the degree  
of Master of Science in

***Networking and System Administration***

Rochester Institute of Technology

B. Thomas Golisano College of Computing and Information Sciences

## **Abstract**

Cloud computing is becoming more essential day by day. The allure of the cloud is the significant value and benefits that people gain from it, such as reduced costs, increased storage, flexibility, and more mobility. Flexibility is one of the major benefits that cloud computing can provide in terms of scaling up and down the infrastructure of a network. Once traffic has increased on one server within the network, a load balancer instance will route incoming requests to a healthy instance, which is less busy and less burdened. When the full complement of instances cannot handle any more requests, past research has been done by Chieu et. al. that presented a scaling algorithm to address a dynamic scalability of web applications on a virtualized cloud computing environment based on relevant indicators that can increase or decrease servers, as needed. In this project, I implemented the proposed algorithm, but based on CPU Utilization threshold. In addition, two tests were run exploring the capabilities of different metrics when faced with ideal or challenging conditions. The results did find a superior metric that was able to perform successfully under both tests.

## **Dedication**

I lovingly dedicate this thesis to my gracious and devoted mother for her unwavering  
love and for always believing in me.

### **Acknowledgments**

This thesis would not have been possible without the support of many people.

My wish is to express humble gratitude to the committee chair, Prof. Sharon Mason, who was perpetually generous in offering her invaluable assistance, support, and guidance. Deepest gratitude is also due to the members of my supervisory committee, Prof. Lawrence Hill and Prof. Jim Leone, without whose knowledge and direction this study would not have been successful. Special thanks also to Prof. Charles Border for his financial support of this thesis and priceless assistance.

Profound gratitude to my mother, Moneerah, who has been there from the very beginning, for her support and endless love. I would also like to convey thanks to my wife for her patient and unending encouragement and support throughout the duration of my studies; without my wife's encouragement, I would not have completed this degree. I wish to express my gratitude to my beloved sister and brothers for their kind understanding throughout my studies. Special thanks to my friend, Mohammed Almathami, for his invaluable assistance.

## ***Table of Contents***

<b>LIST OF FIGURES .....</b>	<b>6</b>
<b>LIST OF TABLES .....</b>	<b>9</b>
<b>1. INTRODUCTION .....</b>	<b>10</b>
1.1 Essential Characteristics: .....	10
1.2 Service Models: .....	11
1.3 Deployment Models:.....	11
<b>2. BACKGROUND .....</b>	<b>13</b>
<b>3. LITERATURE REVIEW .....</b>	<b>13</b>
<b>4. MATERIALS &amp; METHODS:.....</b>	<b>15</b>
<b>4.1 Metrics for the EC2 instances:.....</b>	<b>16</b>
4.1.1 CPU Utilization Metric:.....	16
4.1.2 Network Out Metric:.....	17
<b>4.2 Metrics for the load balancer:.....</b>	<b>17</b>
4.2.1 Latency Metric: .....	17
4.2.2 Request Count Metric: .....	17
<b>4.3 The Repetition of the Tests:.....</b>	<b>20</b>
<b>4.4 The project's architecture:.....</b>	<b>22</b>
4.4.1 Launching Amazon EC2 instances:.....	22
4.4.2 Creating an Amazon Machine Image (AMI): .....	31
4.4.3 Creating an Amazon Load Balancer:.....	32
4.4.4 Configuring Amazon Auto Scaling:.....	35
4.4.5 Creating CloudWatch Alarms: .....	41
<b>5. DEMONSTRATIONS &amp; RESULTS: .....</b>	<b>45</b>
<b>5.1 Auto Scaling Testing:.....</b>	<b>45</b>
<b>5.2 Tests' Results:.....</b>	<b>47</b>
5.2.1 CPU Utilization Metric Results:.....	48
5.2.2 Network Out Metric Results: .....	55
5.2.3 Latency Metric Results: .....	62
5.2.4 Request Counts Metric Results: .....	69
5.2.5 Discussion: .....	76
<b>6. CONCLUSION &amp; FUTURE WORK.....</b>	<b>81</b>
<b>7. REFERENCES.....</b>	<b>82</b>

## LIST OF FIGURES

<b>Figure</b>		<b>Page</b>
1.1	Typical Cloud Computing Environment	12
4.1	Dynamic Scaling Algorithm for Virtual Machine Instances in the Cloud	16
4.1.1	Policies of scaling up and down for all tested metrics as shown in AWS management console	18
4.2	Architecture to Scale Web Applications in a Cloud	22
4.3	Amazon Web Services Management Console	23
4.4	Amazon EC2 Console Dashboard	23
4.5	Create a New Instance page	24
4.6	Choose the AMI from the Quick Start tap	25
4.7	Choose the number and the type of the instance	25
4.8	Naming the instance	26
4.9	Key pair creation	26
4.10	Security Group Creation	27
4.11	The web server instance	27
4.12	The web server started	28
4.13	Apache2 included in the instance services	28
4.14	Apache Service Monitor	29
4.15	The home web page of the website	31
4.16	The AMI from the web server	32
4.17	Web server Load Balancer	33
4.18	The servers under the load balancer	33
4.19	Adding Webserver (1) & (2) to the Load Balancer	33
4.20	Advanced Options of the load balancer's configuration	34
4.21	Java version 1.7 installed on the web server	35
4.22	Testing the Auto Scaling command line installation	36
4.23	MyAutoConfig creation	37
4.24	MyGroup2 creation	39
4.25	Verifying the Auto Scaling Group creation	39
4.26	Verifying the Auto Scaling instances	39
4.27	Scale Up Policy	40
4.28	Scale Down Policy	41
4.29	Testing the CloudWatch command line installation	42
4.30	HighCPU-alarm	43
4.31	LowCPU-alarm	43
4.32	The alarms of the CloudWatch in AWS management console	44
5.1	Auto Scaling creates an instance	45
5.2	CloudWatch Alarms	46
5.3	A notification email in aljebreen50@gmail.com inbox	46
5.4	Shutting the temporary instance down	47
5.5	Terminating the temporary instance	47
5.1.1	CPU Utilization Metric's results for the first test <i>Weekday mornings at RIT System Lab</i>	48
5.1.2	<i>Weekday evenings at RIT System Lab</i>	49
5.1.3	<i>Weekend mornings at RIT System Lab</i>	49
5.1.4	<i>Weekend evenings at RIT System Lab</i>	49
5.1.5	<i>Weekday mornings from home</i>	49
5.1.6	<i>Weekday evenings from home</i>	50
5.1.7	<i>Weekend mornings from home</i>	50
5.1.8	<i>Weekend evenings from home</i>	50

5.1.9	<i>From a server within Amazon datacenter in Virginia web server</i>	50
5.1.10	<i>From a server within Amazon in Oregon</i>	51
5.1.11	<i>From a server within Amazon in Northern California</i>	51
CPU Utilization Metric's results for the second test		
5.2.1	<i>Weekday mornings at RIT System Lab</i>	52
5.2.2	<i>Weekday evenings at RIT System Lab</i>	52
5.2.3	<i>Weekend mornings at RIT System Lab</i>	52
5.2.4	<i>Weekend evenings at RIT System Lab</i>	53
5.2.5	<i>Weekday mornings from home</i>	53
5.2.6	<i>Weekday evenings from home</i>	53
5.2.7	<i>Weekend mornings from home</i>	53
5.2.8	<i>Weekend evenings from home</i>	54
5.2.9	<i>From a server within Amazon datacenter in Virginia web server</i>	54
5.2.10	<i>From a server within Amazon in Oregon</i>	54
5.2.11	<i>From a server within Amazon in Northern California</i>	55
Network Out Metric's results for the first test		
5.3.1	<i>Weekday mornings at RIT System Lab</i>	56
5.3.2	<i>Weekday evenings at RIT System Lab</i>	56
5.3.3	<i>Weekend mornings at RIT System Lab</i>	56
5.3.4	<i>Weekend evenings at RIT System Lab</i>	56
5.3.5	<i>Weekday mornings from home</i>	57
5.3.6	<i>Weekday evenings from home</i>	57
5.3.7	<i>Weekend mornings from home</i>	57
5.3.8	<i>Weekend evenings from home</i>	57
5.3.9	<i>From a server within Amazon datacenter in Virginia web server</i>	58
5.3.10	<i>From a server within Amazon in Oregon</i>	58
5.3.11	<i>From a server within Amazon in Northern California</i>	58
Network Out Metric's results for the second test		
5.4.1	<i>Weekday mornings at RIT System Lab</i>	59
5.4.2	<i>Weekday evenings at RIT System Lab</i>	59
5.4.3	<i>Weekend mornings at RIT System Lab</i>	60
5.4.4	<i>Weekend evenings at RIT System Lab</i>	60
5.4.5	<i>Weekday mornings from home</i>	60
5.4.6	<i>Weekday evenings from home</i>	60
5.4.7	<i>Weekend mornings from home</i>	61
5.4.8	<i>Weekend evenings from home</i>	61
5.4.9	<i>From a server within Amazon datacenter in Virginia web server</i>	61
5.4.10	<i>From a server within Amazon in Oregon</i>	61
5.4.11	<i>From a server within Amazon in Northern California</i>	62
Latency Metric's results for the first test		
5.5.1	<i>Weekday mornings at RIT System Lab</i>	63
5.5.2	<i>Weekday evenings at RIT System Lab</i>	63
5.5.3	<i>Weekend mornings at RIT System Lab</i>	63
5.5.4	<i>Weekend evenings at RIT System Lab</i>	63
5.5.5	<i>Weekday mornings from home</i>	64
5.5.6	<i>Weekday evenings from home</i>	64
5.5.7	<i>Weekend mornings from home</i>	64
5.5.8	<i>Weekend evenings from home</i>	64
5.5.9	<i>From a server within Amazon datacenter in Virginia web server</i>	65
5.5.10	<i>From a server within Amazon in Oregon</i>	65
5.5.11	<i>From a server within Amazon in Northern California</i>	65
Latency Metric's results for the second test		
5.6.1	<i>Weekday mornings at RIT System Lab</i>	66
5.6.2	<i>Weekday evenings at RIT System Lab</i>	66
5.6.3	<i>Weekend mornings at RIT System Lab</i>	66
5.6.4	<i>Weekend evenings at RIT System Lab</i>	67
5.6.5	<i>Weekday mornings from home</i>	67

5.6.6	<i>Weekday evenings from home</i>	67
5.6.7	<i>Weekend mornings from home</i>	67
5.6.8	<i>Weekend evenings from home</i>	68
5.6.9	<i>From a server within Amazon datacenter in Virginia web server</i>	68
5.6.10	<i>From a server within Amazon in Oregon</i>	68
5.6.11	<i>From a server within Amazon in Northern California</i>	69
Request Counts Metric's results for the first test		
5.7.1	<i>Weekday mornings at RIT System Lab</i>	70
5.7.2	<i>Weekday evenings at RIT System Lab</i>	70
5.7.3	<i>Weekend mornings at RIT System Lab</i>	70
5.7.4	<i>Weekend evenings at RIT System Lab</i>	70
5.7.5	<i>Weekday mornings from home</i>	71
5.7.6	<i>Weekday evenings from home</i>	71
5.7.7	<i>Weekend mornings from home</i>	71
5.7.8	<i>Weekend evenings from home</i>	71
5.7.9	<i>From a server within Amazon datacenter in Virginia web server</i>	72
5.7.10	<i>From a server within Amazon in Oregon</i>	72
5.7.11	<i>From a server within Amazon in Northern California</i>	72
Request Counts Metric's results for the second test		
5.8.1	<i>Weekday mornings at RIT System Lab</i>	73
5.8.2	<i>Weekday evenings at RIT System Lab</i>	73
5.8.3	<i>Weekend mornings at RIT System Lab</i>	74
5.8.4	<i>Weekend evenings at RIT System Lab</i>	74
5.8.5	<i>Weekday mornings from home</i>	74
5.8.6	<i>Weekday evenings from home</i>	74
5.8.7	<i>Weekend mornings from home</i>	75
5.8.8	<i>Weekend evenings from home</i>	75
5.8.9	<i>From a server within Amazon datacenter in Virginia web server</i>	75
5.8.10	<i>From a server within Amazon in Oregon</i>	75
5.8.11	<i>From a server within Amazon in Northern California</i>	76
5.9.1	The Average of time exceeded the metrics' boundary in the First Test	76
5.9.2	The Average of time exceeded the metrics' boundary in the Second Test	77

## LIST OF TABLES

<b>Table</b>		<b>Page</b>
1	Policies of scaling up and down for all tested metrics	17
2	CPU Utilization Metric Results of the First Test	48
3	CPU Utilization Metric Results of the Second Test	51
4	Network Out Metric Results of the First Test	55
5	Network Out Metric Results of the Second Test	58
6	Latency Metric Results of the First Test	62
7	Latency Metric Results of the Second Test	65
8	Request Counts Metric Results of the First Test	69
9	Request Counts Metric Results of the Second Test	72

## **1. INTRODUCTION**

Over the years, the National Institute of Standards and Technology (NIST) has worked on defining cloud computing. In November 2009, the first draft of a cloud computing definition was created. The NIST has recently published the 16th and final draft. According to the NIST definition, "cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction."<sup>[3]</sup>

The NIST states that the cloud model is composed of five essential characteristics, three service models, and four deployment models. <sup>[4]</sup>

### ***1.1 Essential Characteristics:***

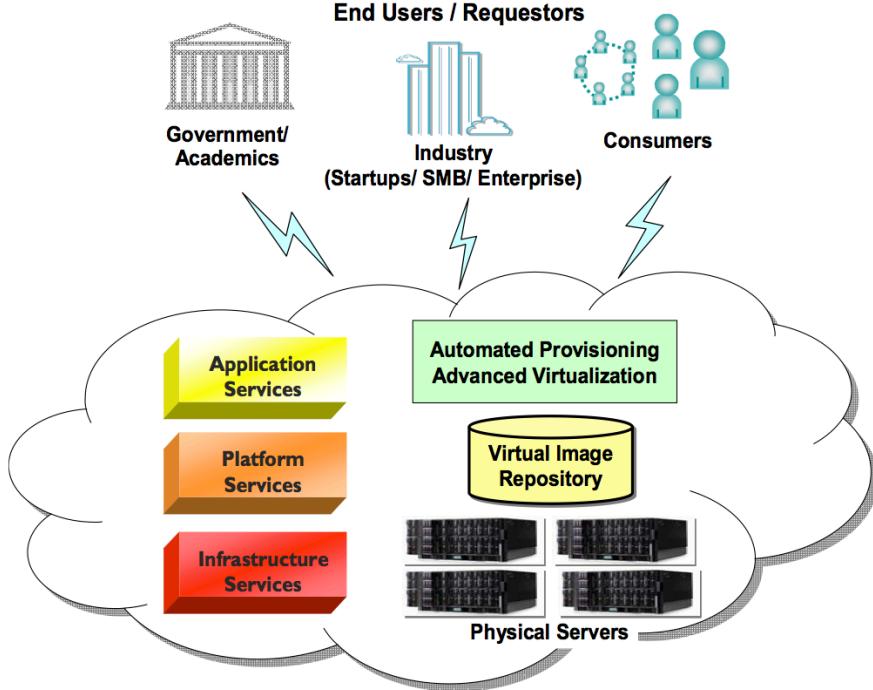
- 1.1.1** On-demand self-service: A consumer can be provided with computing capabilities, such as a server and storage, as needed without interaction with a cloud provider.
- 1.1.2** Broad network access: Capabilities are available over the network from anywhere via different platforms such as laptops, mobiles, PCs, etc.
- 1.1.3** Resource pooling: The resources of the cloud provider are pooled to serve customers.
- 1.1.4** Rapid elasticity: The resources can be elastically provisioned. Also, they appear to the consumer to be unlimited and can be provided at any time.
- 1.1.5** Measured service: The cloud provider and its consumer can monitor usage of the resources and control aspects such as storage, processing, active users, etc.

### ***1.2 Service Models:***

- 1.2.1** Cloud Software as a Service (SaaS): The capability of using the provider's applications running on a cloud infrastructure. The provider is responsible for managing the underlying cloud infrastructure.
- 1.2.2** Cloud Platform as a Service (PaaS): The capability given to the consumer for deploying acquired applications onto the cloud infrastructure using tools supported by the provider.
- 1.2.3** Cloud Infrastructure as a Service (IaaS): The capability given to the consumer to get storage, networks, and other computing resources that the consumer can deploy, which includes any software, whether operating system or application.

### ***1.3 Deployment Models:***

- 1.3.1** Private cloud: In this model, the infrastructure of the cloud can be only operated for an organization, and it can be managed by the organization itself or by a third party.
- 1.3.2** Community cloud: In this model, the cloud infrastructures are shared between certain organizations, which support a certain community that has shared concerns.
- 1.3.3** Public cloud: Here, the infrastructure is available to the public and is owned by a certain organization that sells services.
- 1.3.4** Hybrid cloud: Here in this model, two or more clouds (private, community, or public) compose the cloud infrastructure.



**Figure 1.1: Typical Cloud Computing Environment [2]**

As many organizations and industries are moving toward having their work done in the cloud, cloud computing becomes more essential day by day. The allure of the cloud is the significant value and benefits that people gain from it, such as reduced costs, increased storage, flexibility, and more mobility [1]. Flexibility is one of the major benefits that cloud computing can provide in terms of scaling up and down the infrastructure of a network. For example, when a server receives a glut of requests and cannot handle such demands, a new server will be added to help handle these requests. Chieu et. al. [2] presented a scaling scenario to address the dynamic scalability of web applications on a virtualized cloud computing environment. The scenario is based on using a load balancer to route user requests to web servers that provide web service. The load balancer routes incoming requests to servers that host the web application. The question is: what if the full complement of instances cannot handle more requests? Chieu et. al. proposed a dynamic scaling algorithm that the number of the instances should automatically scale based on the threshold on the number of current active sessions in each web server instance.

## **2. BACKGROUND**

Cloud computing history dates to the 1960s, when McCarthy posited his idea that “computation may someday be organized as a public utility” [5]. The concept of cloud computing was taken from telecommunications companies in the 1990s, making a radical shift from point-to-point data circuits to Virtual Private Network (VPN) services. The term “cloud computing” was used academically for the first time in 1997 by Professor Chellappa in a lecture titled “Intermediaries in Cloud-Computing”. He suggested that this would be a new *“computing paradigm where the boundaries of computing will be determined by economic rationale rather than technical limits alone.”*<sup>1</sup>

Companies began to move to the cloud as early as 1999. Salesforce.com was one of the first. The company introduced the concept of “delivering enterprise applications via a simple website” [5]. In 2002, Amazon.com launched Amazon Web Services [6]. Google Docs came next, in 2006<sup>2</sup>. Eucalyptus came in 2008, and was the first open source option for deploying a private cloud. Microsoft entered the cloud by launching Microsoft Azure in 2009. Today, there are many companies involved in cloud computing solutions, such as Oracle, Dell, IBM, Fujitsu, Teradata, and HP [5].<sup>3</sup>

## **3. LITERATURE REVIEW**

Mao et. al [7] presented a mechanism that dynamically scales cloud infrastructure up and down based on deadline and budget information. The proposed mechanism scales using a virtual machine (VM) that takes into consideration the performance and budget of the application. In terms of the performance, the mechanism provides enough VM instances to finish submitted jobs within the defined

---

<sup>1</sup> From his personal website: <http://www.bus.emory.edu/ram/>

<sup>2</sup> Google Docs belongs to a Software as a Service (SaaS) under the service models.

<sup>3</sup> Amazon Web Services & Microsoft Azure are considered cloud providers. They offer management consoles that provide all the service models mentioned above: SaaS, PaaS, IaaS.

deadline. In terms of budget, the mechanism runs VM types<sup>4</sup> based on the applications' needs. For example, if the submitted job does not need a huge processor, the mechanism will provide a small VM type instead of a large one, and vice versa. This mechanism could cut costs when compared to using only one type of instance.

Cushing et. al. [8] proposed a method for auto scaling data-centric workflow tasks using a prediction-based approach. The authors achieved scaling using a prediction mechanism in which the input data load on a workflow's task is used to compute the estimated time for execution of the task. Through this prediction, a framework can make decisions to scale multiple tasks independently to improve output and reduce congestion.

Dean. et. al [9] developed a programming model called MapReduce which more closely relates to this project. MapReduce is a software framework—introduced by Google—which supports distributed computing on large data sets on computer clusters [18]. MapReduce processes data that is distributed across servers in a cluster using three operations [19]. The first is called Map, which is a set of tasks that process in parallel by each node within a cluster separated from other nodes within that cluster. Users determine a map function, which processes a key/value pair to generate a set of intermediate key/pair values [9]. Every map task is assigned a part of the input file that is called a split. Each split has a single HDFS block by default [20]. The second involves data that is distributed across all nodes within the cluster, and the third, known as reduce, is a set of tasks that each node executes in parallel. Here, the values that are associated with same key are merged.

Venugopal et. al [8] used features of the Session Initiation Protocol (SIP), in addition to using the Amazon Elastic Compute Cloud service (EC2), to implement a mechanism which scales a VoIP-based call center to respond to emergency calls.

---

<sup>4</sup> For example, Amazon Elastic Compute instances (EC2) are grouped into seven groups of types [13]: Standard, Micro, High-Memory, High-CPU, Cluster Compute, Cluster GPU, and High I/O.

They developed a control system that dynamically scales up or down according to the increasing or decreasing call volumes. Their results show that the control system was able to respond to increasing call volumes by providing additional servers as needed.

#### 4. MATERIALS & METHODS:

The algorithm proposed (see figure 4.1) in the paper titled “Dynamic Scaling of Web Applications in a Virtualized Cloud Computing Environment” written by Chieu et. al, [2] uses the number of active sessions in each server to scale up and down as a threshold. However, in this thesis, the same algorithm has been implemented on an Amazon Web Services cloud using different thresholds (metrics) to scale up or down.

```

For an instance i in NInstance
    If (Ai/SMax >= TUpper) then
        Increment NExceed
    If (Ai/SMax < TLower) then
        Increment NBelow
    Record and sort all indexes J in ascending of Ai/SMax
If (NExceed == NInstance) then
    Provision and start a new instance
    Increment number of instances: NInstance
    Add new instance to Load-Balancer
If (NBelow >= 2) then
    Set m equal first index in J
    If (Am == 0) then
        Remove instance m from Load-Balancer
        Shutdown instance m
        Decrement number of instances: NInstance
        Decrement NBelow and remove index m from J
    If (NBelow >= 2) then
        Set n equal first index in J
        Remove instance n temporary
        Decrement number of instances: NInstance
        Set normalized load factor Ln = 0
For an instance i in NInstance
    Evaluate normalized load factor:
    Li =  $(1 - A_i/S_{Max}) / \sum_{k=1, N_{Instance}}^{N_{Instance}} [1 - A_k/S_{Max}]$ 
    Apply new load factors Li to Load Balancer

```

where  $A_i$ : Number of active sessions in instance  $i$   
 $S_{Max}$ : Maximum sessions per instance (e.g. 40,000)  
 $T_{Upper}$ : Session upper-threshold (e.g. 80% or 0.8)  
 $T_{Lower}$ : Session lower-threshold (e.g. 60% or 0.6)  
 $N_{Instance}$ : Number of existing instances  
 $N_{Exceed}$ : Number of instances exceeding session upper-threshold  
 $N_{Below}$ : Number of instances below session lower-threshold

**Figure 4.1: Dynamic Scaling Algorithm for Virtual Machine Instances in the Cloud. [1]**

The CPU utilization metric will be used as an example for how to set auto scaling for practical purposes (see section 4.1.4), however the actual experiment will be done without auto scaling action and will compare several metrics to decide which is the most effective. These tests (explained below in the Tests' Results Section) would be performed through various substantial analyses.

When debating which metrics to use, those of the load balancer should also be taken into consideration as well as the EC2 Instances. Thus, a Network Out metric would be used as a metric of the EC2 instances besides the CPU Utilization metric. In regards to the load balancer metrics, two metrics were chosen: Latency and Request Count.

The explanation of each metric is as follows:

#### ***4.1 Metrics for the EC2 instances:***

##### ***4.1.1 CPU Utilization Metric:***

As explained earlier, this is the percentage of an “allocated EC2 [which] computes units that are currently in use on the instance. This metric identifies the processing power required to run an application upon a selected instance.”<sup>5</sup>

---

<sup>5</sup> Amazon Web Services documentations

#### **4.1.2 Network Out Metric:**

“The number of bytes sent out on all network interfaces by the instance. This metric identifies the volume of outgoing network traffic to an application on a single instance.”<sup>3</sup>

#### **4.2 Metrics for the load balancer:**

##### **4.2.1 Latency Metric:**

“Time elapsed after the load balancer receives a request until it receives the corresponding response.”<sup>3</sup>

##### **4.2.2 Request Count Metric:**

“The number of requests handled by the load balancer.”<sup>3</sup>

Due to the huge variety of applications, website workloads, and web servers type, there is no universal policy for when a certain metric should be scaled. The policies featured below in Table 1 were created through tests using an experimental website, T1.micro type:

Metric	Policy of Scaling Up	Policy of Scaling Down
CPU Utilization	CPU Utilization $\geq 80\%$ for 5 minutes	CPU Utilization $\leq 20\%$ for 5 minutes
Network Out	Network Out $\geq 2\text{ GB}$ for 5 minutes	Network Out $\leq 5\text{ MB}$ for 5 minutes
Latency	Latency $\geq 0.3\text{ second}$ for 5 minutes	Latency $\leq 0.002\text{ second}$ for 5 minutes
Request Count	Request Counts $\geq 150$ requests for 5 minutes	Request Counts $\leq 2$ requests for 5 minutes

Table 1: Policies of scaling up and down for all tested metrics

State	Name	Threshold
🔴 ALARM	Low-CPU-Utilization	CPUUtilization <= 20 for 5 minutes
🔴 ALARM	Low_RequestCount	RequestCount <= 2 for 5 minutes
🟢 OK	Low_NetworkOut	NetworkOut <= 5,000,000 for 5 minutes
🟢 OK	High_RequestCount	RequestCount >= 150 for 5 minutes
🟢 OK	Low_Latency	Latency <= 0.002 for 5 minutes
🟢 OK	High_Latency	Latency >= .3 for 5 minutes
🟢 OK	High_CPU_Utilization	CPUUtilization >= 80 for 5 minutes
🟢 OK	High_NetworkOut	NetworkOut >= 2.0E9 for 5 minutes

**Figure 4.1.1: Policies of scaling up and down for all tested metrics as shown in AWS management console**

Since the metrics differ in how they reach their upper and lower boundaries, two tests had to be performed to ensure proper results. The first test, designed specifically for the EC2 instances metrics (CPU Utilization & Network Out), involves three users running a bash script that downloading a gallery of photos (sized 1.6 GB) from the T1Micro server. The script has two loops: the outer loop runs 5 times, with the inner loop running twice per outer loop. When the inner loop runs, it starts by downloading the gallery. When the download is finished, it stops for two minutes, before repeating the process. When the loop is completed, it stops for seven minutes, after which the second loop of the outer loop would start. This test, called FIRST TEST, is in the results section. The bash scrip is as follows:

```

#!/bin/bash

START=$(date +%s.%N)

command

counting=0;

for i in {1..5}

do

for j in { 1..2}

do

wget myloadbalancer-1942790311.us-east-
1.elb.amazonaws.com/Photos.gz;

sleep 2m;

let "counting += 1";

done

sleep 7m;

echo "$counting of galleries have been downloaded";

done

END=$(date +%s.%N)

DIFF=$(echo "$END - $START" | bc)

echo $DIFF

```

The second test, coded for the Load Balancer metrics (Latency & Request Count), involves a bash script that downloads the index page (3KB). The bash script also has two loops, with the outer loop running 100 times, and the inner running 30 times per outer loop. When the inner loop runs, it downloads the index page thirty times, with a one-second waiting period between each connection. When completed, it stops for 15 seconds before the second loop of the outer loop starts. This continues until the end of the script. This test is called SECOND TEST, and can also be found in the results sections. Below is the bash script:

```

#!/bin/bash

START=$(date +%s.%N)

command

counting=0;

for i in {1..100}

do

for j in {1..30}

do

wget myloadbalancer-1942790311.us-east-
1.elb.amazonaws.com/index.html;

sleep 1;

let "counting += 1";

done

echo "$counting times I connected to the
loadbalancer";

sleep 15;

done

END=$(date +%s.%N)

DIFF=$(echo "$END - $START" | bc)

echo $DIFF

```

#### ***4.3 The Repetition of the Tests:***

The tests were conducted multiple times in different times and places to better compare the metrics. Five different locations were chosen. The first was in a home setting, since this is where most people access the Internet. The second was in the Systems Lab in Golisano College of Computing and Information Sciences at Rochester Institute of Technology. This location offered vastly different results than

the first, due to more network bandwidth and closer virtual proximity to the web server in Virginia. The third location was within the Amazon datacenter in the Virginia region to see if results would differ if the web server and the tester machine were within the same physical place. The fourth and fifth locations of testing were at Amazon cloud datacenter in the Oregon and North California state region, respectively.

For the first two locations, the tests were run at these times:

- Weekdays in the morning
- Weekdays in the evening
- Weekends in the morning
- Weekends in the evening

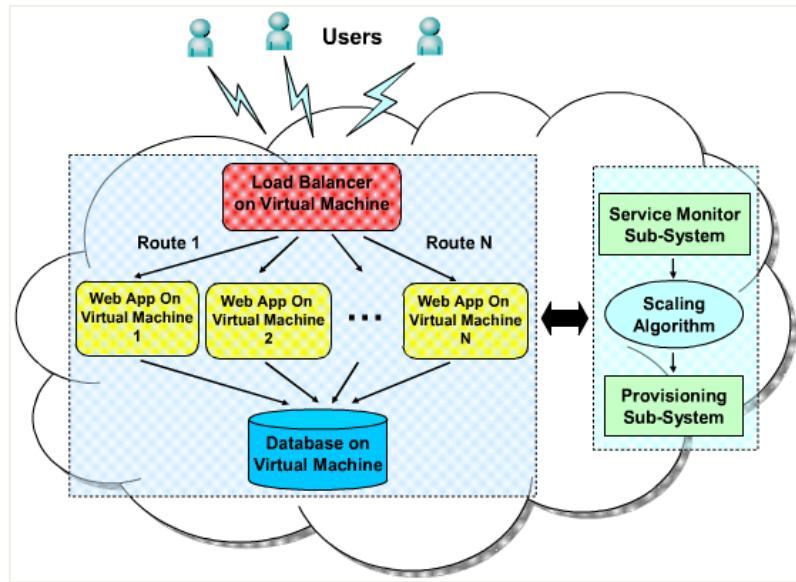
The purpose of these times was to see how the results would vary based on Internet traffic. For example, would the results taken from weekday mornings, when there is less Internet traffic, be strikingly different than results taken during weekend evenings, when there is high Internet traffic?

*The tests were both run eleven times as follows:*

- 1- Weekday mornings at RIT System Lab
- 2- Weekday evenings at RIT System Lab
- 3- Weekend mornings at RIT System Lab
- 4- Weekend evenings at RIT System Lab
- 5- Weekday mornings from home
- 6- Weekday evenings from home
- 7- Weekend mornings from home

- 8- Weekend evenings from home
- 9- From a server within Amazon cloud in the same geographic region of the Virginia web server
- 10- From a server within Amazon cloud in the Oregon geographic region
- 11- From a server within Amazon cloud in the North California geographic region

#### **4.4 The project's architecture:**



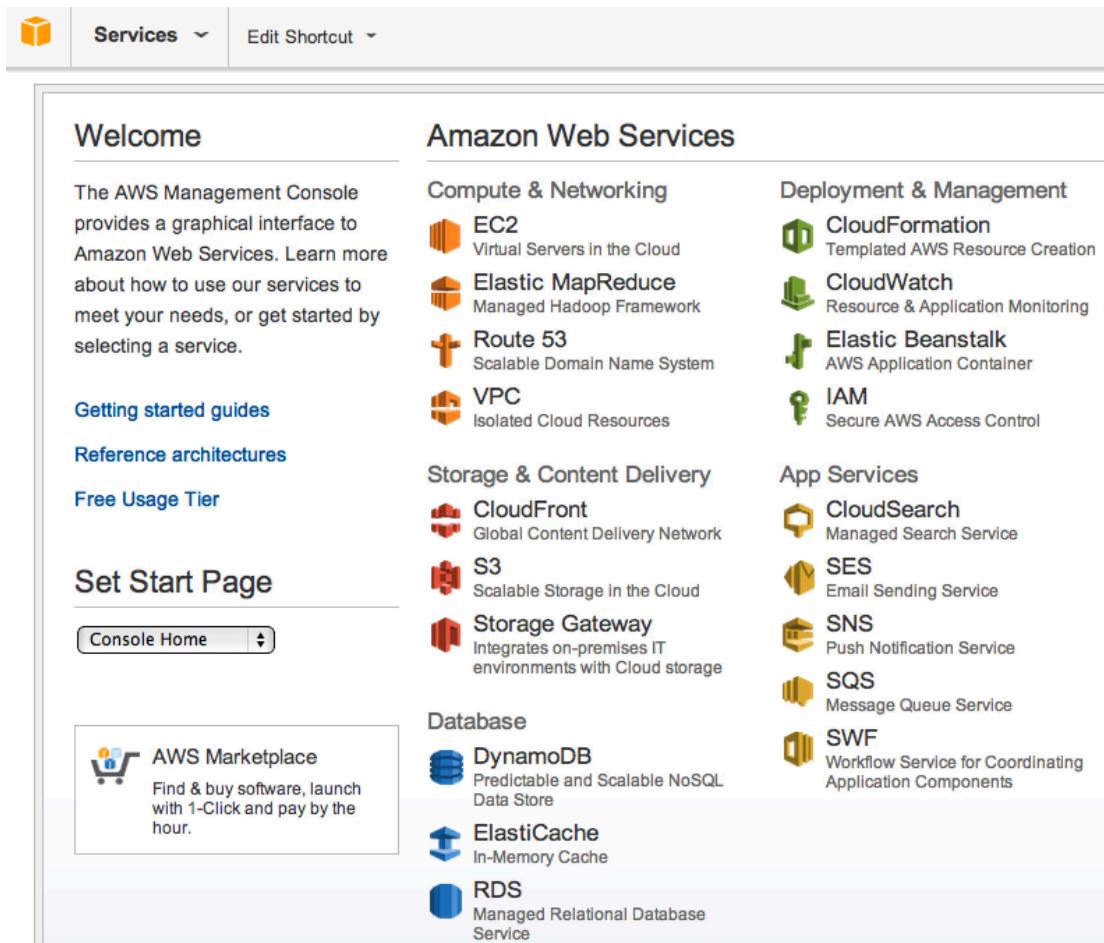
**Figure 4.2: Architecture to Scale Web Applications in a Cloud [1]**

In this thesis, some of Amazon Web Services have been used such as Amazon Elastic Compute Cloud instances (EC2), Amazon AMIs, Amazon Load Balancer, Amazon Auto Scaling, and Amazon CloudWatch.

##### **4.4.1 Launching Amazon EC2 instances:**

A web server (instance) was launched on Amazon Cloud with micro type using 613 MB memory and a 64-bit platform for the purpose of reaching the defined CPU utilization easily. The following steps were used to launch the Amazon EC2 Windows instance:

- Signed in to Amazon Web Services Management Console, and clicked on EC2 (see figure 4.3):



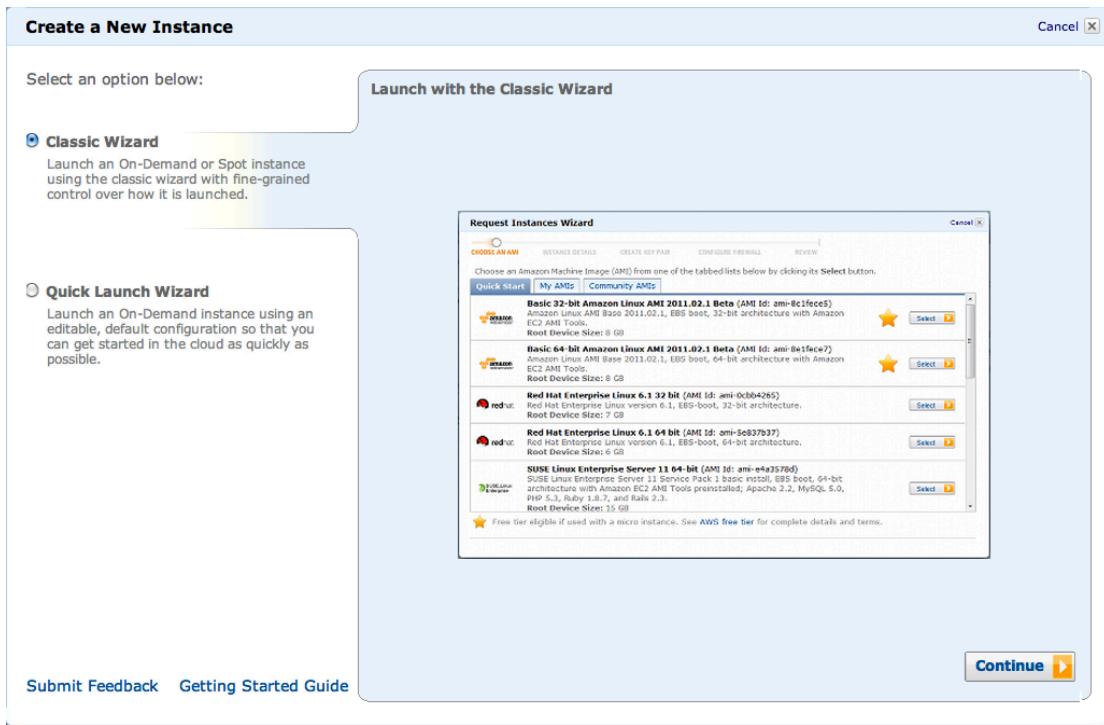
**Figure 4.3: Amazon Web Services Management Console**

- Clicked “Launch Instance” from the Amazon EC2 Console Dashboard (figure 4.4):



**Figure 4.4: Amazon EC2 Console Dashboard**

- When clicking on Launch Instance, a “Create a new instance” page was opened, as in this figure 4.5:



**Figure4.5:** Create a New Instance page

In this page, there were two ways to launch an instance:

- **Classic Wizard**, which offered more control and advance settings that might be wanted.
  - **Quick Launch Wizard**, which simplified the process to launch an instance and also automatically configured many selections. In this project, Classic Wizard was used, as it allowed the configuration of some specific choices.
- In the Quick Start tab, Microsoft Windows Server 2008 R2 Base was chosen (see figure 4.6).

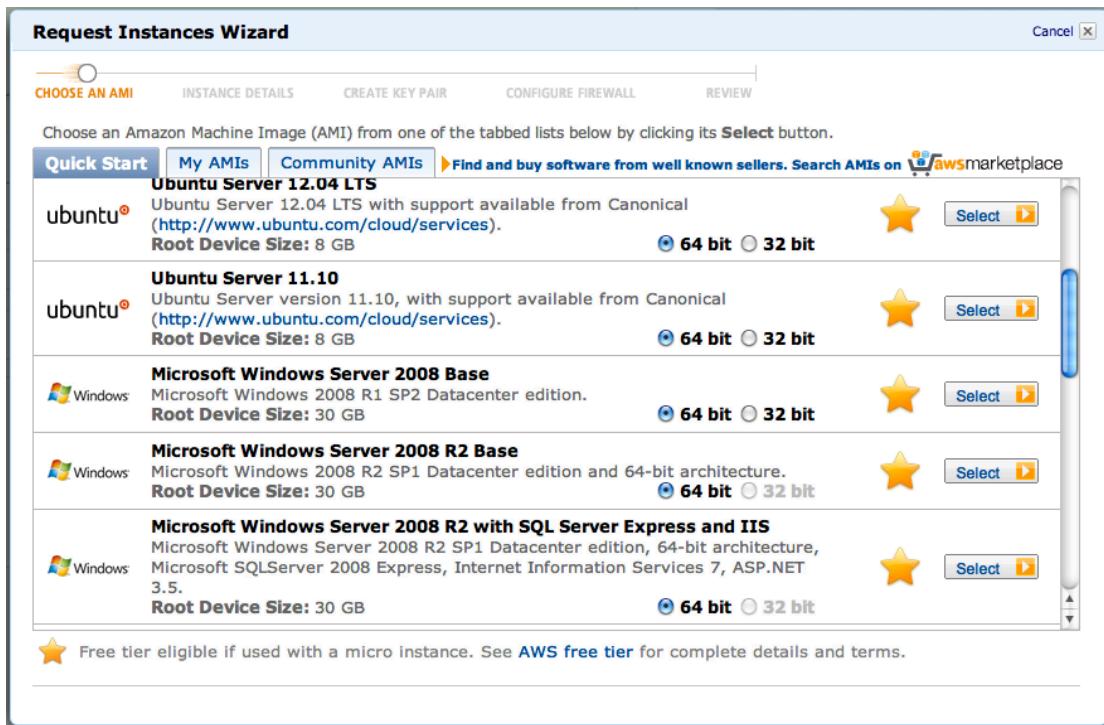


Figure 4.6: Choose the AMI from the Quick Start tap

- Micro type was selected (figure 4.7):

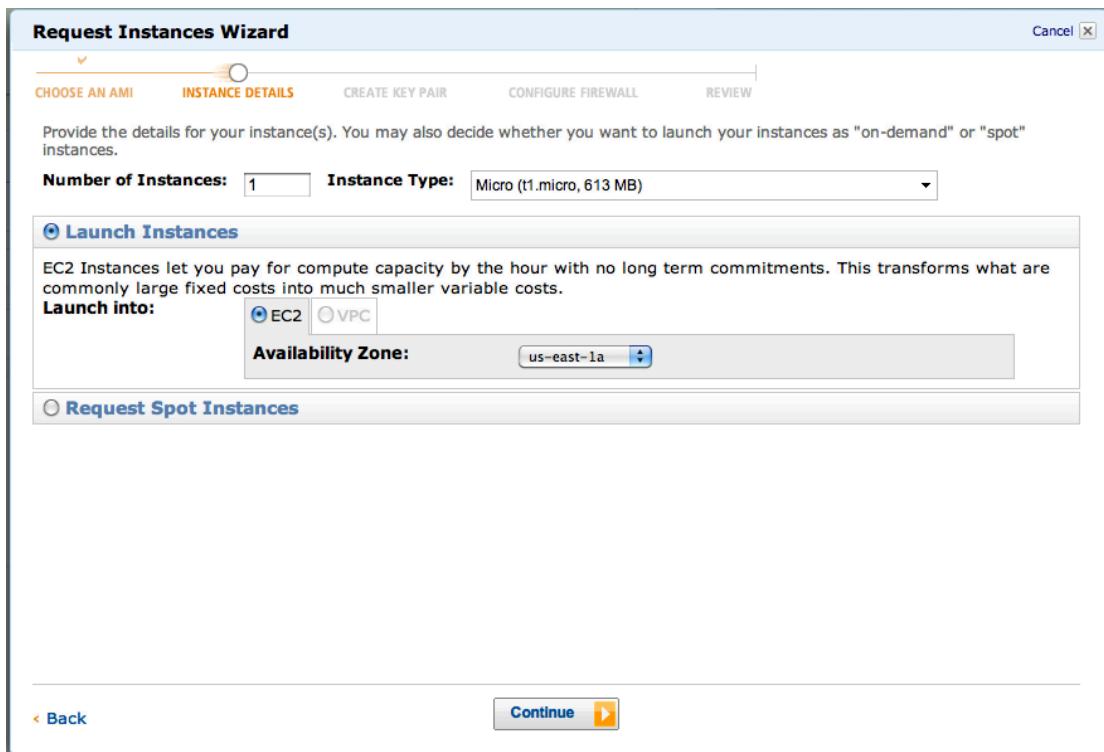


Figure 4.7: Choosing the number and the type of the instance

- The name of the web server and a short description were entered (figure 4.8):

**Request Instances Wizard**

CHOOSE AN AMI    INSTANCE DETAILS    CREATE KEY PAIR    CONFIGURE FIREWALL    REVIEW

Add tags to your instance to simplify the administration of your EC2 infrastructure. A form of metadata, tags consist of a case-sensitive key/value pair, are stored in the cloud and are private to your account. You can create user-friendly names that help you organize, search, and browse your resources. For example, you could define a tag with key = Name and value = Webserver. You can add up to 10 unique keys to each instance along with an optional value for each key. For more information, go to [Using Tags in the EC2 User Guide](#).

Key (127 characters maximum)	Value (255 characters maximum)	Remove
Name	Webserver	X
Purpose	To create an AMI instance	X
		X

Add another Tag. (Maximum of 10)

< Back    Continue ➔

Figure 4.8: Naming the instance

- A key pair was created, securing the connection to the instance (figure 4.9):

**Request Instances Wizard**

CHOOSE AN AMI    INSTANCE DETAILS    CREATE KEY PAIR    CONFIGURE FIREWALL    REVIEW

Public/private key pairs allow you to securely connect to your instance after it launches. To create a key pair, enter a name and click **Create & Download your Key Pair**. You will then be prompted to save the private key to your computer. Note, you only need to generate a key pair once - not each time you want to deploy an Amazon EC2 instance.

Choose from your existing Key Pairs

Create a new Key Pair

1. Enter a name for your key pair\*:  (e.g., jdoekey)
2. Click to create your key pair: [Create & Download your Key Pair](#)

Save this file in a place you will remember.  
You can use this key pair to launch other instances in the future or visit the Key Pairs page to create or manage existing ones.

Proceed without a Key Pair

< Back    Continue ➔

Figure 4.9: Key pair creation

- Then, a security group was created, defining the connection ports to the machine. 80 (HTTP) and 443 (HTTPS) ports were added (see figure 4.10):

**Request Instances Wizard**

CHOOSE AN AMI   INSTANCE DETAILS   CREATE KEY PAIR   **CONFIGURE FIREWALL**   REVIEW

Security groups determine whether a network port is open or blocked on your instances. You may use an existing security group, or we can help you create a new security group to allow access to your instances using the suggested ports below. Add additional ports now or update your security group anytime using the Security Groups page.

Choose one or more of your existing Security Groups  
 Create a new Security Group

**Group Name:** Final  
**Group Description:** Mohammed Aljebreen's

**Inbound Rules:**

Create a new **Custom TCP rule**  
 rule:  
 Port range:   
 (e.g., 80 or 49152-65535)  
 Source:   
 (e.g., 192.168.2.0/24, sg-47ad482e, or 1234567890/default)

**TCP**

Port (Service)	Source	Action
80 (HTTP)	0.0.0.0/0	Delete

**Add Rule**

**Continue**

Figure 4.10: Security Group Creation

- The web server launched and took a few minutes to initialize (see figures 4.11, 4.12):

Name	Instance	AMI ID	Root Device	Type	State	Status Checks
Webserver (1)	i-5478542d	ami-4c309225	ebs	t1.micro	running	2/2 checks passed

Figure 4.11: The web server instance

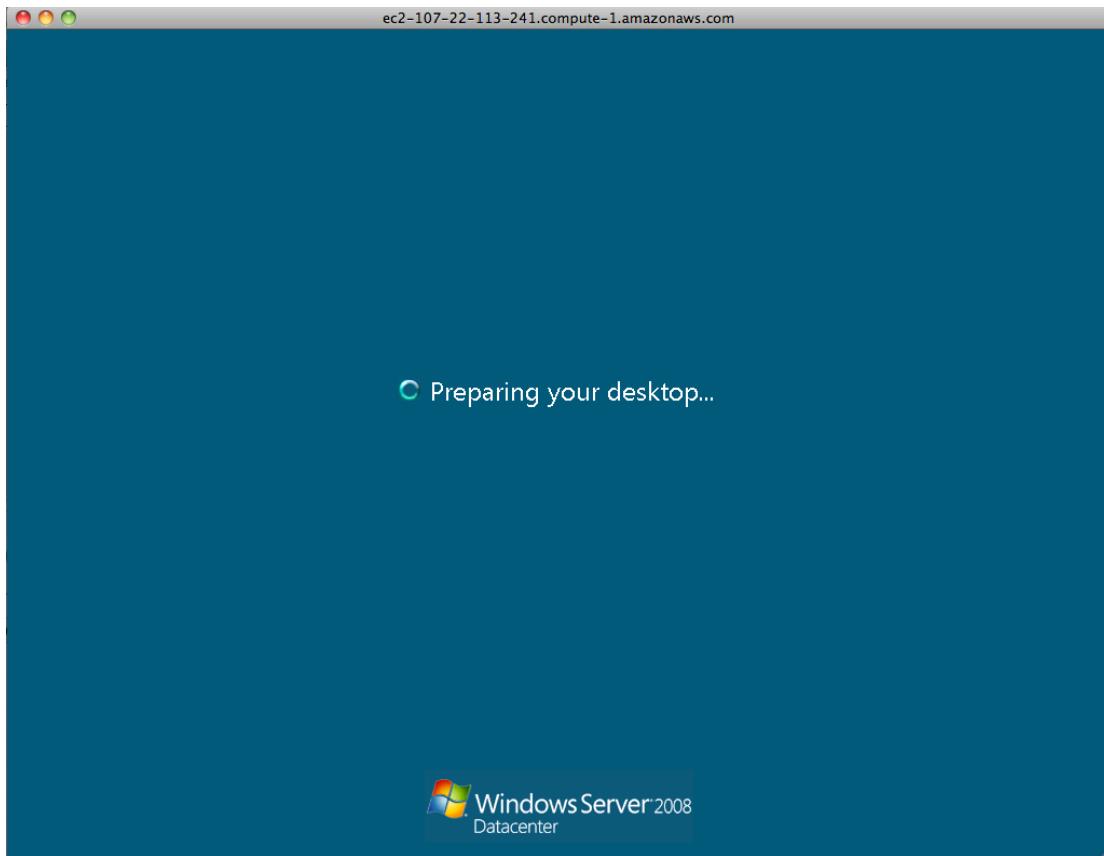


Figure 4.12: the web server started

- Apache web server, version 2.0.64 for Win32 platform was installed. The server's services included Apache web services (see figure 4.13):

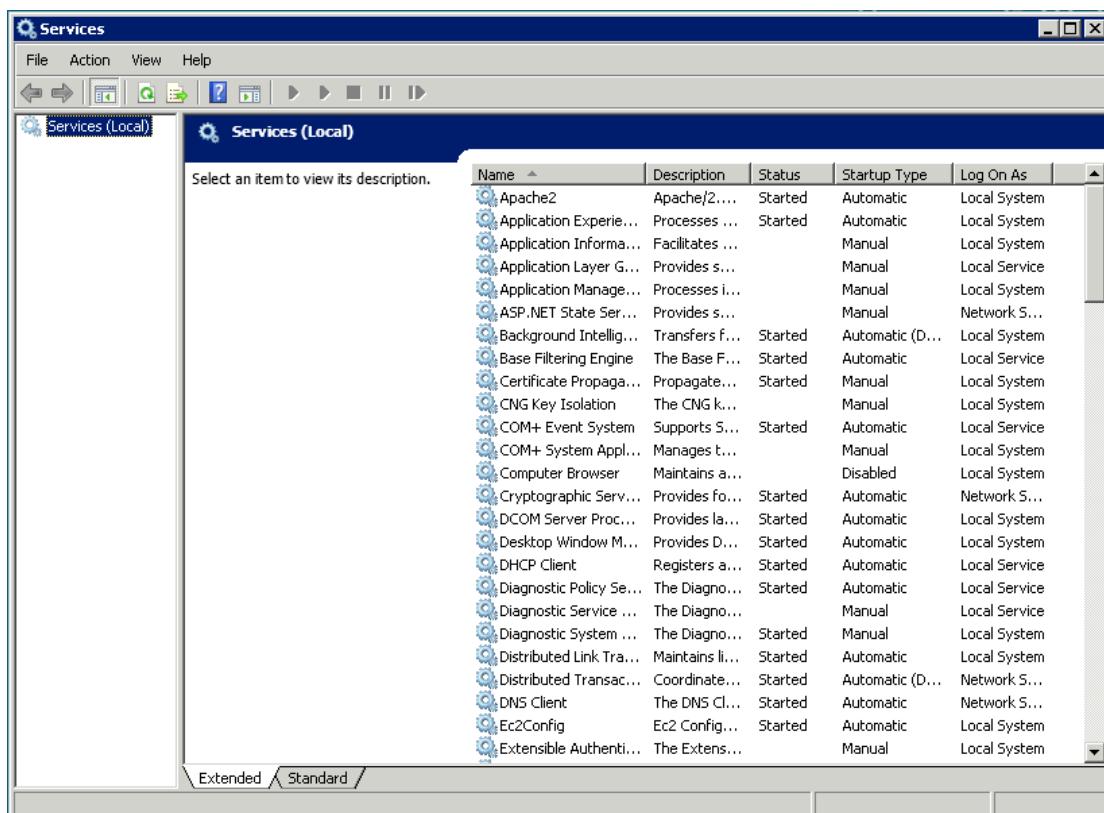
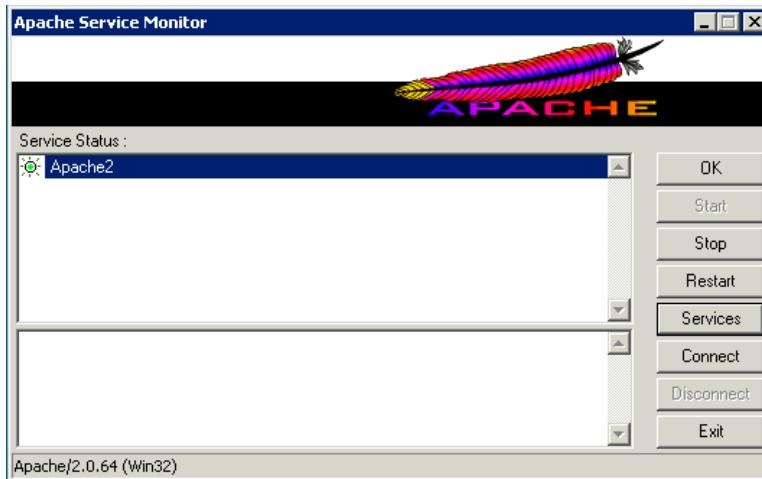


Figure 4.13: Apache2 included in the instance services

- Here is the Apache Service Monitor screen (figure 4.14):



**Figure 4.14: Apache Service Monitor**

After launching a web server and installing all the needed services, a simple web site<sup>6</sup> was designed, presenting services to the public. The page source code is:

```
<title>MOHAMMED ALJEBREEN</title>

<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
<html xmlns="http://www.w3.org/1999/xhtml">
<head>
    <meta name="author" content="Wink Hosting (www.winkhosting.com) " />
    <meta http-equiv="Content-Type" content="text/html; charset=iso-8859-
1" />
    <link rel="stylesheet" href="images/style.css" type="text/css" />
    <title>Webserver</title>
</head>
<body>

    <div id="page" align="center">
        <div id="content" style="width:800px">
            <div id="logo">
                <div style="margin-top:70px"
class="whitetitle">R.I.T</div>
            </div>
            <div id="topheader">
                <div align="left" class="bodytext">
                    <br />
                    <strong>Mohammed Aljebreen </strong><br />
                    475 Countess Dr<br />
                    West Henrietta<br />
                    Phone: (585)309-3487<br />
                    mja7175@erit.edu
                </div>
                <div id="toplanks" class="smallgraytext">
                    <a href="#">Home</a> | <a
href="#">Sitemap</a> | <a href="#">Contact Us</a>
                </div>
            </div>
            <div id="menu">
                <div align="right" class="smallwhitetext"
style="padding:9px;">
```

---

<sup>6</sup> The design from <http://www.w3.org>

```

        <a href="#">Home</a> | <a href="#">About Us</a> | <a href="#">Products</a> | <a href="#">Our Services</a> | <a href="#">Contact Us</a>
    </div>
</div>

<div id="contenttext">
    <div style="padding:20px">
        <span class="titletext">Final Project!</span>
    </div>
    <div class="bodytext" style="padding:12px;" align="justify">
        <strong>Hi! This is Mohammed Aljebreen. I'm going to present my final project which is about Implementing a Dynamic Scaling of Web Applications in a Virtualized Cloud Computing Environment.
        <br />
    <br />

    The committee members:<br />
    Sharon Mason (chair)<br />
    Lawrence Hill (committee member)<br />
    Jim Leone (committee member)<br />
</div>
</div>

<span class="titletext">Permanently Server</span>
<div style="padding:75px">

    <div class="bodytext" style="padding:1px;" align="middle">
        <strong>Here are three different galleries of pictures: </strong> <br />
    </div>
    <table align="center">

        <tr>
            <td></td>
            <td></td>
            <td></td>
        </tr>

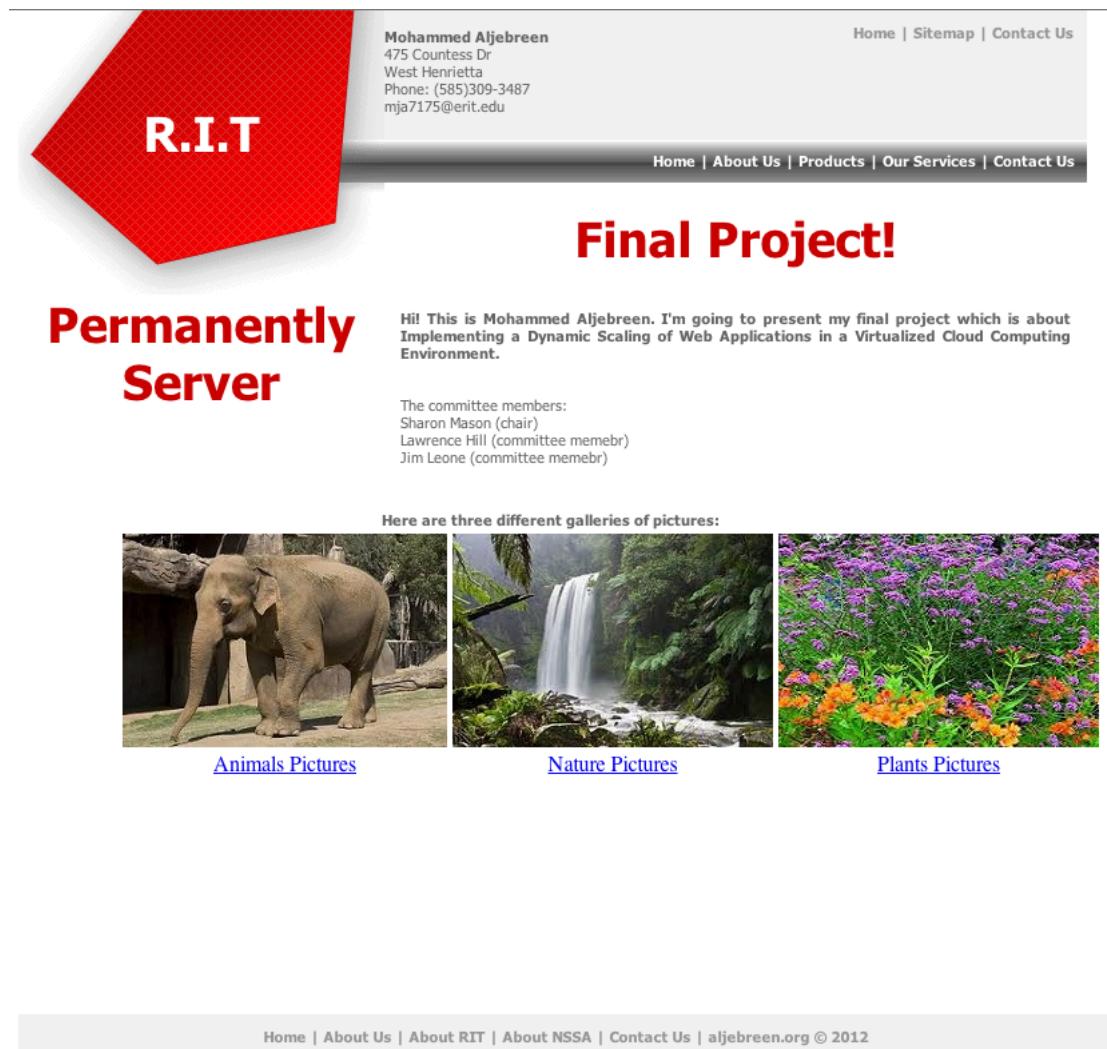
        <tr>
            <td align="center"><a href="Photos.gz" target="_blank">Animals Pictures</a></td>
            <td align="center"><a href="Photos1.gz" target="_blank">Nature Pictures</a></td>
            <td align="center"><a href="Photos2.gz" target="_blank">Plants Pictures</a></td>
        </tr>
    </table>
</div>

    <div style="padding:50px">
        </div>
        <div id="footer" class="smallgraytext">
            <a href="#">Home</a> | <a href="#">About Us</a> |
            <a href="#">About RIT</a> | <a href="#">About NSSA</a> | <a href="#">Contact Us</a>
            | aljebreen.org
            &copy; 2012 <a href="http://www.winkhosting.com" target="_blank"></a>
        </div>
    </div>
</body>
</html>

```

- The home web page of the website (see figure 4.15)

(<http://loadbalancer.aljebreen.org>):



**Figure 4.15: The home page of the website**

The purpose of this web page was to present services to the users. The website will be used to test the proposed algorithm when applying a lot of load traffic to that website.

#### ***4.4.2 Creating an Amazon Machine Image (AMI):***

What is the AMI? As defined by the Amazon Web Services website, “an Amazon Machine Image (AMI) is a special type of pre-configured operating system and virtual application software which is used to create a virtual machine within the Amazon Elastic Compute Cloud (EC2). It serves as the basic unit of deployment for

services delivered using EC2” [10]. It is like a template of a computer’s root volume [11]. For example, an AMI could act as a web server or other data base server. An AMI was created from the previously made instance in order to have an image of the webserver that would be needed to scale up and down later when auto scaling takes action.

- The created AMI, named “clone”:

	Name	AMI ID	Source	Owner	Visibility	Status	Platform
<input type="checkbox"/>	Clone	 ami-1975da70	131511199887/Clone	131511199887	Private	 available	 Windows

Figure4.16: The AMI from the web server

#### ***4.4.3 Creating an Amazon Load Balancer:***

Amazon Elastic Load Balancing (see figure 4.17) consists of two elements: the load balancer and the controller service. The purpose of the load balancer is to monitor the traffic and handle the requests coming from the Internet. The purpose of the controller service is to monitor the load balancer, verify its proper functioning, and add and remove its capacity as needed.

The Elastic Load Balancing service on Amazon cloud will be used because it has many useful features [12]:

- Distributing incoming traffic across Amazon EC2 instances.
- Automatically scaling the request in response to the incoming traffic.
- Creating and managing security groups to provide more networking and security options when the Elastic Load Balancing is used in Virtual Private Cloud.

- Detecting the health<sup>7</sup> of the instances. Once an unhealthy load balance instance has been detected, it will cease to route the traffic, and the rest of the healthy instances will take care of routing instead.

	Load Balancer Name	DNS Name	Port Configuration
<input type="checkbox"/>	Webserver-Load-Balance	Webserver-Load-Balancer-1342417435.us-eas	80 (HTTP) forwarding to 80 (HTTP), 443 (HTTP) forwarding to 443 (HTTPS)

Figure 4.17: Web server Load Balancer

- Below, two servers are running under the load balancer (figure 4.18):

Instances				
Instance	Name	Availability Zone	Status	Actions
i-5478542d	Webserver (1)	us-east-1a	In Service	<a href="#">Remove from Load Balancer</a>
i-cedeb1b6	empty	us-east-1a	In Service	<a href="#">Remove from Load Balancer</a>
i-24d0f85d	Webserver (2)	us-east-1a	Out of Service ( <a href="#">why?</a> )	<a href="#">Remove from Load Balancer</a>

Availability Zones			
Availability Zone	Instance Count	Healthy?	Actions
us-east-1a	3	Yes	-

Figure 4.18: The servers under the load balancer

- Webserver (1) and webserver (2) were added manually to the load balancer (see figure 4.19); however, any new server will be added automatically.

Select	Instance	Name	State	Security Groups	Availability Zone
<input type="checkbox"/>	i-50d6bd29	DC (ARS7919)	● stopped	DefaultG	us-east-1a
<input type="checkbox"/>	i-2a0d7d53	TMG (ARS7919)	● stopped	DefaultG	us-east-1a
<input type="checkbox"/>	i-b2773bcb	Webserver (temporary)	● running	Final	us-east-1a
<input checked="" type="checkbox"/>	i-5478542d	Webserver (1)	● running	Final	us-east-1a
<input type="checkbox"/>	i-84200efd	ahmed	● stopped	WindowsCB	us-east-1a
<input type="checkbox"/>	i-7a391003	ahmed2	● stopped	WindowsCB	us-east-1d
<input checked="" type="checkbox"/>	i-24d0f85d	Webserver (2)	● running	Final	us-east-1a

Figure 4.19: Adding Webserver (1) & (2) to the Load Balancer

---

<sup>7</sup> The load balancer automatically performs health checks on EC2 instances using the protocol, URL, timeout, and interval specified when configuring load balancer. It routes traffic to only the instances that pass the health check. If an instance fails the health check, it is automatically removed from the load balancer [13].

- In figure 4.20, HTTP is chosen as the ping protocol, and 80 as the ping port to allow the load balancer to receive incoming traffic.

**Ping Protocol:** HTTP

**Ping Port:** 80

**Ping Path:** /index.html

**Advanced Options:**

**Response Timeout:** 5 Seconds

**Health Check Interval:** 0.5 Minutes

**Unhealthy Threshold:**  2

**Healthy Threshold:**  10

Figure 4.20: Advanced Options of the load balancer's configuration

- As the above figure shows, the configurations are set as follows:
  - Ping Protocol: HTTP
  - Ping Port: 80
  - Ping Path: /index.html
  - Response Timeout: 5 seconds, which is wait time before receiving a response from the health check.
  - Health Check Interval: 0.5 minutes, which is the amount of time between health check.
  - Unhealthy Threshold: 2 checks failures, which are the number of health check failures before an EC2 instance is declared unhealthy.
  - Healthy Threshold: 10 succeed checks, which is the number of health check successes before an EC2 instance is declared healthy.

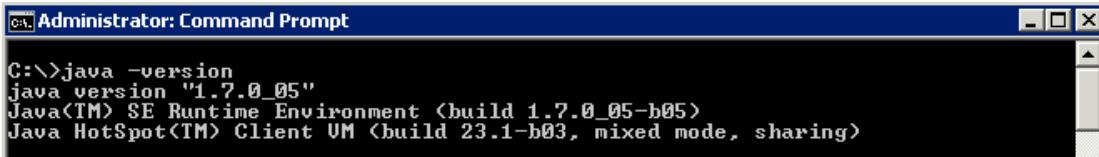
#### **4.4.4 Configuring Amazon Auto Scaling:**

The auto scaling was created using the CPU Utilization metric as an example.

The policies were set to:

- Create a new temporary webserver when the CPU utilization of the web server exceeded 80% for five minutes.
- Remove the temporary webserver when the CPU utilization of web server dropped less than 20% for five minutes.

Before creating an Auto Scaling, the proper command line tools must be downloaded from the Amazon Web Services site. Java version 1.7 (see figure 4.21) must also be installed on the instance used to create auto scaling.



```
C:\>java -version
java version "1.7.0_05"
Java(TM) SE Runtime Environment (build 1.7.0_05-b05)
Java HotSpot(TM) Client VM (build 23.1-b03, mixed mode, sharing)
```

**Figure 4.21: Java version 1.7 installed on the web server**

- The JAVA\_HOME environment variable was set to point to Java installation via this command:

```
C:\> set JAVA_HOME=C:\Program Files\Java\jre7
```

- The %JAVA\_HOME%\bin to %PATH environment variable was added:

```
C:\> set PATH=%PATH%;%JAVA_HOME%\bin
```

- The AWS\_AUTO\_SCALING\_HOME environment variable was set where the Auto Scaling folder was unpacked:

```
C:\> set AWS_AUTO_SCALING_HOME=
C:\Users\Administrator\Desktop\AutoScaling-2011-01-01\AutoScaling-
1.0.61.0
```

- The %AWS\_AUTO\_SCALING\_HOME%\bin to %PATH environment variable was added:

```
C:\> set PATH=%PATH%;%AWS_AUTO_SCALING_HOME%\bin
```

- The tools for the security credentials which were saved in a .template file were configured. The file contained two security credentials: AWS access key ID, and AWS secret key.

```
C:\>set AWS_CREDENTIAL_FILE=C:\Users\Administrator\Desktop\  
AutoScaling-2011-01-01\AutoScaling-1.0.61.0
```

- The configuration was tested by using `as-cmd` command, and the resulting output (see figure 4.22), confirmed the installation:

Command Name	Description
<code>as-create-auto-scaling-group</code>	Create a new Auto Scaling group.
<code>as-create-launch-config</code>	Creates a new launch configuration.
<code>as-create-or-update-tags</code>	Create or update tags.
<code>as-delete-auto-scaling-group</code>	Deletes the specified Auto Scaling group.
<code>as-delete-launch-config</code>	Deletes the specified launch configuration.
<code>as-delete-notification-configuration</code>	Deletes the specified notification configuration.
<code>as-delete-policy</code>	Deletes the specified policy.
<code>as-delete-scheduled-action</code>	Deletes the specified scheduled action.
<code>as-delete-tags</code>	Delete the specified tags.
<code>as-describe-adjustment-types</code>	Describes all policy adjustment types.
<code>as-describe-auto-scaling-groups</code>	Describes the specified Auto Scaling groups.
<code>as-describe-auto-scaling-instances</code>	Describes the specified Auto Scaling instances.
<code>as-describe-auto-scaling-notification-types</code>	Describes all Auto Scaling notification types.

Figure 4.22: Testing the Auto Scaling command line installation

- After testing the Auto Scaling installation and configuration, a launch config (called MyAutoConfig), was created (see figure 4.23):

```
as-create-launch-config MyAutoConfig --image-id ami-1975da70 --  
instance-type t1.micro --access-key-id AKIAIQS44P5ZBP244KQA --  
secret-key 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
```

- `--image-id: ami-1975da70`

The AMI's ID, which was created earlier, called Clone.

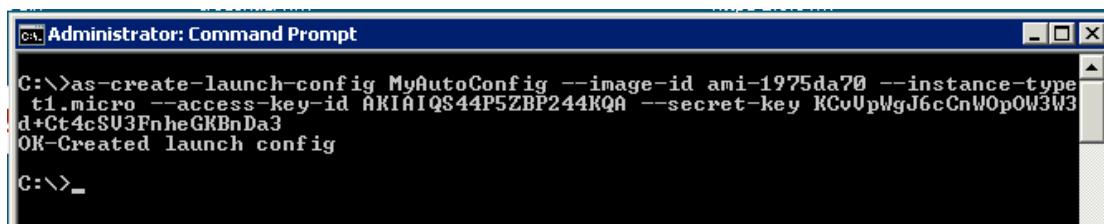
- `-- instance-type: t1.micro`

This is the type of instance that the image cloned from.

Amazon Elastic Compute instances (EC2) are grouped into seven groups [13]: Standard, Micro, High-Memory, High-CPU, Cluster Compute, Cluster GPU, and High I/O. Each type has different characteristics that meet the user's needs. The usage charge is different from one group to another as well.

- Micro instances features:

- 613 MB memory
  - Up to 2 EC2 Compute Units (for short periodic bursts)
  - EBS storage only
  - 32-bit or 64-bit platform ( I chose 64-bit)
  - I/O Performance: Low
  - API name: t1.micro
- --access-key-id & --secret-key: these are the credential security keys.



The screenshot shows a Windows Command Prompt window titled "Administrator: Command Prompt". The command entered is:

```
C:\>as-create-launch-config MyAutoConfig --image-id ami-1975da70 --instance-type t1.micro --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
OK-Created launch config
C:\>_
```

Figure 4.23: MyAutoConfig creation

- An Auto Scaling group was then created and named MyGroup2 (see figure 4.24):

```
as-create-auto-scaling-group MyGroup2 -launch-configuration
MyAutoConfig -availability-zones us-east-1a -min-size 1 -max-size
4 -load-balancers Webserver-Load-Balancer -access-key-id
AKIAIQS44P5ZBP244KQA -secret-key
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
```

- –launch-configuration: MyAutoConfig , which is the name of the launch configuration that was previously created.
- –availability-zones: us-east-1a , defines which zone that the instances will be launched.
  - Amazon has seven different physical locations that can be used to launch instances conveniently and be closer to your customers. For the experiment, the Northern Virginia region was chosen, as it was the closest region to Rochester, NY. The seven regions are [14]:
- US East (Northern Virginia) Region
- US West (Northern California) Region
- US West (Oregon) Region
- EU (Ireland) Region
- Asia Pacific (Singapore) Region
- Asia Pacific (Tokyo) Region
- South America (Sao Paulo) Region
- –min-size 1: the minimum number of instances that the Auto Scaling will not lower.
- –max-size 4: the maximum number of instances that the Auto Scaling will not exceeded.
- --load-balancers :Webserver-Load-Balancer, the name of the load balancer that was created. All new instances will belong to it.
- –access-key-id & --secret-key: the credential security keys.

```
C:\>as-create-auto-scaling-group MyGroup2 --launch-configuration MyAutoConfig --availability-zones us-east-1a --min-size 1 --max-size 4 --load-balancers Webservlet-Load-Balancer --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
OK-Created AutoScalingGroup
C:\>_
```

Figure 4.24: MyGroup2 creation

- Verifying the Auto Scaling Group creation as in figure 4.25:

```
as-describe-auto-scaling-groups -headers -access-key-id
AKIAIQS44P5ZBP244KQA -secret-key
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
```

```
C:\>as-describe-auto-scaling-groups MyGroup2 --headers --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
AUTO-SCALING-GROUP GROUP-NAME LAUNCH-CONFIG AVAILABILITY-ZONES LOAD-BALANCERS
AUTO-SCALING-GROUP MyGroup2 MyAutoConfig us-east-1a Webserver-Load-Balancer 1 4 1
INSTANCE INSTANCE-ID AVAILABILITY-ZONE STATE STATUS LAUNCH-CONFIG
INSTANCE i-1eed7766 us-east-1a InService Healthy MyAutoConfig
C:\>_
```

Figure 4.25: Verifying the Auto Scaling Group creation

- Also, verifying the Auto Scaling instances, which has only one instance (see figure 4.26):

```
as-describe-auto-scaling-instances -headers -access-key-id
AKIAIQS44P5ZBP244KQA -secret-key
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
```

```
C:\>as-describe-auto-scaling-instances --headers --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
INSTANCE INSTANCE-ID GROUP-NAME AVAILABILITY-ZONE STATE STATUS LAUNCH-CONFIG
INSTANCE i-dc861ba4 MyGroup2 us-east-1a InService HEALTHY MyAutoConfig
C:\>_
```

Figure 4.26: Verifying the Auto Scaling instances

- The first scale up policy created was named ScaleUpPolicy (see figure 4.27):

```
as-put-scaling-policy ScaleUpPolicy --auto-scaling-group MyGroup2 --
adjustment=1 --type ChangeInCapacity --cooldown 60 --access-key-id
AKIAIQS44P5ZBP244KQA --secret-key
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
```

- --auto-scaling-group: MyGroup2 , the Auto Scaling group.
- --adjustment=1: adds one instance when the CPU utilization reaches its defined upper level.
- --type ChangeInCapacity: This option was chosen because it could adjust by constant increment. The alternative option was “PercentChangeInCapacity”, which adjusts by the percentage of the current capacity.
- --cooldown 60 seconds: the period that helps prevent the Auto Scaling from initiating more activities before the previous activities are visible.

**Figure 4.27: Scale Up Policy**

This command generates an ARN, which represents the action of the scaling up policy. This policy would use ARN to associate Auto Scaling policies with CloudWatch alarms for scaling up, which will be explained later:

```
arn:aws:autoscaling:us-east-1:131511199887:scalingPolicy:b8d78a4a-
04a5-4f9b-b175-c26ce7aa1569:autoScalingGroupName/
MyGroup2:policyName/ScaleUpPolicy
```

- A second policy, this time to scale down, was created and named ScaleDownPolicy (see figure 4.28):

```
as-put-scaling-policy ScaleDownPolicy --auto-scaling-group MyGroup2 "--adjustment=-1" --type ChangeInCapacity --cooldown 300 --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
```

All these parameters are explained in the previous step.

- --adjustment=-1: deletes one instance when the CPU utilization reaches its defined lower level.

```
C:\>as-put-scaling-policy ScaleDownPolicy --auto-scaling-group MyGroup2 "--adjustment=-1" --type ChangeInCapacity --cooldown 300 --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
arn:aws:autoscaling:us-east-1:131511199887:scalingPolicy:d215e8d8-b5d3-4825-a5a5-a9622c9fb10e:autoScalingGroupName/MyGroup2:policyName/ScaleDownPolicy
```

**Figure 4.28: Scale Down Policy**

In this step, the ARN action associated the policy with CloudWatch alarms for scaling down:

```
arn:aws:autoscaling:us-east-1:131511199887:scalingPolicy:d215e8d8-b5d3-4825-a5a5-a9622c9fb10e:autoScalingGroupName/MyGroup2:policyName/ScaleDownPolicy
```

#### **4.4.5 Creating CloudWatch Alarms:**

The Cloud Watch Alarms were created to notify when any of the Auto Scaling actions were taking place. The cloud watch alarms were set to follow the main server, and when the CPU utilization reaches any of the triggers that were set for the Auto Scaling, it takes action accordingly. In addition, the alarm sends a notification email each time.

To create the CloudWatch alarms, its command line tools must be installed. The environment for the command line was set as follows:

- The AWS\_CLOUDWATCH\_HOME environment variable was set where the Auto Scaling folder was unpacked:

```
C:\> set AWS_AUTO_SCALING_HOME= C:\Users\Administrator\Desktop\  
CloudWatch-2010-08-01\CloudWatch-1.0.12.1
```

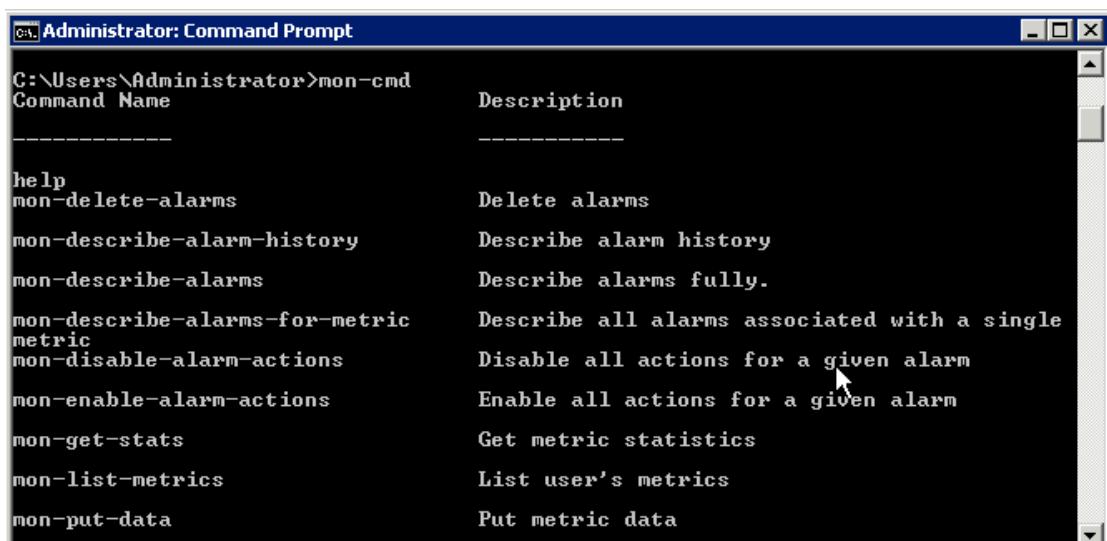
- The %AWS\_AUTO\_SCALING\_HOME%\bin to %PATH environment variable was added:

```
C:\> set PATH=%PATH%;%AWS_COULDWATCH_HOME%\bin
```

- The tools for security credentials were configured and saved in a .template file. The file contains two security credentials: AWS access key ID, and AWS secret key.

```
Set AWS_CREDENTIAL_FILE=C:\>set AWS_CREDENTIAL_FILE=C:\Users\  
Administrator\Desktop\CloudWatch-2010-08-01\CloudWatch-1.0.12.1
```

- The configurations were tested by using a mon-cmd command, and the resulting output (see figure 4.29) confirmed the installation:



Command Name	Description
help	
mon-delete-alarms	Delete alarms
mon-describe-alarm-history	Describe alarm history
mon-describe-alarms	Describe alarms fully.
mon-describe-alarms-for-metric metric	Describe all alarms associated with a single metric
mon-disable-alarm-actions	Disable all actions for a given alarm
mon-enable-alarm-actions	Enable all actions for a given alarm
mon-get-stats	Get metric statistics
mon-list-metrics	List user's metrics
mon-put-data	Put metric data

Figure 4.29: Testing the CloudWatch command line installation

- The first alarm created was called “HighCPU-alarm” (see figure 4.30). Once the CPU reaches its defined high level ( $\geq 80$ ), the alarm runs the ARN policy action and sends a notification when the Auto Scaling creates a new web server instance:

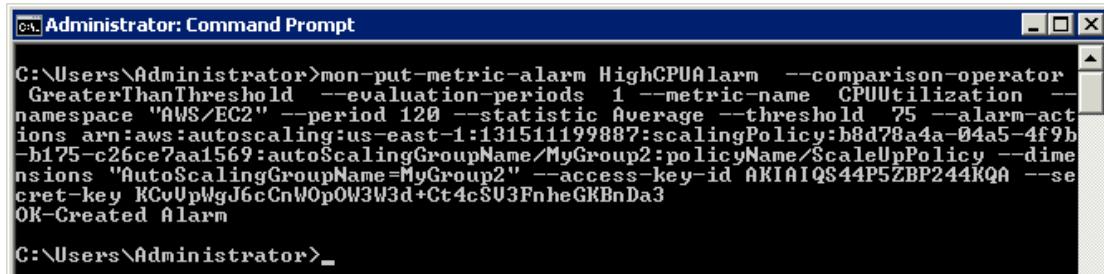
```
mon-put-metric-alarm HighCPUAlarm --comparison-operator  
GreaterThanOrEqualTo --evaluation-periods 1 --metric-name
```

```

CPUUtilization --namespace "AWS/EC2" --period 300 --statistic
Average --threshold 80 --alarm-actions arn:aws:autoscaling:us-east-
1:131511199887:scalingPolicy:b8d78a4a-04a5-4f9b-b175-
c26ce7aa1569:autoScalingGroupName/MyGroup2:policyName/ScaleUpPolicy -
--dimensions "AutoScalingGroupName=MyGroup2" --access-key-id
AKIAIQS44P5ZBP244KQA --secret-key
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

```

- --comparison-operator GreaterThanThreshold: the operator used to scale up, in this case, the GreaterThanThreshold.
- --evaluation-periods 1: Number of consecutive periods for which the value of the metric needs to be compared to the threshold [15].
- --metric-name CPUUtilization: The metric used.
- --namespace "AWS/EC2": the name space for the alarms. Here, it is the default name.
- --period 300: after 300 seconds, the metric would set off the alarm.
- --statistic Average: The average between periods. The minimum or the maximum can also be used.
- --threshold 80: The alarm would take place when the CPU utilization reaches 80%.
- --alarm-actions arn:aws:autoscaling:us-east-1:131511199887:scalingPolicy:b8d78a4a-04a5-4f9b-b175-c26ce7aa1569:autoScalingGroupName/MyGroup2:policyName/ScaleUpPolicy: The connection between the alarm action and the Auto Scaling.
- --dimensions "AutoScalingGroupName=MyGroup2": all available instances were filtered by the name of the group for Auto Scaling: MyGroup2.



```
C:\Users\Administrator>mon-put-metric-alarm HighCPUAlarm --comparison-operator GreaterThanThreshold --evaluation-periods 1 --metric-name CPUUtilization --namespace "AWS/EC2" --period 120 --statistic Average --threshold 75 --alarm-actions arn:aws:autoscaling:us-east-1:131511199887:scalingPolicy:b8d78a4a-04a5-4f9b-b175-c26ce7aa1569:autoScalingGroupName/MyGroup2:policyName/ScaleUpPolicy --dimensions "AutoScalingGroupName=MyGroup2" --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
OK-Created Alarm
C:\Users\Administrator>_
```

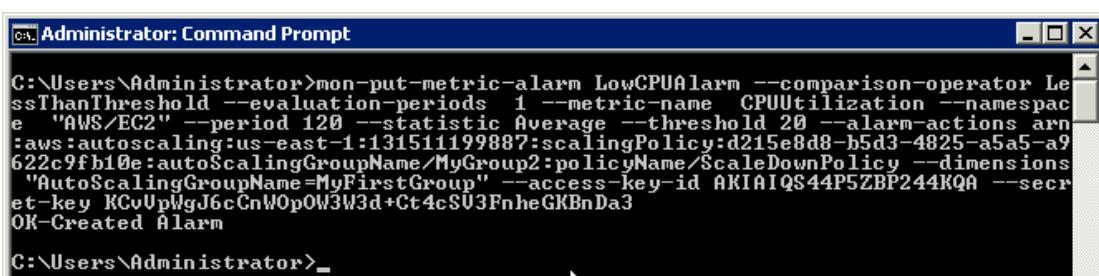
Figure 4.30: HighCPU-alarm

- The second alarm created was called “LowCPU-Alarm” (see figure 4.31).

Once the CPU reaches its defined low level ( $\leq 20$ ), the alarm runs the ARN policy action and sends a notification when the Auto Scaling deletes an instance:

```
mon-put-metric-alarm LowCPUAlarm --comparison-operator LessThanThreshold --evaluation-periods 1 --metric-name CPUUtilization --namespace "AWS/EC2" --period 300 --statistic Average --threshold 20 --alarm-actions arn:aws:autoscaling:us-east-1:131511199887:scalingPolicy:d215e8d8-b5d3-4825-a5a5-a9622c9fb10e:autoScalingGroupName/MyGroup2:policyName/ScaleDownPolicy --dimensions "AutoScalingGroupName=MyFirstGroup" --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
```

All the parameters were explained in the previous step. In the --comparison-operator the LessThanThreshold was used as an operator to scale down.



```
C:\Users\Administrator>mon-put-metric-alarm LowCPUAlarm --comparison-operator LessThanThreshold --evaluation-periods 1 --metric-name CPUUtilization --namespace "AWS/EC2" --period 120 --statistic Average --threshold 20 --alarm-actions arn:aws:autoscaling:us-east-1:131511199887:scalingPolicy:d215e8d8-b5d3-4825-a5a5-a9622c9fb10e:autoScalingGroupName/MyGroup2:policyName/ScaleDownPolicy --dimensions "AutoScalingGroupName=MyFirstGroup" --access-key-id AKIAIQS44P5ZBP244KQA --secret-key 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
OK-Created Alarm
C:\Users\Administrator>_
```

Figure 4.31: LowCPU-alarm

- The alarms of the CloudWatch in AWS management console (see figure 4.32):

Name	Threshold
LowCPU-Alarm	CPUUtilization <= 20 for 2 minutes
HighCPU-Alarm	CPUUtilization >= 75 for 2 minutes

Figure 4.32: The alarms of the CloudWatch in AWS management console

## 5. DEMONSTRATIONS & RESULTS:

### 5.1 Auto Scaling Testing:

After testing the Auto Scaling policies for scaling up and scaling down, they worked perfectly. The Wget tool was used to download the entire website content using three different shells. As a result, the Auto Scaling created a new web server:

Name	Instance	AMI ID	Root Device	Type	State	Status Checks
empty	i-3208954a	ami-1975da70	ebs	t1.micro	<span style="color: green;">●</span> running	<span style="color: gray;">✖</span> initializing...

Figure 5.1: Auto Scaling creates an instance

As shown under “Status Checks” in Figure 5.1, the new instance was initializing. It took about two minutes until it was ready to use.

The testing of the creation of the new instance showed that scaling up features worked. To test the scaling down features, the newly created instance was deleted when the CPU utilization reached its lower level. The download was halted from the webserver (1) and, as a result, the CPU utilization became less than 20% for five minutes, prompting the CloudWatch LowCPU-Alarm to take action (see figure 5.2):

	<b>State</b>	<b>Name</b>	<b>Threshold</b>
<input type="checkbox"/>	ALARM	LowCPU-Alarm	CPUUtilization <= 20 for 2 minutes
<input type="checkbox"/>	OK	HighCPU-Alarm	CPUUtilization >= 75 for 2 minutes

Figure5.2: CloudWatch Alarms<sup>8</sup>

- The alarm sent a notification email:

ALARM: "Proceed the Minimum" in US - N. Virginia Inbox x

AWS Notifications no-reply@sns.amazonaws.com via bounces.amazon.com to me 8:40 PM (0 minutes ago)

You are receiving this email because your Amazon CloudWatch Alarm "Proceed the Minimum" in the US - N. Virginia region has entered the ALARM state, because "Threshold Crossed: 2 datapoints were less than or equal to the threshold (20.0). The most recent datapoints: [2.08, 2.296000000000003]." at "Thursday 12 July, 2012 00:40:16 UTC".

View this alarm in the AWS Management Console:  
<https://console.aws.amazon.com/cloudwatch/home?region=us-east-1#s=Alarms&alarm=Proceed%20the%20Minimum>

Alarm Details:

- Name: Proceed the Minimum
- Description: Removing an instances when the CPU utilization becomes less than 20 % for 10 minutes
- State Change: INSUFFICIENT\_DATA -> ALARM
- Reason for State Change: Threshold Crossed: 2 datapoints were less than or equal to the threshold (20.0). The most recent datapoints: [2.08, 2.296000000000003].
- Timestamp: Thursday 12 July, 2012 00:40:16 UTC
- AWS Account: 131511199887

Threshold:

- The alarm is in the ALARM state when the metric is LessThanOrEqualToThreshold 20.0 for 300 seconds.

Monitored Metric:

- MetricNamespace: AWS/EC2
- MetricName: CPUUtilization
- Dimensions: [AutoScalingGroupName = MyGroup]
- Period: 300 seconds
- Statistic: Average
- Unit: not specified

State Change Actions:

- OK:
- ALARM: [arn:aws:sns:us-east-1:131511199887:NotifyMe] [arn:aws:autoscaling:us-east-1:131511199887:scalingPolicy:be968e44-bdd1-49af-865d-606b6b59ed1c:autoScalingGroupName/MyAutoScalingGroup;policyName/MyScaleDownPolicy]
- INSUFFICIENT\_DATA:

--

If you wish to stop receiving notifications from this topic, please click or visit the link below to unsubscribe:  
<https://sns.us-east-1.amazonaws.com/unsubscribe.html?SubscriptionArn=arn:aws:sns:us-east-1:131511199887:NotifyMe:4d8184e5-a3b4-4570-bc8e-5be2cccb6b8&Endpoint=aljebreen50@gmail.com>

Please do not reply directly to this e-mail. If you have any questions or comments regarding this email, please contact us at [sns-question@amazon.com](mailto:sns-question@amazon.com)

Click here to [Reply](#) or [Forward](#)

Figure 5.3: A notification email in aljebreen50@gmail.com inbox

<sup>8</sup> As it appears, the LowCPU-alarm is reached its goal, which means that the CPU utilization in that moment became less than 20% for two minutes, which allowed the Auto Scaling to remove the temporary web server. The HighCPU-Alarm is still in its range, which means the CPU utilization was still fewer than 80%, so I can see the green sign next to the alarm.

- After the CPU utilization became less than 20% for five minutes, the ScalingDown policy took place and shut the temporary instance down (figure 5.4):

	Name	Instance	AMI ID	Root Device	Type	State	Status Checks	Alarm Status	Monitoring	Security Groups
!	empty	i-2630555e	ami-1975da70	ebs	t1.micro	shutting-down		none	detailed	default

**Figure 5.4: shutting the temporary instance down**

- The instance was then terminated (see figure 5.5):

	Name	Instance	AMI ID	Root Device	Type	State	Status Checks	Alarm Status	Monitoring	Security Groups
!	empty	i-2630555e	ami-1975da70	ebs	t1.micro	terminated		none	detailed	default

**Figure 5.5: Terminating the temporary instance**

- Afterwards, the experiment was successfully scaled up and down of the proposed algorithm (see figure 5.4).

## **5.2 Tests' Results:**

In these next tests, the auto scaling actions did not take place. The reason was to monitor the web server's performance during the tests without scaling up (adding a new instance) or scaling down (deleting an instance), which would help decide which metric is the best overall. To clarify this point, let's take this example: perhaps the auto scaling action would take place during the test time, and the CPU Utilization exceeded the high boundary (which was 80%) only for five minutes and returned back to 60% for the rest of the test period. The auto scaling would create a new instance because it exceeded 80% for 5 minutes, and it would not delete it because it did not exceed the low boundary (20%), thus preventing the web server from exceeding the upper boundary alone without additional instances. However, placing an alarm action, which simply alerts CloudWatch when the web server has reached a boundary,

instead of auto scaling action is more beneficial to the experiment as it is non-intrusive.

After testing, the results were analyzed using tables and graphs, as shown below:

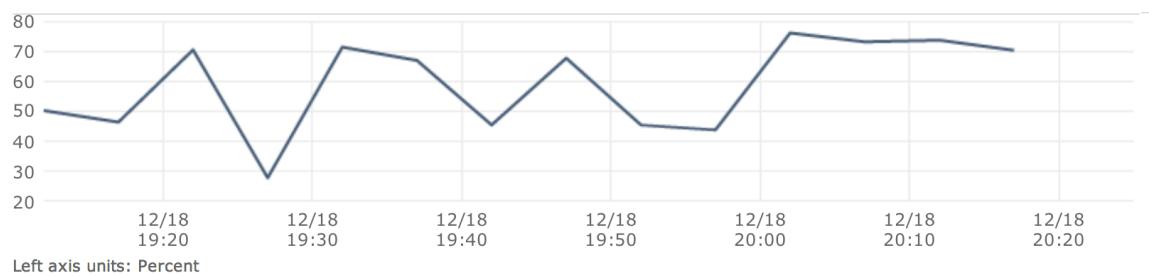
### **5.2.1 CPU Utilization Metric Results:**

#### **5.2.1.1 First Test:**

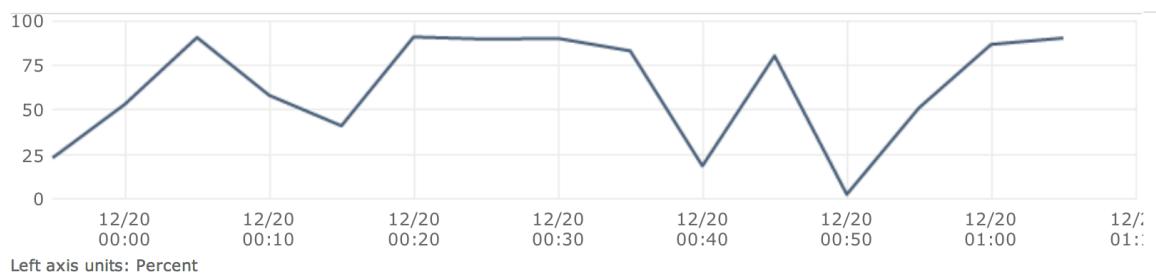
Test Place	High Alarm Period (~ minutes) ( $\geq 80\%$ )	Low Alarm Period (~minutes) ( $\leq 20\%$ )
<i>Morning Weekdays RIT</i>	0	0
<i>Evening Weekdays RIT</i>	12	4
<i>Morning Weekends RIT</i>	25	0
<i>Evening Weekends RIT</i>	17	2
<i>Morning Weekdays Home</i>	10	8
<i>Evening Weekdays Home</i>	0	19
<i>Morning Weekends Home</i>	0	22
<i>Evening Weekends Home</i>	0	27
<i>Virginia</i>	30	0
<i>Oregon</i>	8	12
<i>North California</i>	2	6

**Table 2: CPU Utilization Metric Results of the First Test**

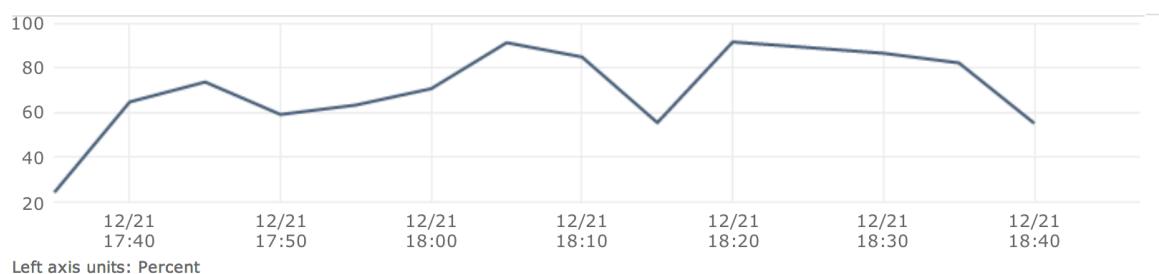
- *Weekday mornings at RIT System Lab*



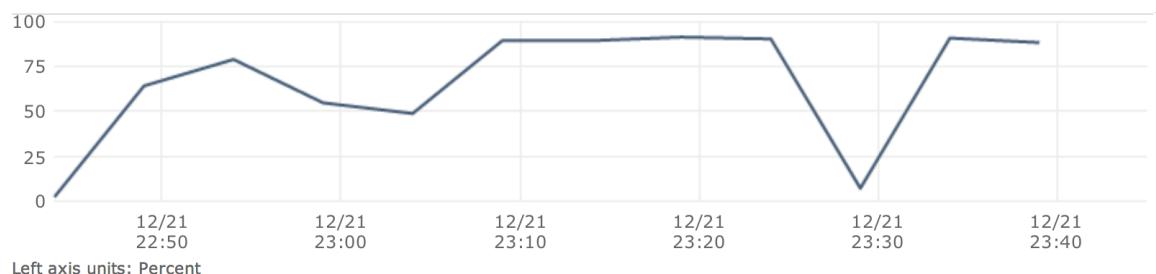
○ *Weekday evenings at RIT System Lab*



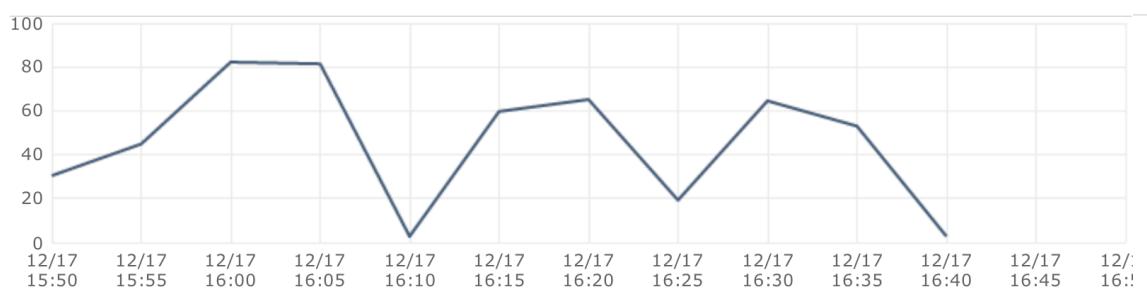
○ *Weekend mornings at RIT System Lab*



○ *Weekend evenings at RIT System Lab*



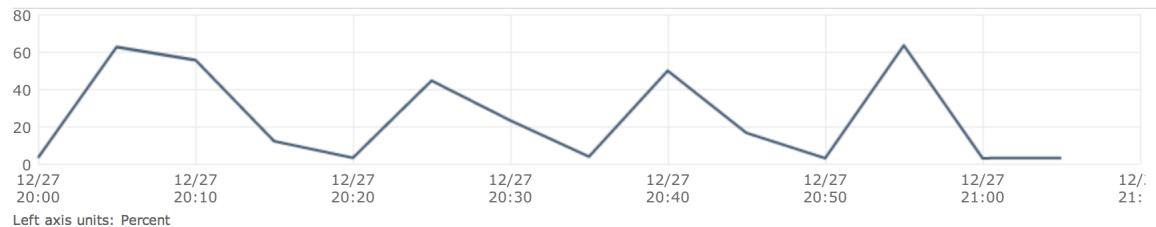
○ *Weekday mornings from home*



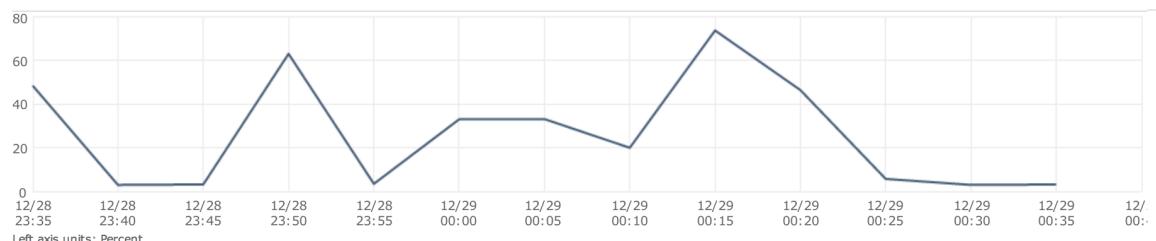
○ *Weekday evenings from home*



○ *Weekend mornings from home*

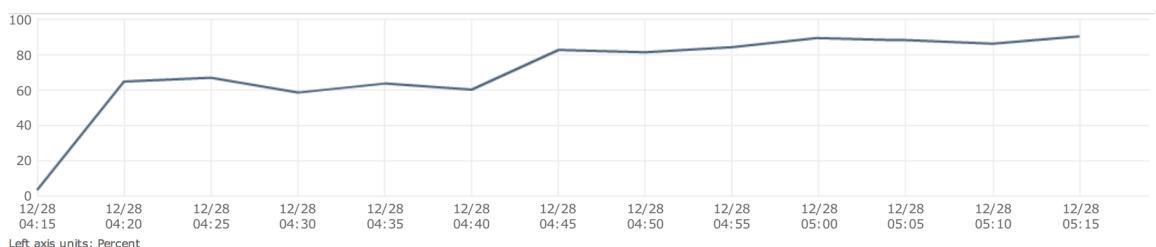


○ *Weekend evenings from home*

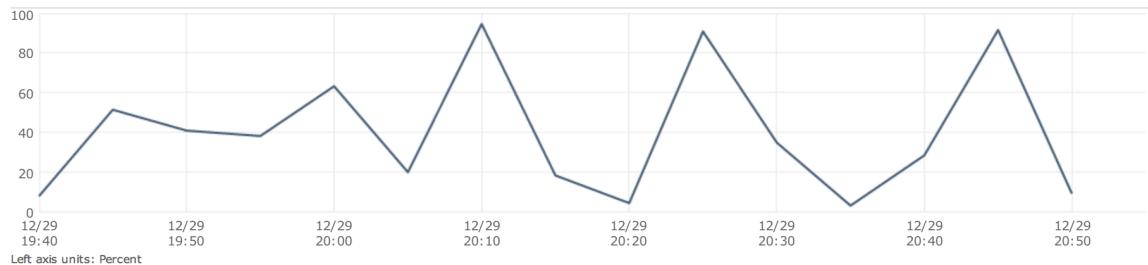


○ *From a server within Amazon cloud in the same geographic region,*

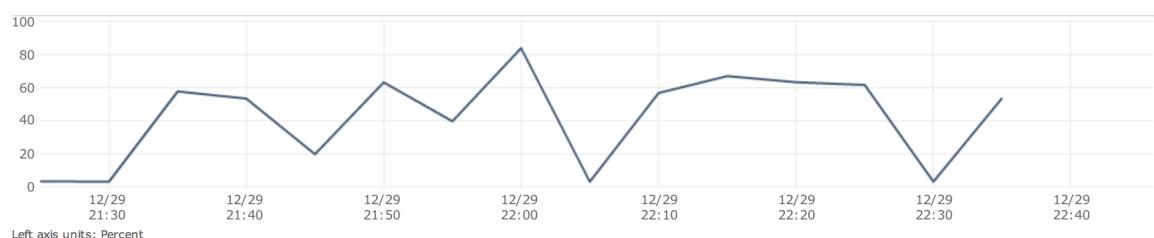
*Virginia web server*



- From a server within Amazon cloud in Oregon geographic region



- From a server within Amazon cloud in Northern California geographic region



### 5.2.1.2 Second Test:

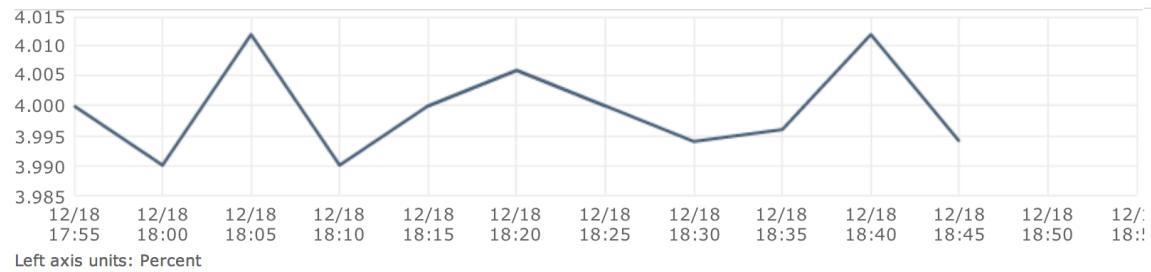
CPU Utilization in the second test was always around 4% for most of the tests, even when the server was receiving around 45 connections per minute.

Test Place	High Alarm Period (~ minutes) ( $\geq 80\%$ )	Low Alarm Period (~minutes) ( $\leq 20\%$ )
Morning Weekdays RIT	0	During the whole test period
Evening Weekdays RIT	0	During the whole test period
Morning Weekends RIT	0	During the whole test period
Evening Weekends RIT	0	During the whole test period
Morning Weekdays Home	0	During the whole test period
Evening Weekdays Home	0	During the whole test period
Morning Weekends Home	0	During the whole test period
Evening	0	Only 15 minutes above 20%

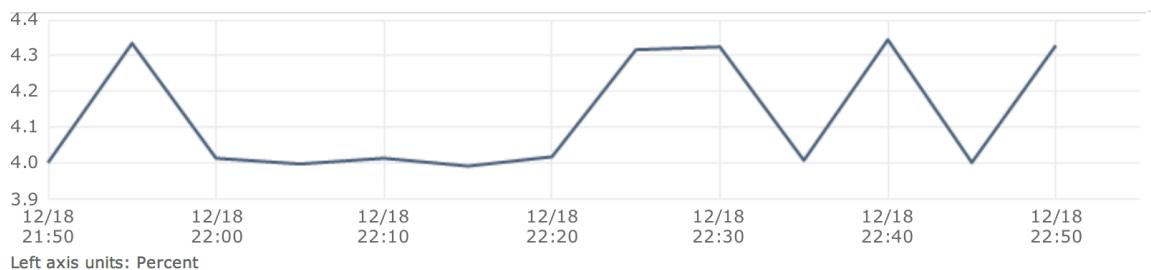
<i>Weekends Home</i>		
<i>Virginia</i>	9	Only 2 minutes above 20%
<i>Oregon</i>	0	During the whole test period
<i>North California</i>	0	During the whole test period

**Table 3: CPU Utilization Metric Results of Second Test**

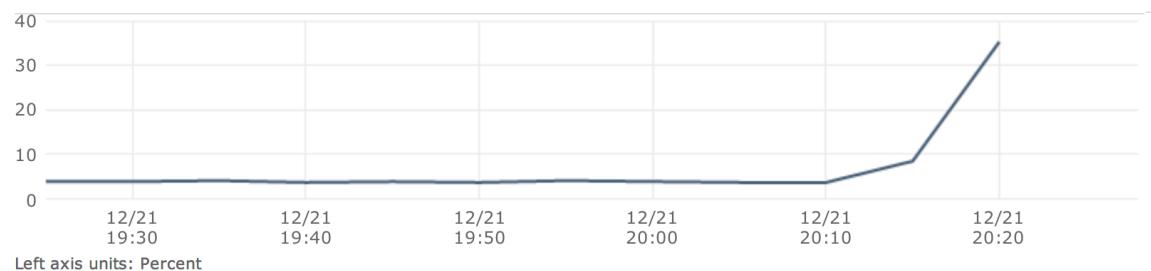
- *Weekday mornings at RIT System Lab*



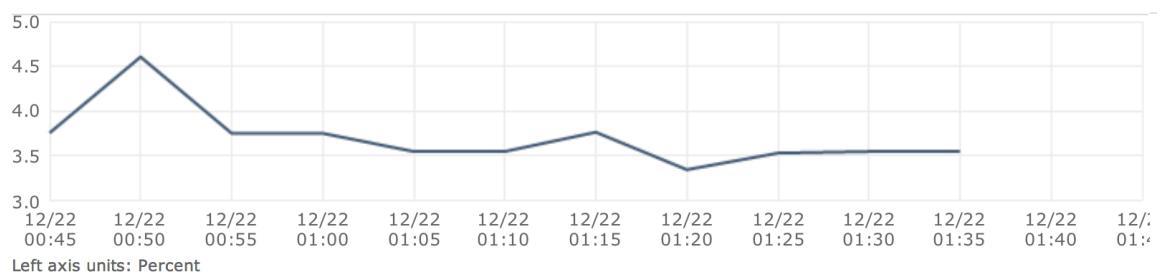
- *Weekday evenings at RIT System Lab*



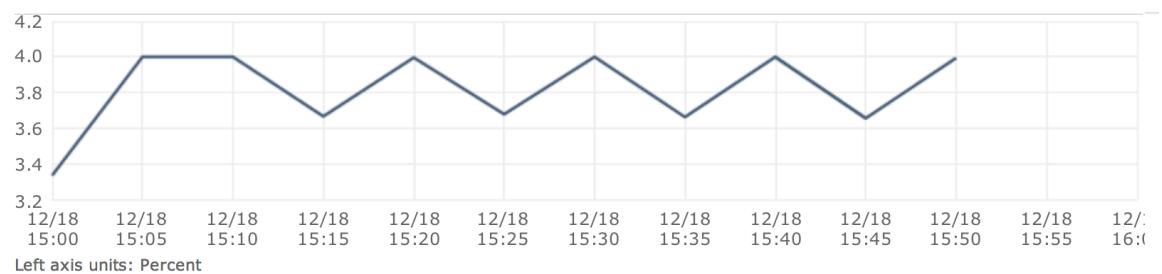
- *Weekend mornings at RIT System Lab*



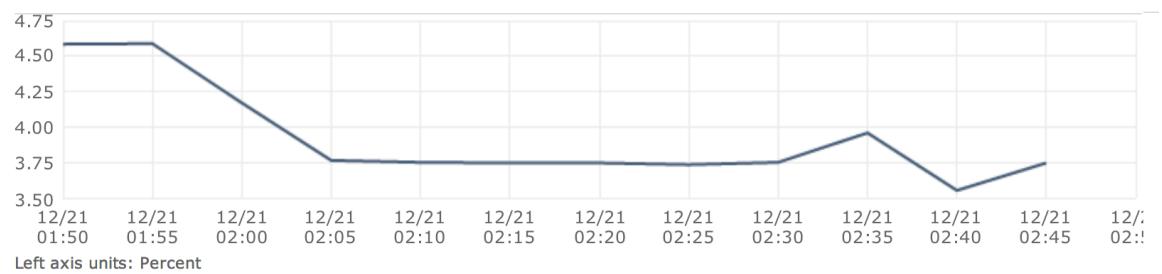
○ *Weekend evenings at RIT System Lab*



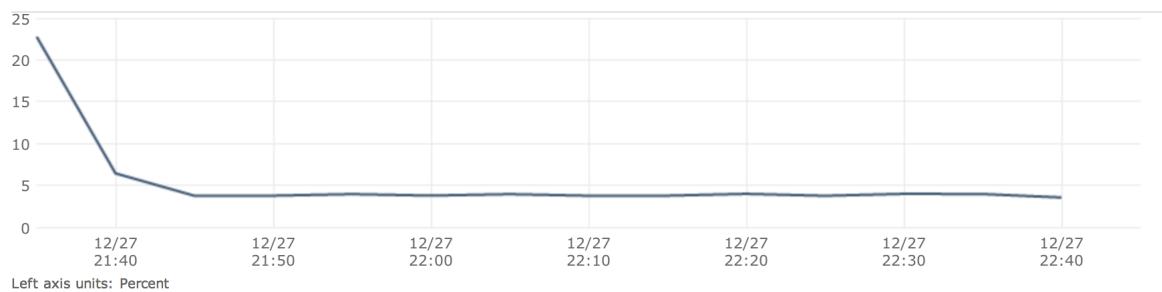
○ *Weekday mornings from home*



○ *Weekday evenings from home*

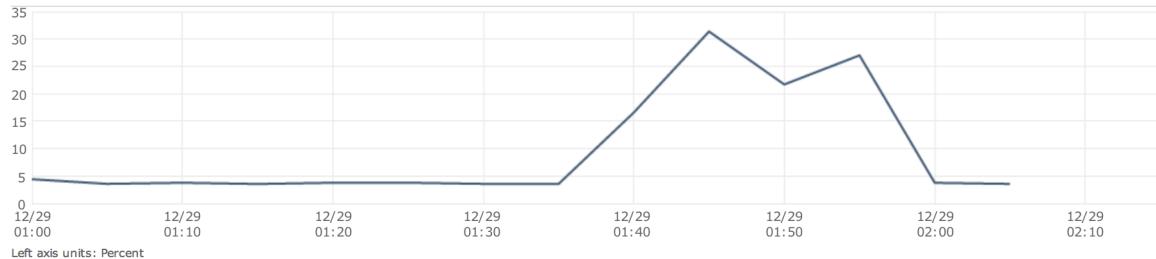


○ *Weekend mornings from home*

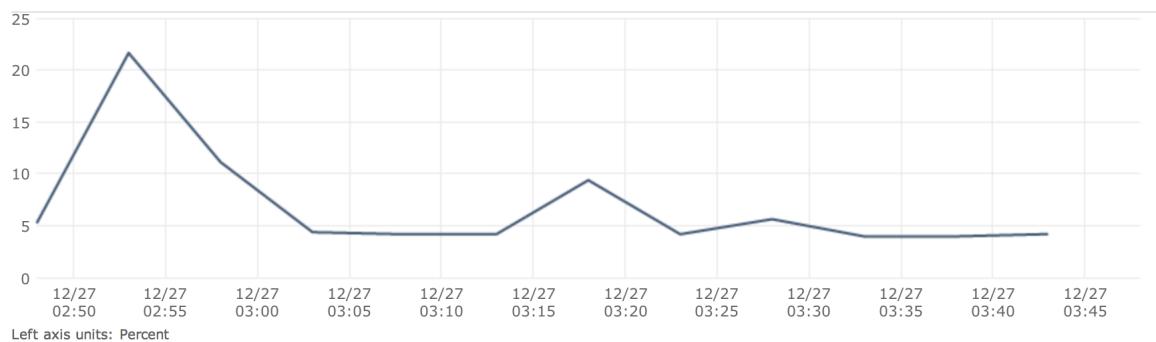


- *Weekend evenings from home*

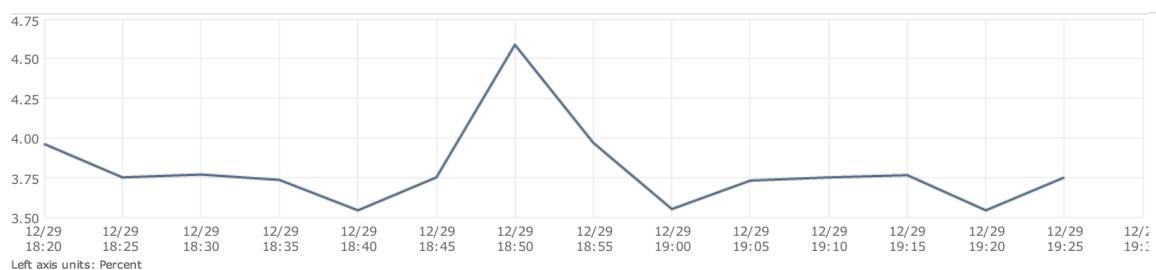
In this test, the performance jumped to 30% for 25 minutes. Although this is an outlier result, it was still taken into consideration.



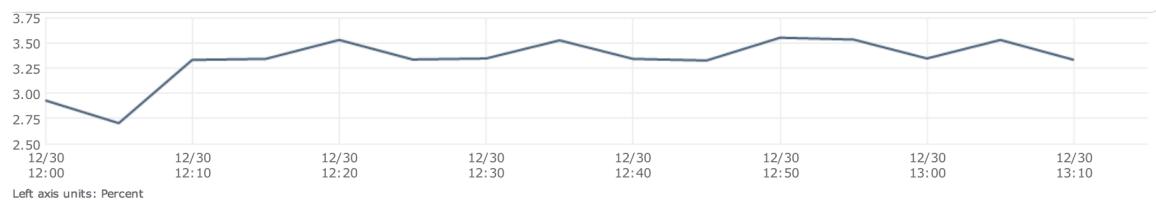
- *From a server within Amazon cloud in the same geographic region,  
Virginia web server*



- *From a server within Amazon cloud in Oregon geographic region*



- From a server within Amazon cloud in Northern California geographic region



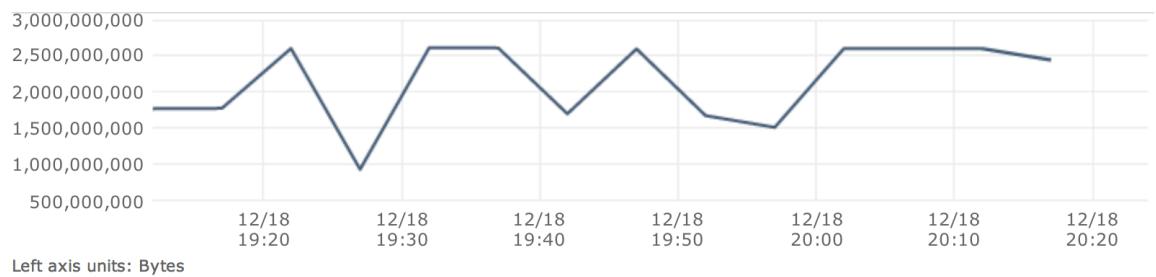
### 5.2.2 Network Out Metric Results:

#### 5.2.2.1 First Test:

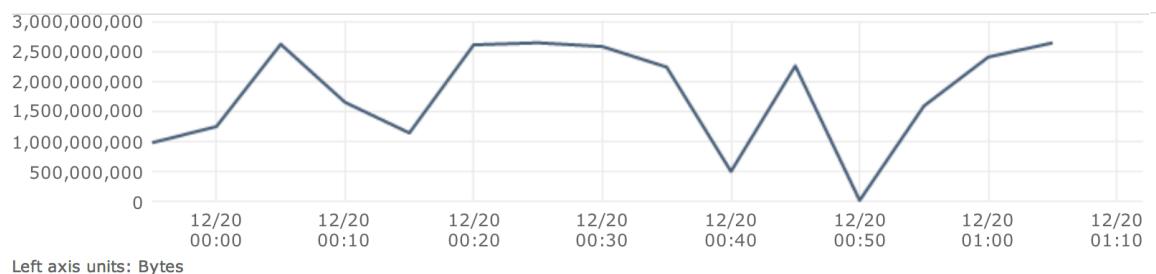
Test Place	High Alarm Period (~ Minutes) ( $\geq 2$ GB)	Low Alarm Period (~ Minutes) ( $\leq 5$ MB)
Morning Weekdays RIT	34	0
Evening Weekdays RIT	25	2
Morning Weekends RIT	35	0
Evening Weekends RIT	32	2
Morning Weekdays Home	9	4
Evening Weekdays Home	0	19
Morning Weekends Home	0	22
Evening Weekends Home	0	31
Virginia	35	0
Oregon	9	13
North California	3	6

Table 4: Network Out Metric Results of First Test

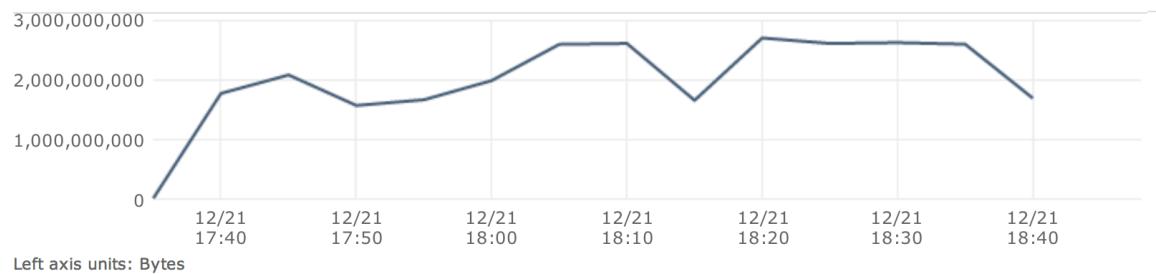
○ Weekday mornings at RIT System Lab



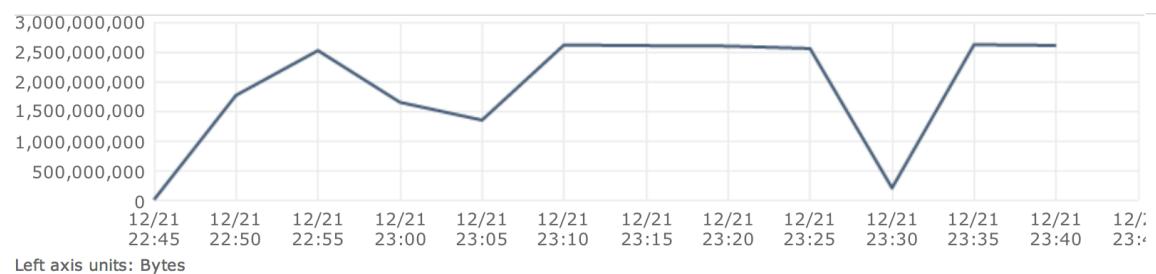
○ Weekday evenings at RIT System Lab



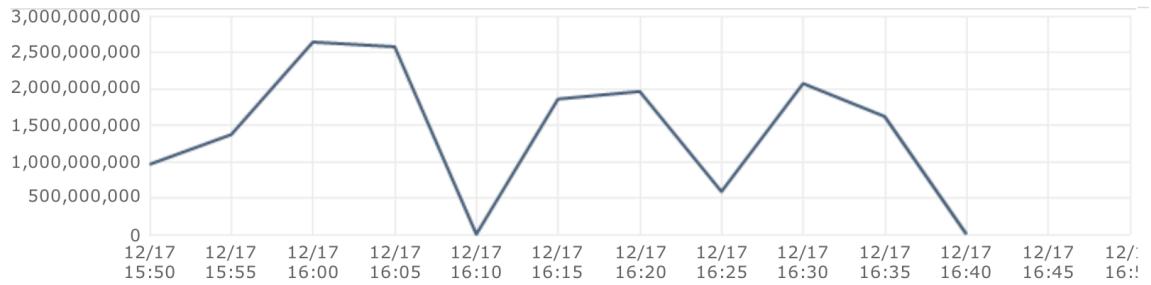
○ Weekend mornings at RIT System Lab



○ Weekend evenings at RIT System Lab



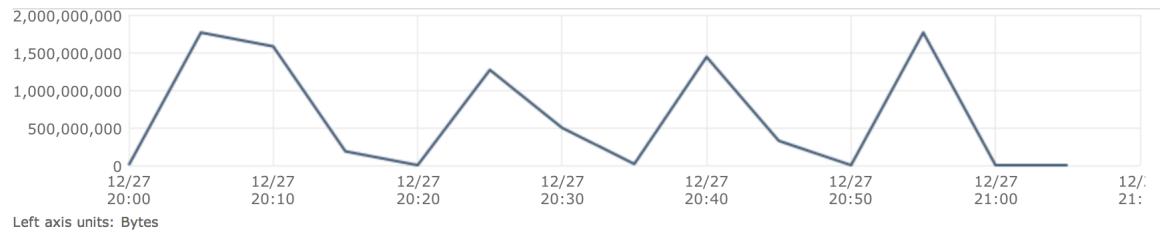
○ *Weekday mornings from home*



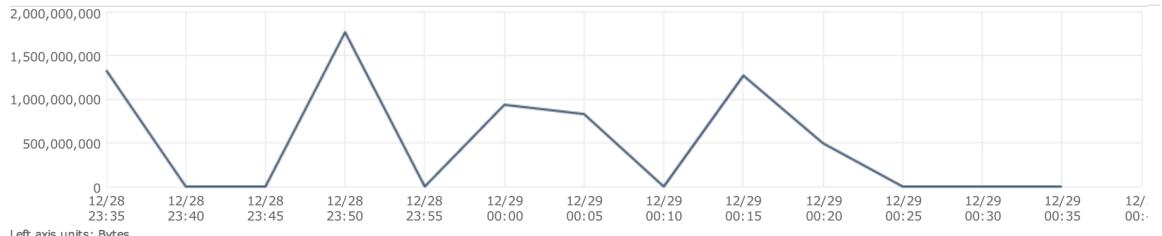
○ *Weekday evenings from home*



○ *Weekend mornings from home*

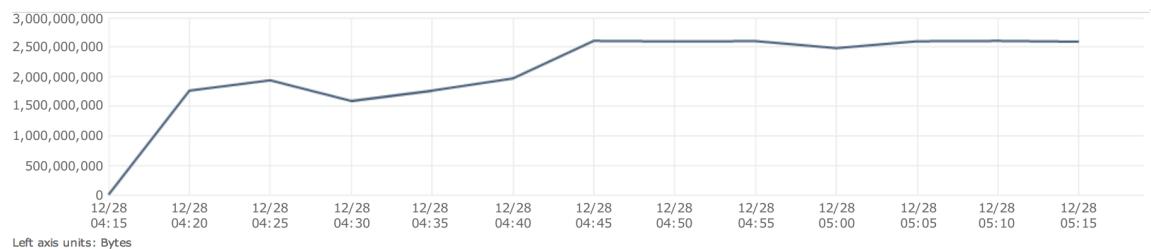


○ *Weekend evenings from home*

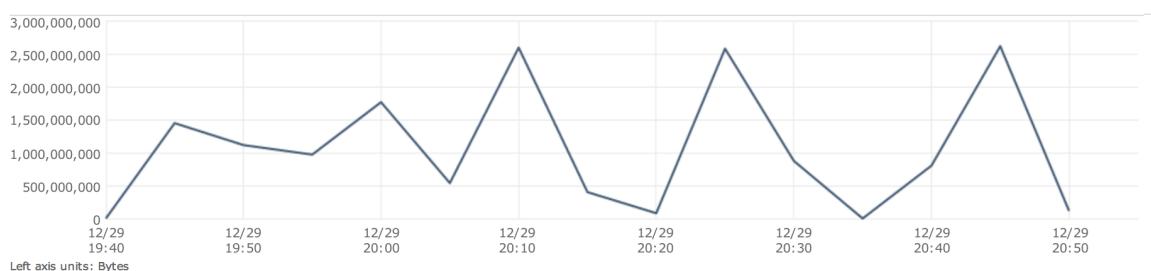


- From a server within Amazon cloud in the same geographic region,

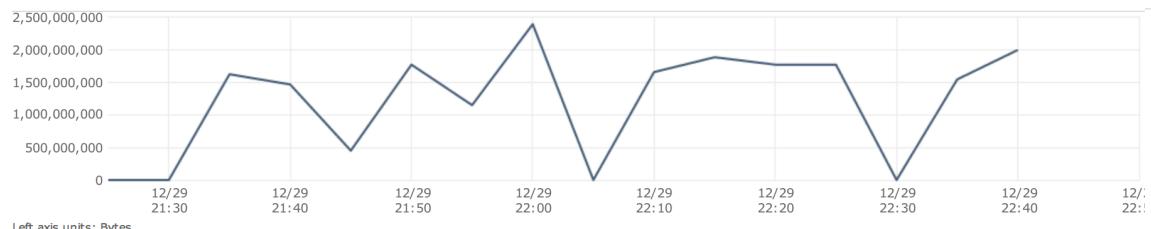
*Virginia web server*



- From a server within Amazon cloud in Oregon geographic region



- From a server within Amazon cloud in Northern California geographic region



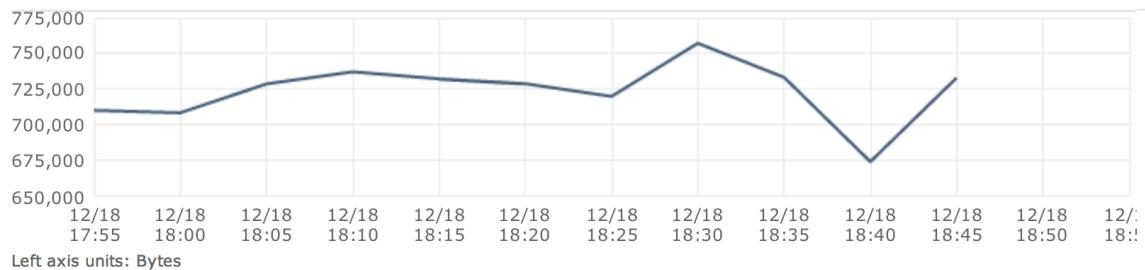
### 5.2.2.2 Second Test:

Test Place	High Alarm Period (~ Minutes) ( $\geq 2$ GB)	Low Alarm Period (~ Minutes) ( $\leq 5$ MB)
Morning Weekdays RIT	0	During the whole test period
Evening Weekdays RIT	0	During the whole test period
Morning Weekends RIT	0	During the whole test period
Evening Weekends RIT	0	During the whole test period

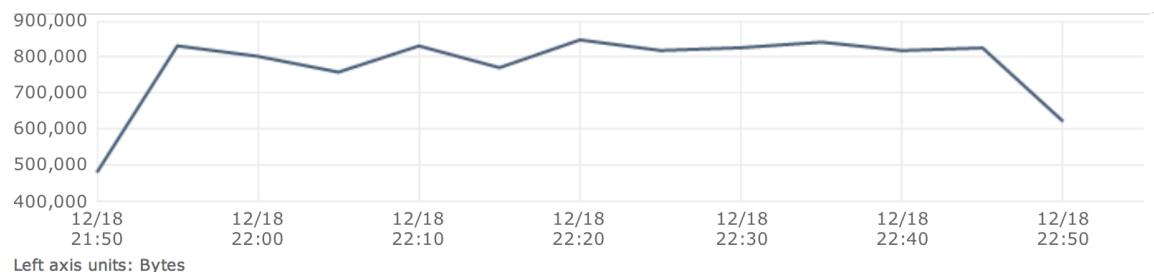
<i>Morning Weekdays Home</i>	0	During the whole test period
<i>Evening Weekdays Home</i>	0	During the whole test period
<i>Morning Weekends Home</i>	0	During the whole test period
<i>Evening Weekends Home</i>	0	During the whole test period
<i>Virginia</i>	0	During the whole test period
<i>Oregon</i>	0	During the whole test period
<i>North California</i>	0	During the whole test period

**Table 5: Network Out Metric Results of Second Test**

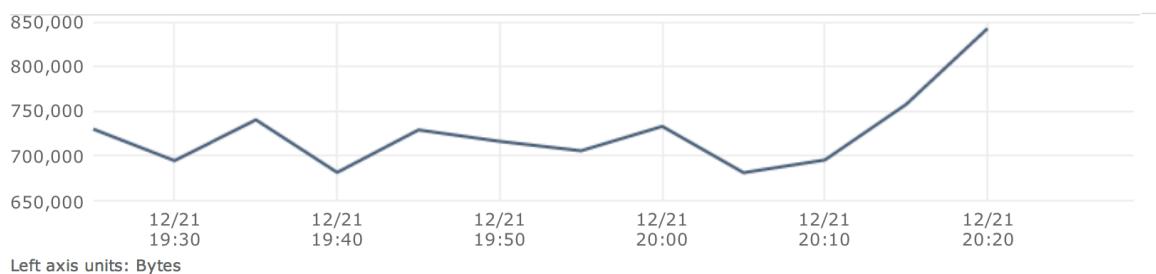
- Weekday mornings at RIT System Lab



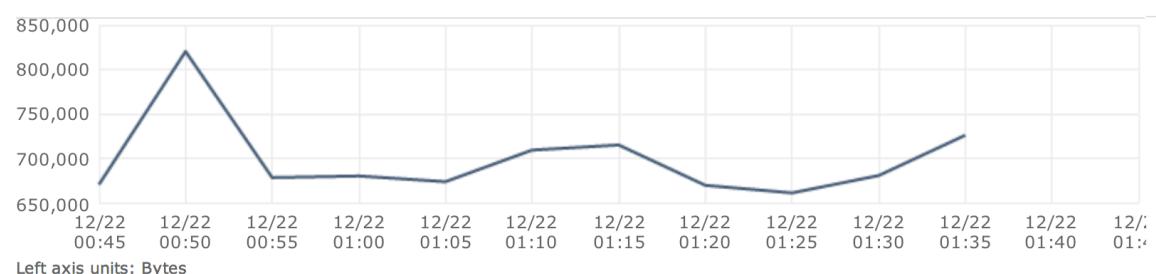
- Weekday evenings at RIT System Lab



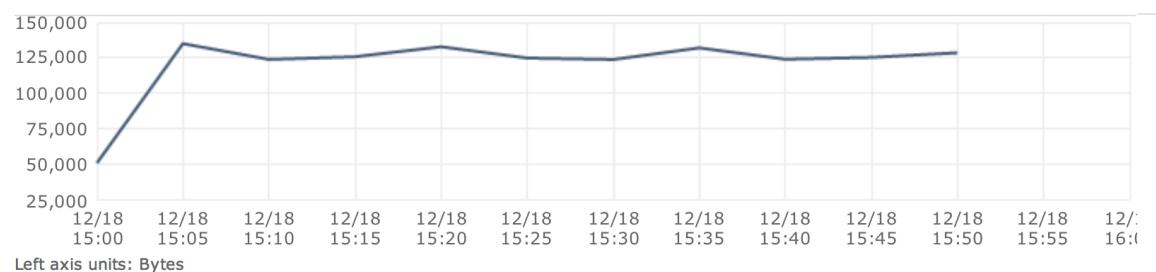
○ *Weekend mornings at RIT System Lab*



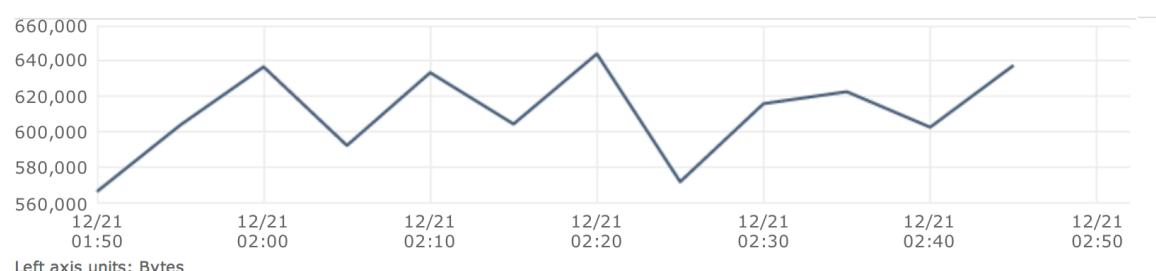
○ *Weekend evenings at RIT System Lab*



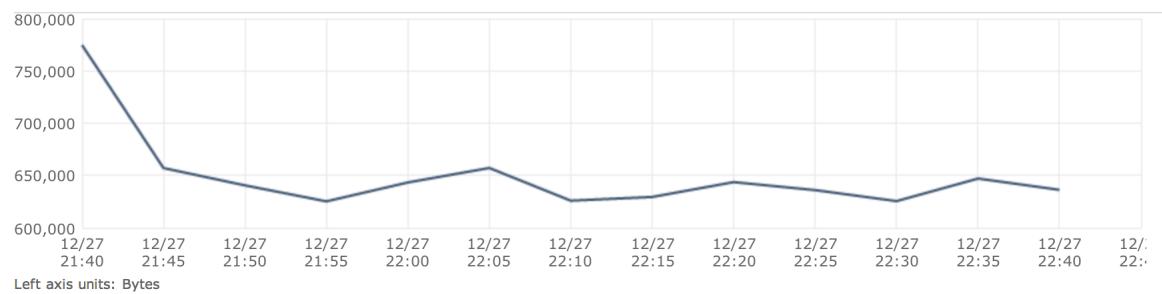
○ *Weekday mornings from home*



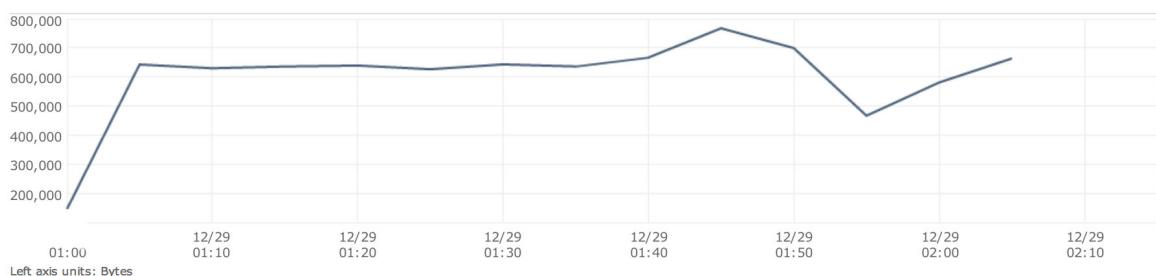
○ *Weekday evenings from home*



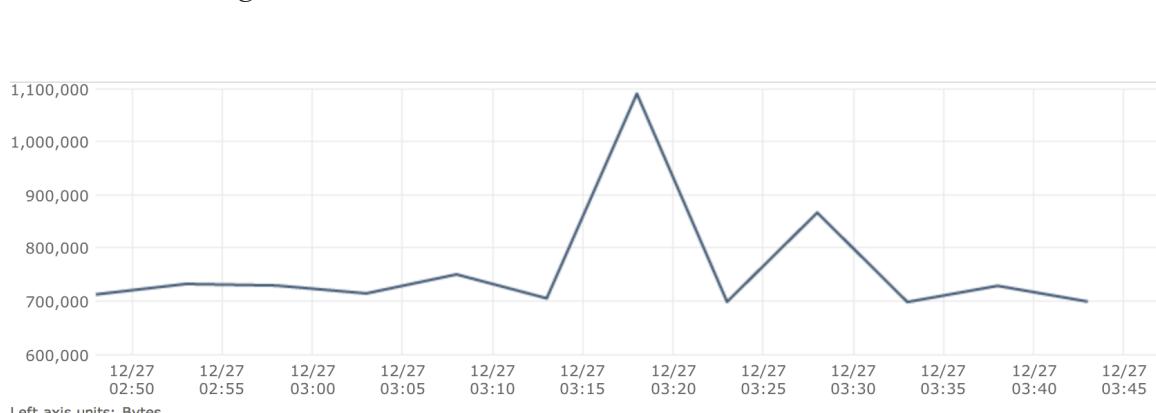
○ *Weekend mornings from home*



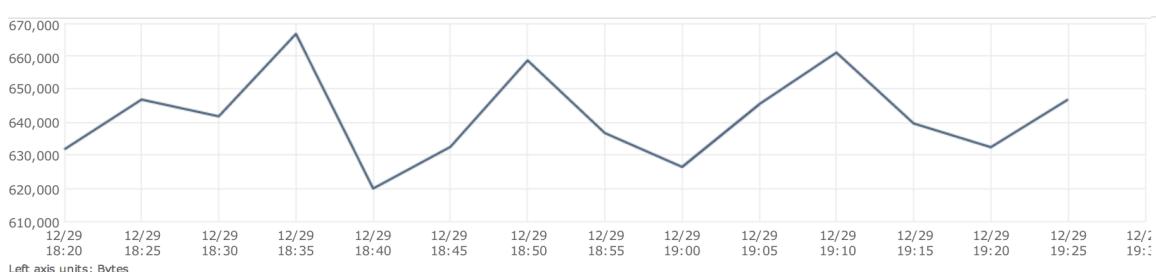
○ *Weekend evenings from home*



○ *From a server within Amazon cloud in the same geographic region, Virginia web server*



○ *From a server within Amazon cloud in Oregon geographic region*



- From a server within Amazon cloud in Northern California geographic region



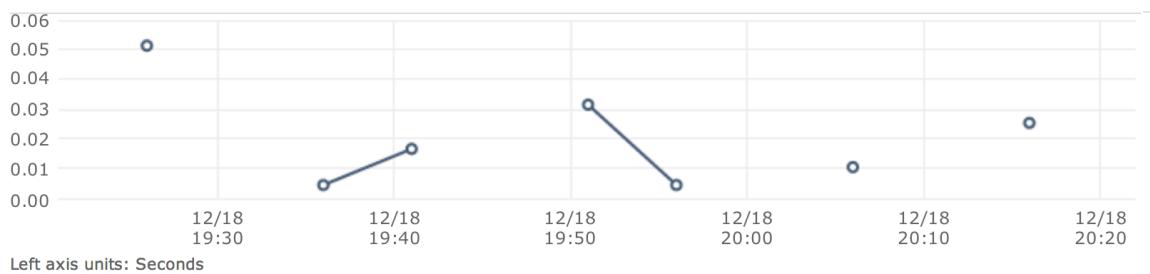
### 5.2.3 Latency Metric Results:

#### 5.2.3.1 First Test:

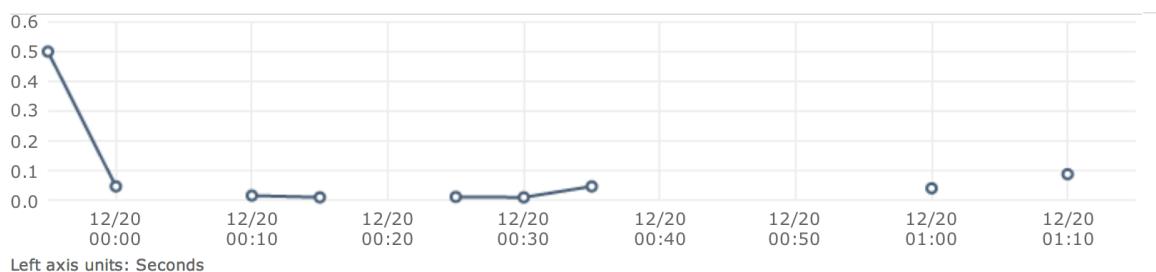
Test Place	High Alarm Period (~ Minutes) (.30 Second)	Low Alarm Period (~ Minutes) (<= .005 Second)
Morning Weekdays RIT	9	0
Evening Weekdays RIT	2	27
Morning Weekends RIT	0	25
Evening Weekends RIT	During the whole test period except 4 minutes	0
Morning Weekdays Home	5	0
Evening Weekdays Home	0	0
Morning Weekends Home	8	8
Evening Weekends Home	1	0
Virginia	13	0
Oregon	0	0
North California	0	0

Table 6: Latency Metric Results of First Test

○ Weekday mornings at RIT System Lab



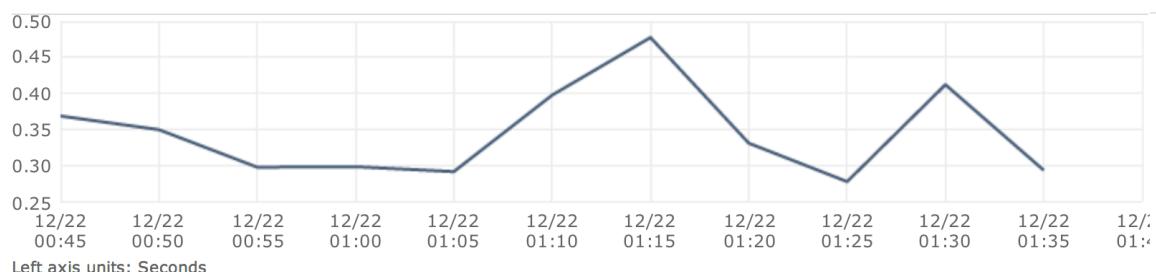
○ Weekday evenings at RIT System Lab



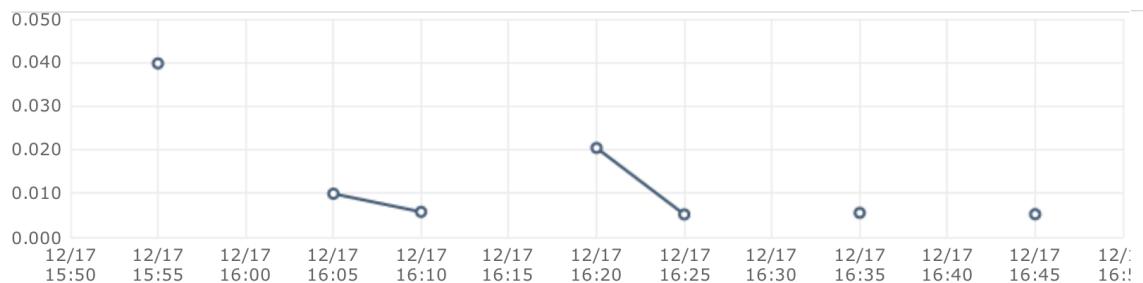
○ Weekend mornings at RIT System Lab



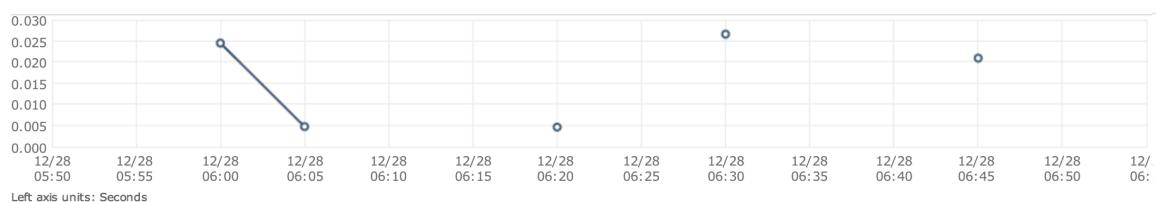
○ Weekend evenings at RIT System Lab



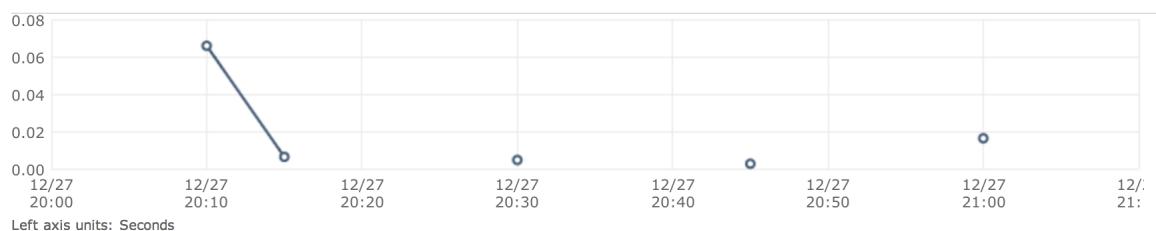
○ Weekday mornings from home



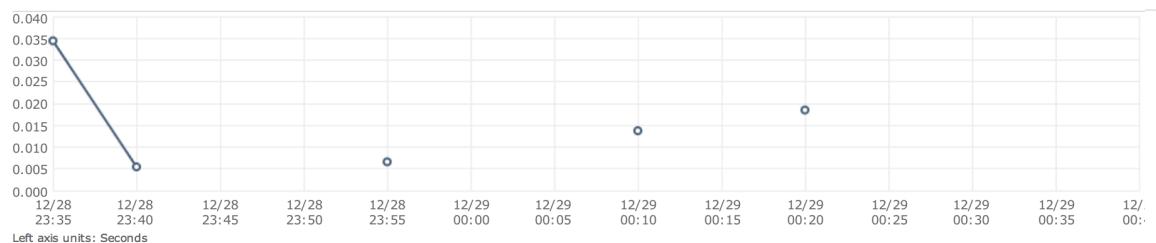
○ Weekday evenings from home



○ Weekend mornings from home

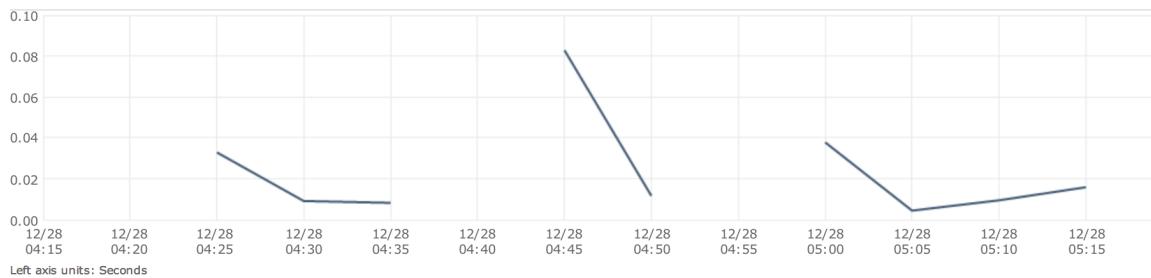


○ Weekend evenings from home

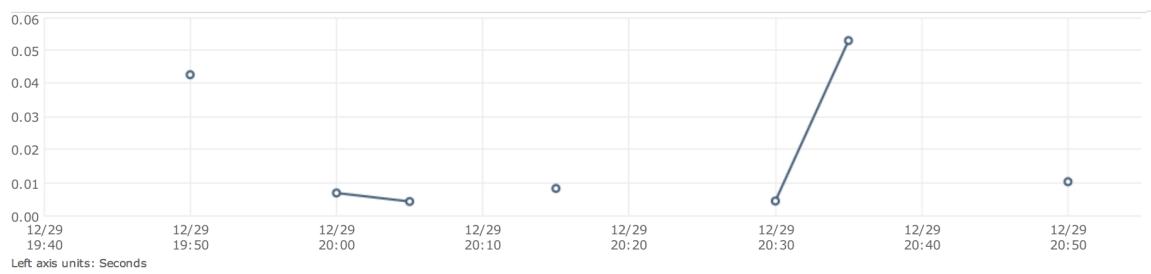


- From a server within Amazon cloud in the same geographic region,

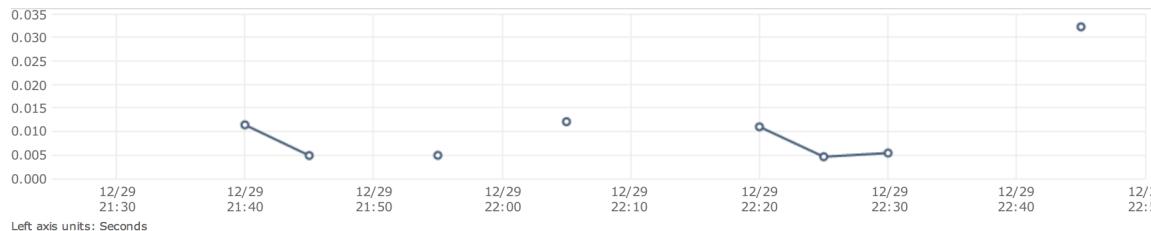
*Virginia web server*



- From a server within Amazon cloud in Oregon geographic region



- From a server within Amazon cloud in Northern California geographic region



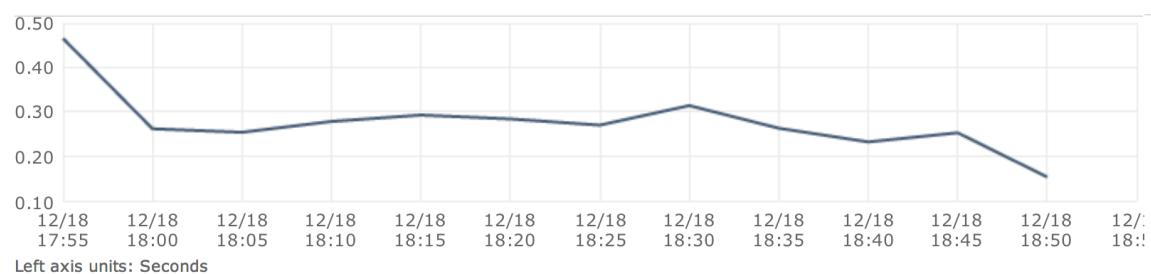
### 5.2.3.2 Second Test:

Test Place	High Alarm Period (~ Minutes) (>.30 Second)	Low Alarm Period (~ Minutes) (<= .005 Second)
Morning Weekdays RIT	6	0
Evening Weekdays RIT	5	0
Morning Weekends RIT	43	0
Evening Weekends RIT	35	0

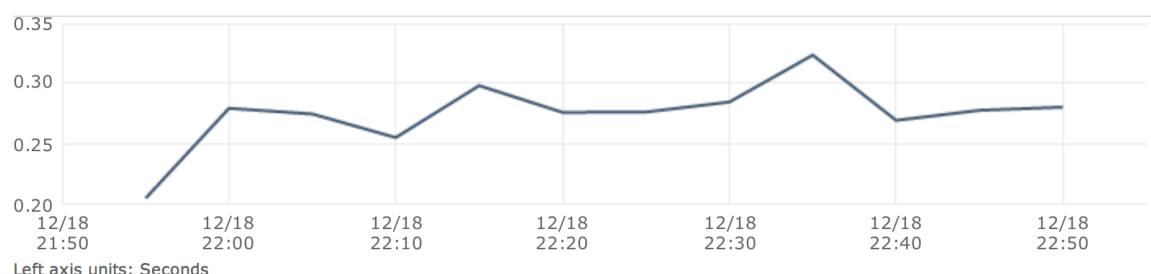
<i>Morning Weekdays Home</i>	18	0
<i>Evening Weekdays Home</i>	47	0
<i>Morning Weekends Home</i>	53	0
<i>Evening Weekends Home</i>	46	0
<i>Virginia</i>	53	0
<i>Oregon</i>	57	0
<i>North California</i>	17	0

Table 7: Latency Metric Results of First Test

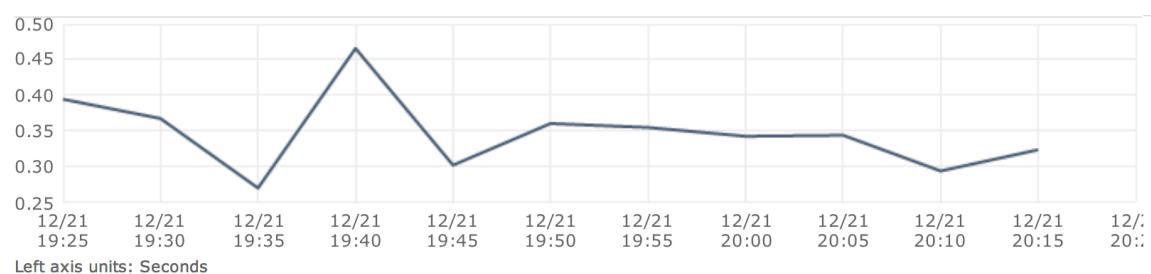
- *Weekday mornings at RIT System Lab*



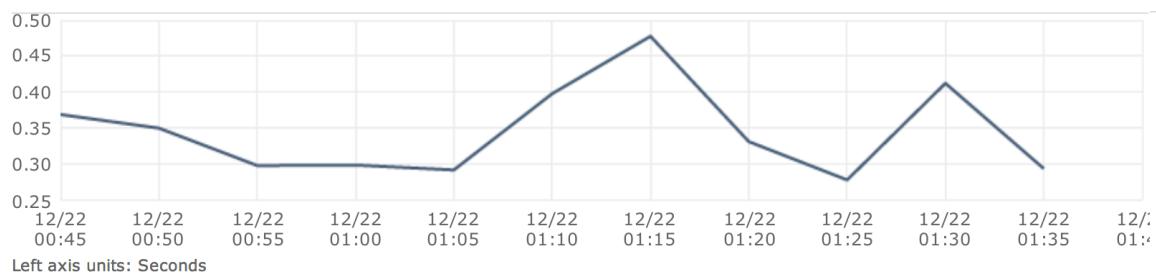
- *Weekday evenings at RIT System Lab*



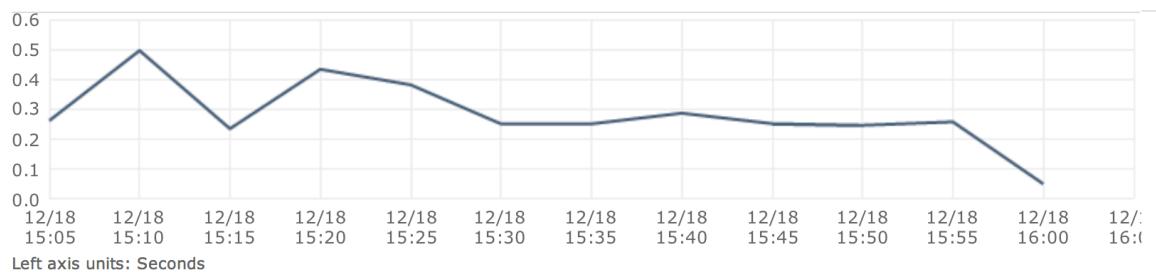
- *Weekend mornings at RIT System Lab*



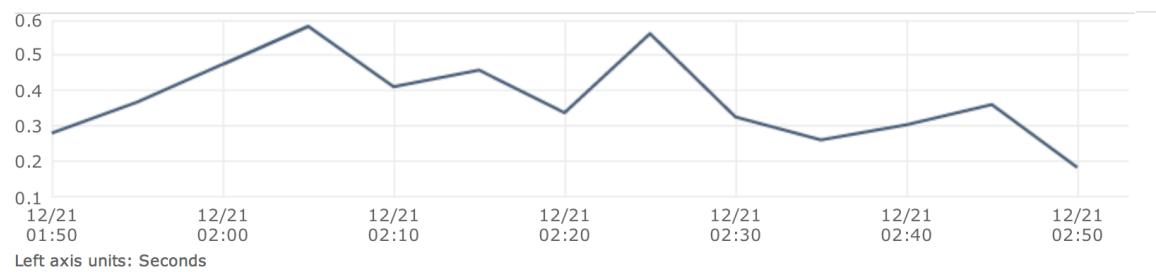
○ *Weekend evenings at RIT System Lab*



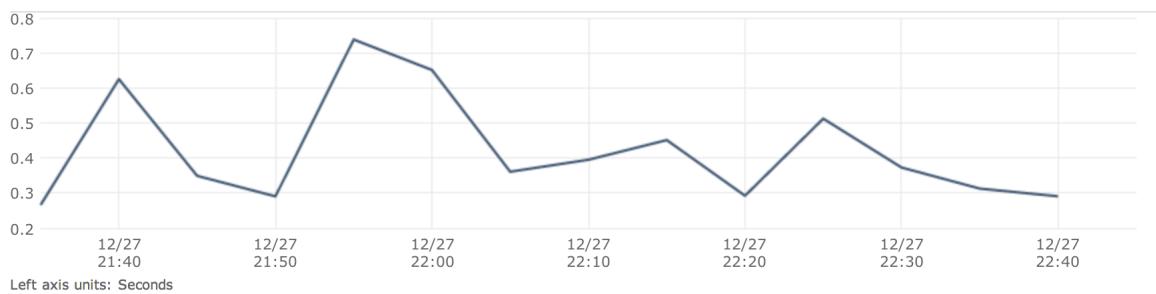
○ *Weekday mornings from home*



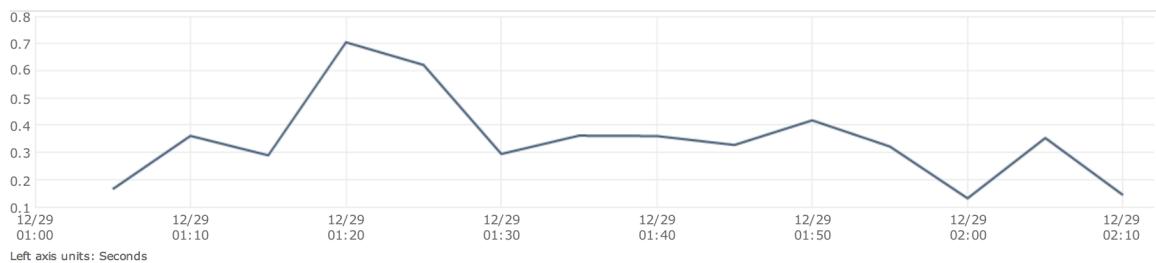
○ *Weekday evenings from home*



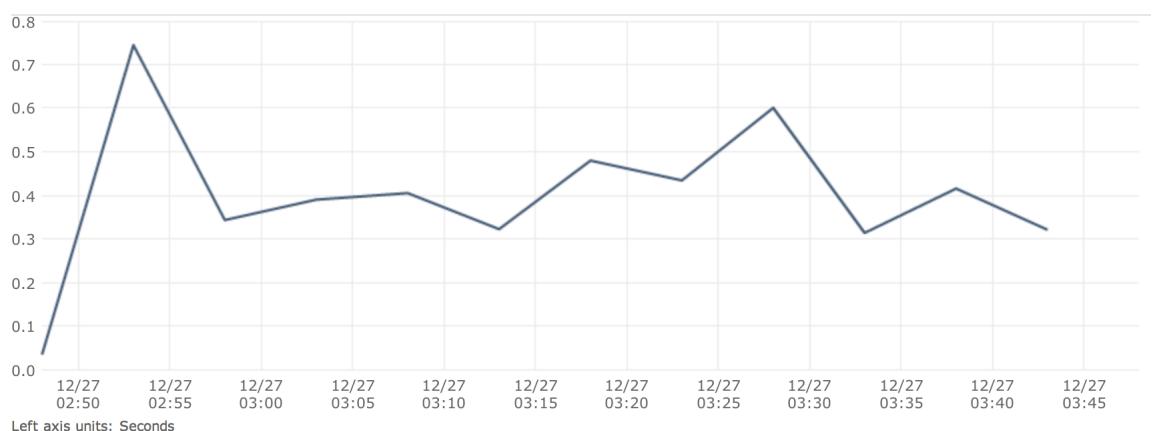
○ *Weekend mornings from home*



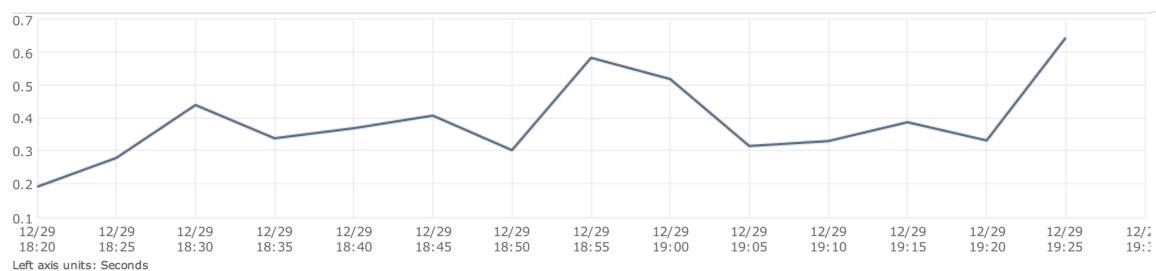
- *Weekend evenings from home*



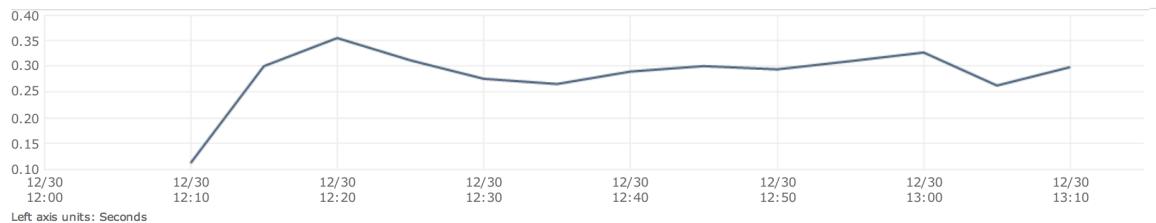
- *From a server within Amazon cloud in the same geographic region,  
Virginia web server*



- *From a server within Amazon cloud in Oregon geographic region*



- From a server within Amazon cloud in Northern California geographic region



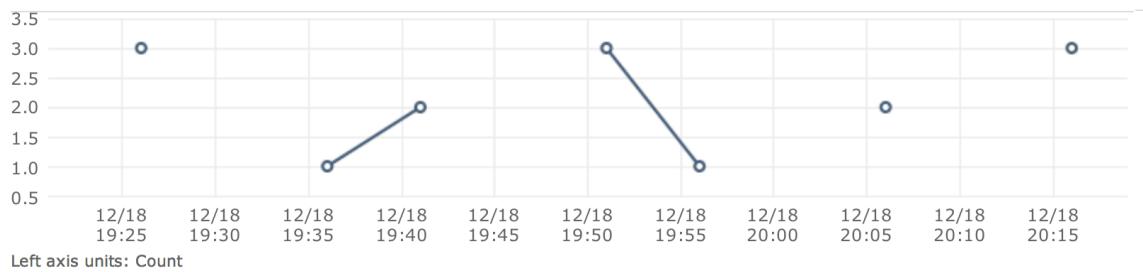
### 5.2.4 Request Counts Metric Results:

#### 5.2.4.1 First Test:

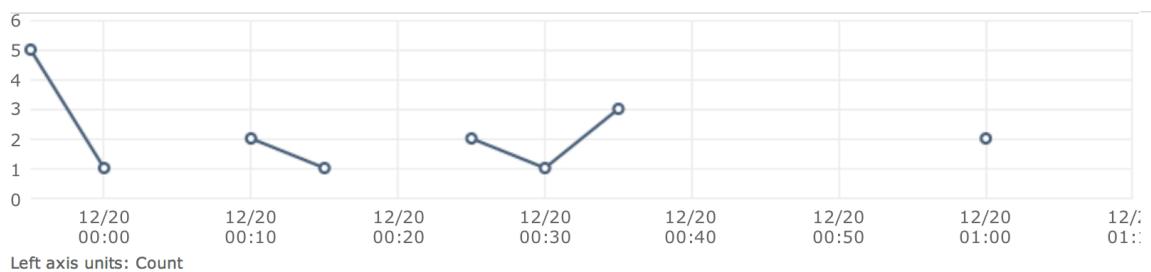
Test Place	High Alarm Period (~ Minutes) (>= 200 Counts)	Low Alarm Period (~ Minutes) (<2 Counts)
Morning Weekdays RIT	0	19
Evening Weekdays RIT	0	18
Morning Weekends RIT	0	34
Evening Weekends RIT	0	27
Morning Weekdays Home	0	35
Evening Weekdays Home	0	During the whole test period
Morning Weekends Home	0	During the whole test period
Evening Weekends Home	0	During the whole test period
Virginia	0	27
Oregon	0	34
North California	0	31

Table 8: Request Counts Metric Results of First Test

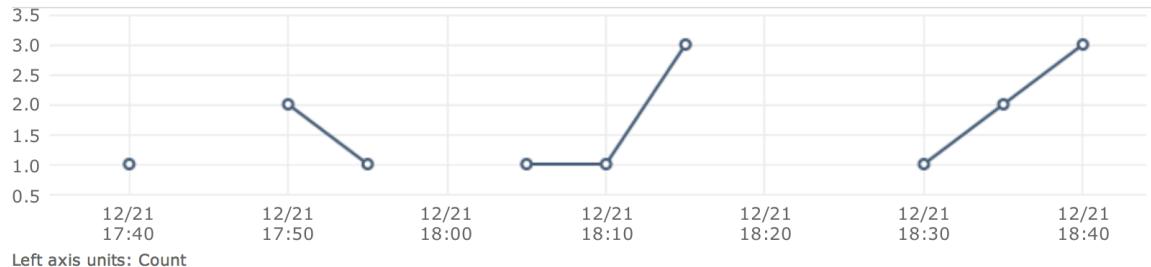
○ Weekday mornings at RIT System Lab



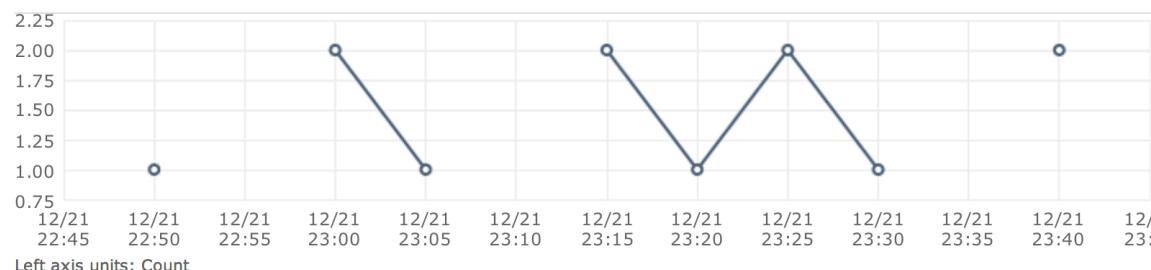
○ Weekday evenings at RIT System Lab



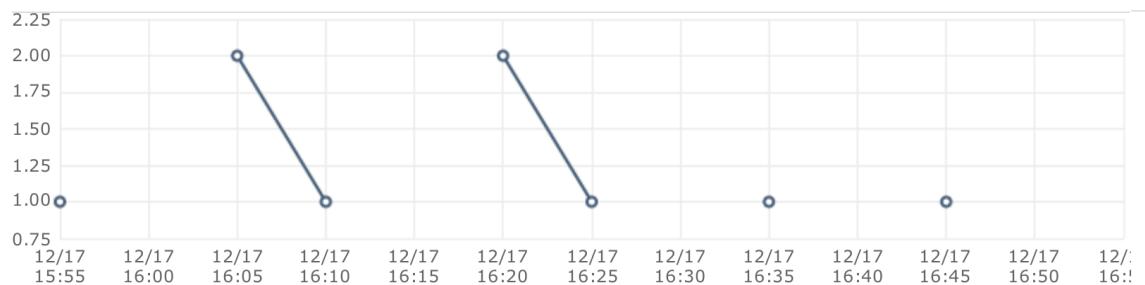
○ Weekend mornings at RIT System Lab



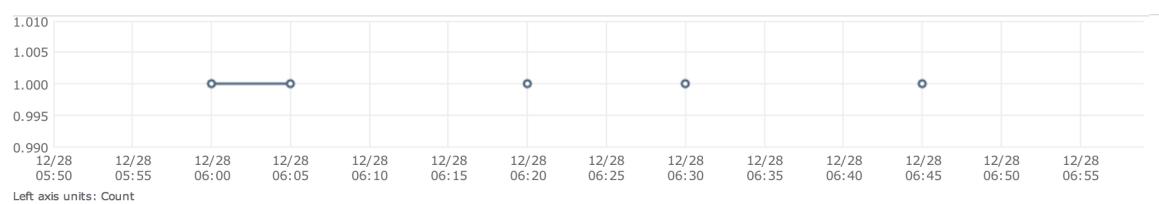
○ Weekend evenings at RIT System Lab



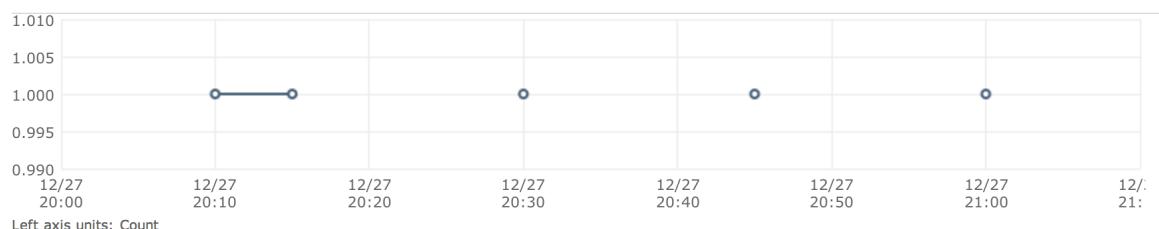
○ *Weekday mornings from home*



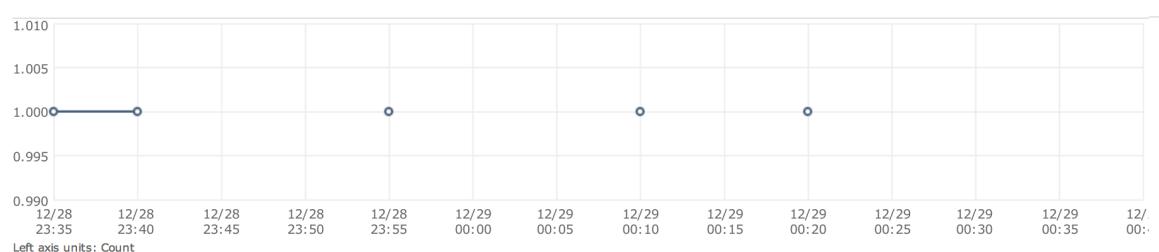
○ *Weekday evenings from home*



○ *Weekend mornings from home*

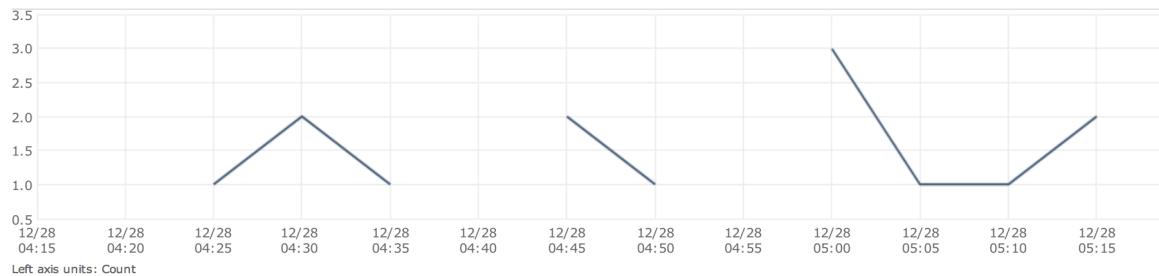


○ *Weekend evenings from home*

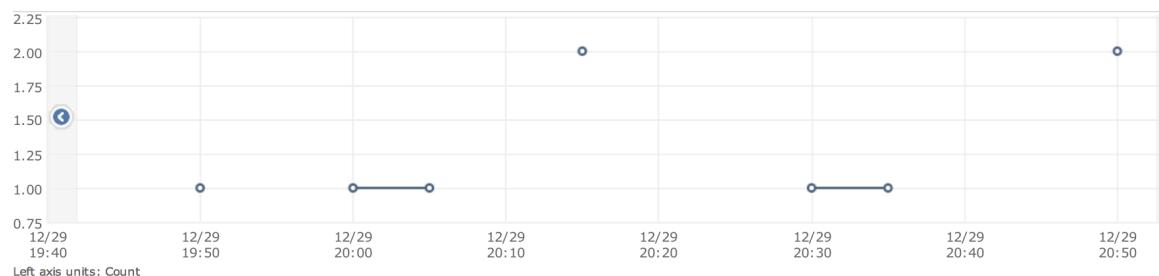


- From a server within Amazon cloud in the same geographic region,

*Virginia web server*



- From a server within Amazon cloud in Oregon geographic region



- From a server within Amazon cloud in Northern California geographic region



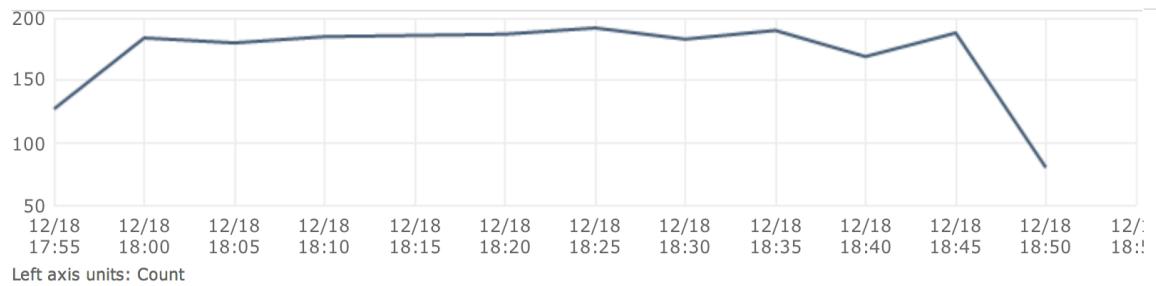
#### 5.2.4.2 Second Test:

Test Place	High Alarm Period (~ Minutes) ( $\geq 150$ Counts)	Low Alarm Period (~ Minutes) (<2 Counts)
Morning Weekdays RIT	51	0
Evening Weekdays RIT	56	0
Morning Weekends RIT	60	0
Evening	60	0

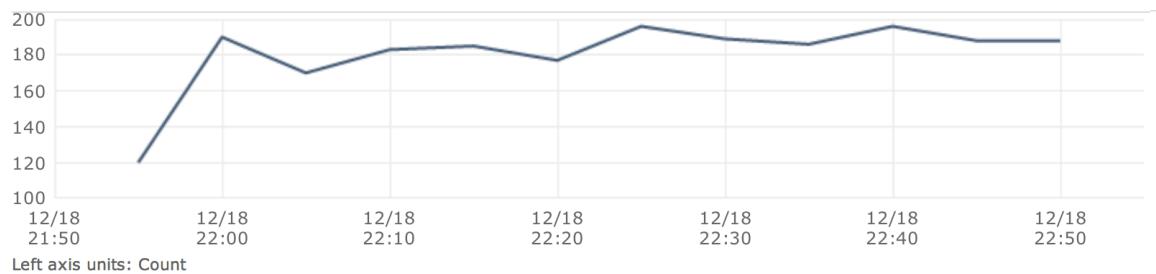
<i>Weekends RIT</i>		
<i>Morning Weekdays Home</i>	60	0
<i>Evening Weekdays Home</i>	54	0
<i>Morning Weekends Home</i>	60	0
<i>Evening Weekends Home</i>	49	0
<i>Virginia</i>	60	0
<i>Oregon</i>	60	0
<i>North California</i>	60	0

**Table 9: Requests Count Metric Results of Second Test**

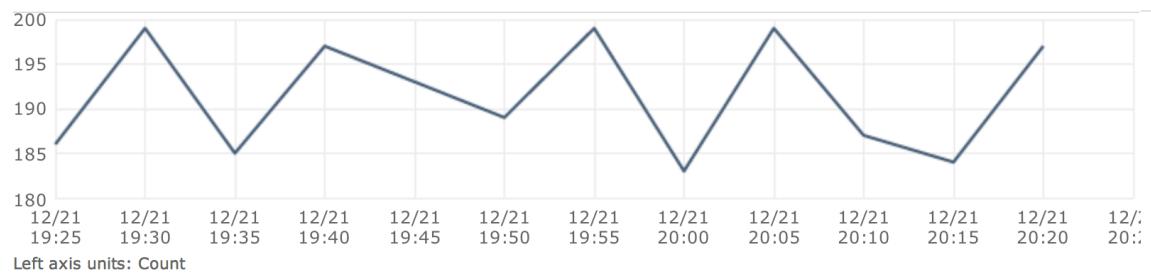
- *Weekday mornings at RIT System Lab*



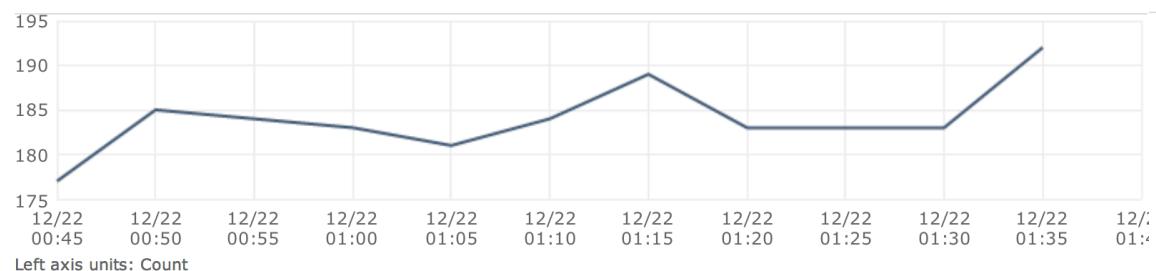
- *Weekday evenings at RIT System Lab*



○ *Weekend mornings at RIT System Lab*



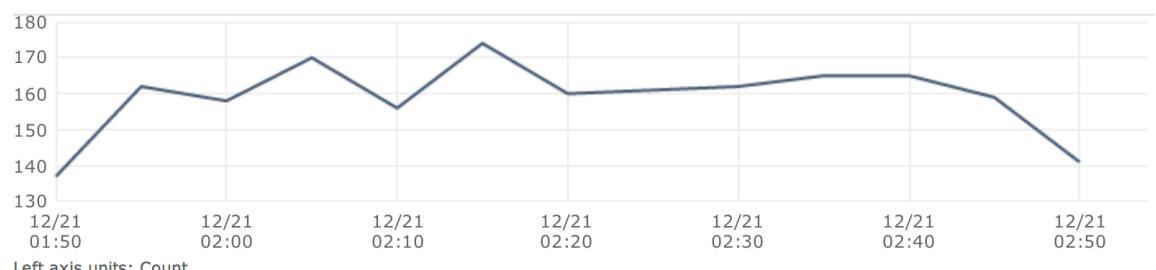
○ *Weekend evenings at RIT System Lab*



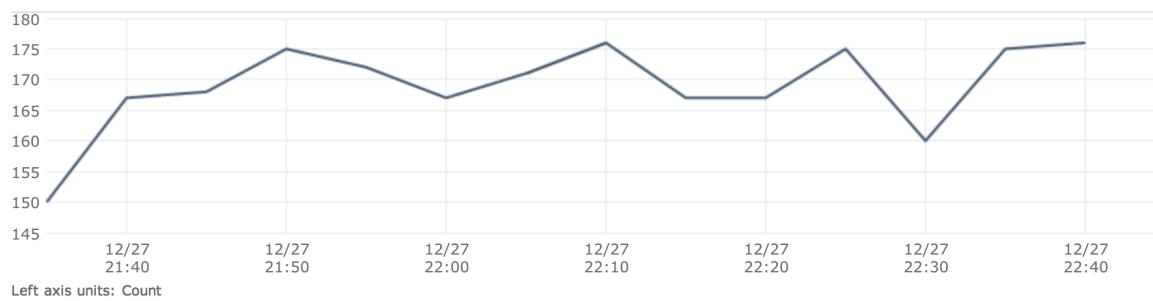
○ *Weekday mornings from home*



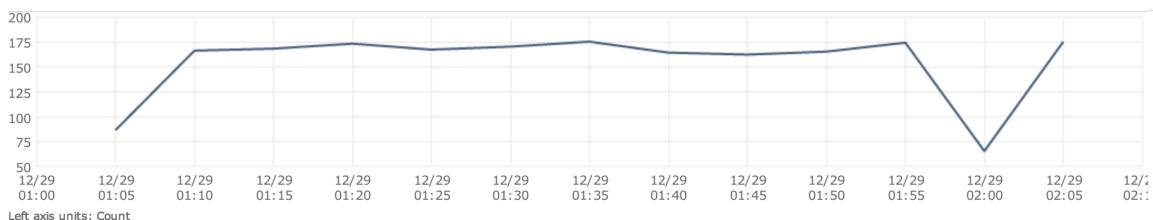
○ *Weekday evenings from home*



○ *Weekend mornings from home*

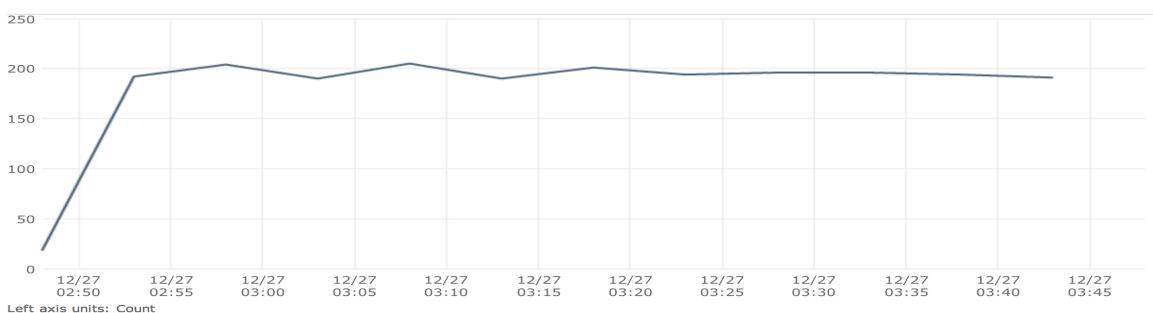


○ *Weekend evenings from home*

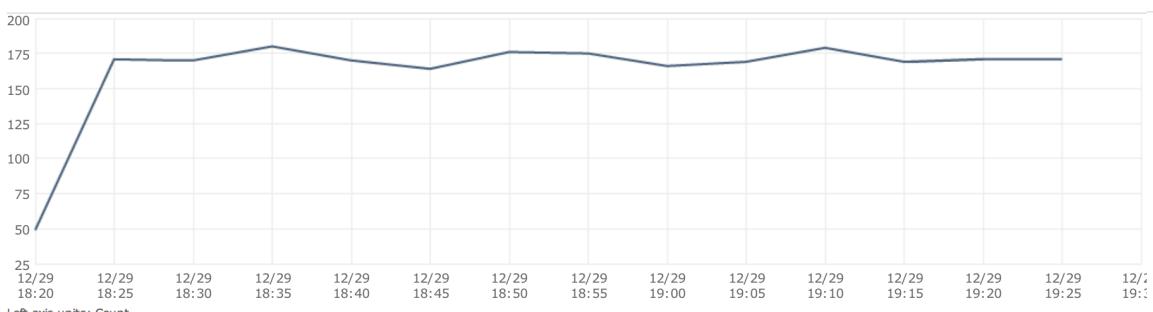


○ *From a server within Amazon cloud in the same geographic region,*

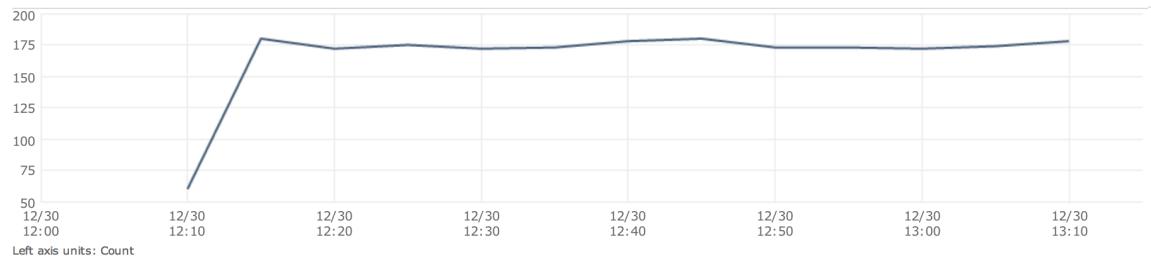
*Virginia web server*



○ *From a server within Amazon cloud in Oregon geographic region*



- From a server within Amazon cloud in Northern California geographic region



### 5.2.5 Discussion:

As explained earlier the FIRST TEST was coded to give better results when examining the EC2 instances metrics (CPU Utilization & Network Out), and the SECOND TEST was coded to give better results when examining the load balancer metrics (Latency & Request Counts). However, each metric was tested using the First and Second Tests in order to get more accurate results to help with the final decision of which metric is the best overall.

Below are explanations of each metric, followed by a full comparison of all the metrics together. Figure 5.9.1 and Figure 5.9.2 (below) show the average time of exceeding the boundaries of each metric for the First and the Second tests, respectively.

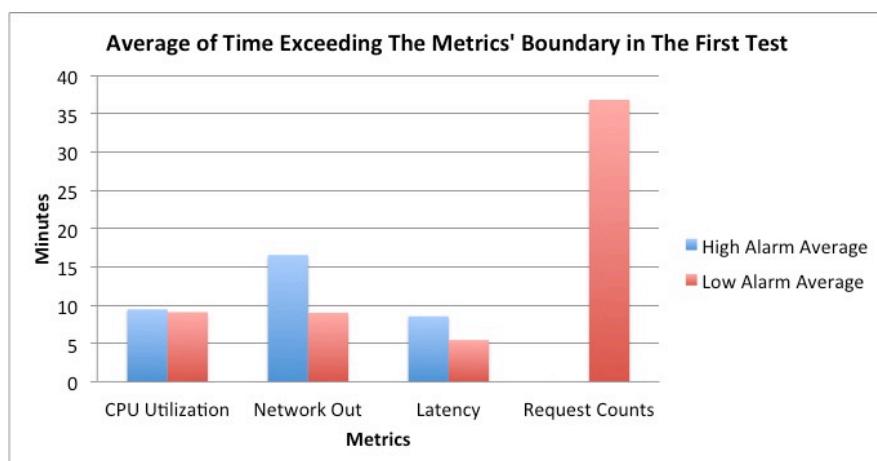
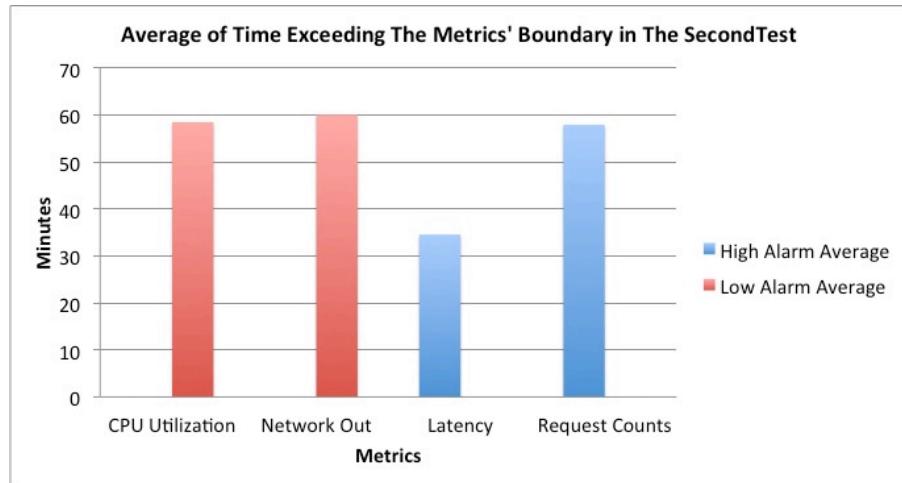


Figure 5.9.1: The Average of time exceeded the metrics' boundary in the First Test



**Figure 5.9.2: The Average of time exceeded the metrics' boundary in the Second Test**

- ***CPU Utilization Metric:***

In the First Test, the high boundary was set at 80%. This metric exceeded the high level in some tests for about half an hour, specifically those tests at RIT labs and Amazon regions (Virginia, Oregon, North California). The tests from home were a different case, would be tested perfectly if the high boundary was about 60% due to the low bandwidth and speed compared to that of RIT's. When looking at the home results in these tests, the metric did not exceed the high level except for 10 minutes in the Morning Weekday test, when it seemed that most people were at work, thus causing less Internet traffic. Staying below the high level alarm for the home test means that the throughput is higher when using the RIT System Lab, meaning that data rates delivered to the web server are also higher. The average of exceeding the high level for all tests was 9.45 minutes, which is about 1/6 of an hour. This long period of exceeding the high level in this short time of testing (one hour) should scale up with a new instance. For the lower level alarm, RIT's results showed that the metric exceeded the low level for short times, whereas the home's results showed that the metric exceeded about half an hour in some tests. The reasons for this are again, bandwidth and Internet speed.

In regards to the second test, the CPU Utilization metric performed poorly. The CPU performance was about 5% for almost the whole test period, meaning that auto scaling would be needed to scale down. However, this provided a conflict because the Second Test needed to scale up for other metrics such as Latency & Request Counts. Thus, the CPU Utilization was not an appropriate metric for the Second Test, whereas it was a good metric for the First Test.

- ***Network Out Metric:***

This metric's results gave very similar results as found by the CPU Utilization metric, except for the results conducted at RIT lab. As shown in Table 2 (CPU Utilization metric), the total time that the high level alarm was exceeded in RIT lab's tests (Morning Weekdays, Evening Weekdays, Morning Weekends, and Evening Weekends) were 0, 12, 25, and 17 minutes, respectively. Also, as shown in table 4 (Network Out metric), the total time that the high level alarm was exceeded in RIT lab's tests (Morning Weekdays, Evening Weekdays, Morning Weekends, and Evening Weekends) were 34, 25, 35, and 32 minutes, respectively. From this observation of the RIT lab's results for both metrics, the results of the CPU Utilization metric were too varied. On the other hand, the results of the Network Out metric had smaller differences. The deviations themselves resulted from the times that the tests were conducted. The average of all tests where the metric exceeded the high level was 16.5 minutes. The average for this metric was greater than the average for the CPU Utilization metric, mostly due to the unexpected results at the RIT lab. As a result, the Network Out metric is more accurate than the CPU Utilization metric. The time of exceeding the low level was very short in all tests except the home's test due to the same reasons for the CPU Utilization metric. The low level in both metrics (CPU Utilization & Network Out) was exceeded for almost the same length period

(with average  $\approx$  9 minutes). This observation was expected because they were running the First Test, which was coded specifically for these two metrics.

In relation to the second test, the Network Out metric performed as poorly as the CPU Utilization metric. As shown in Table 5, all tests had exceeded the low level for the whole test period. In other words, the auto scaling would need to scale down by one instance every five minutes during the whole test period. However, the web server was experiencing too many requests, and would need to scale up when using the second two metrics (Latency & Request Counts).

In summary, regarding the EC2 instances metrics: the Network Out metric is better than the CPU Utilization metric, due to the lesser variation of data and the greater average of exceeding the high boundary for the former. These two metrics will be compared with the load balancer metrics as well for the purpose of determining which metric is the best overall.

- ***Latency Metric:***

In the First Test, there were 5 out of 11 tests exceeding the high level for five minutes or more, during which auto scaling would take action by adding a new instance, if it was activated. The average of the time that exceeded the high level in the First Test was 8.55 minutes. This, however, is not accurate because the Evening Weekends test (at the RIT lab) was exceeding the high level for almost the entire test period, making it an outlier and obviously increasing the average. However, the metric still performed well in the First Test, even it was not as good as the EC2 metrics. In regards to the low level, the metric did not exceed it except in three tests (Morning Weekends at home, Evening Weekdays and Morning Weekends at RIT lab).

In the Second Test, the average of time exceeding the high level was 34.55 minutes, almost half the test period. This was expected because the Second Test was designed for the Latency & Request Counts metrics.

- ***Request Counts Metric:***

In the First Test, this metric did not exceed its high level of 150 counts. This was because the test was run by only three users, meaning it would not be possible to exceed three counts, hence why the minimum boundary was set at two counts.

In the Second Test, the average of exceeding the high level was 57.90 minutes, which was almost the entire test period. This was expected because the Second Test was coded specifically to stress the Request Counts metric.

The purpose of these experiments was to find which metric is the most suitable for the auto scaling. Overall, the Latency metric is the best as it was the only one successfully affected by both tests. In practical terms, if a few users downloaded large files from the server, the latency would go up as it did in the First Test. Similarly, if several users connected to the server during a short period of time, the latency would also go up as it did in the Second Test. In contrast, the CPU Utilization & Network Out metrics were successfully affected by the First Test only, and the Request Counts metric was successfully affected by the Second Test only.

In conclusion, the Latency metric was the most applicable metric overall, regardless of whether the server was experiencing users downloading large files or handling several smaller requests.

## **6. CONCLUSION & FUTURE WORK**

Cloud computing is becoming a goal for many organizations to use in running their services because of the cloud's benefits, such as reduced costs, increased storage, and flexibility. Cloud flexibility can enable users and companies to use as many instances as needed. When a web server needs more instances to run a particular service, auto scaling can be provided to scale the number of the instances up and down accordingly. In this project, an algorithm was implemented to scale the number of instances up and down based on the CPU Utilization metric. Two separate tests were also run, exploring the capabilities of the EC2 and load balancer metrics when faced with ideal or challenging conditions. The results found the Latency metric as the best, as it was the only metric to perform successfully under both tests.

The project was done on Amazon Elastic Compute Cloud. In the future, it could be implemented on Amazon Virtual Private Cloud (VPC) to examine the same scenario with the same presented scaling algorithm, addressing a dynamic scalability of web applications on a virtualized cloud computing environment based on the CPU utilization. There are other Amazon services that can be added to this project as well, such as Amazon DynamoDB (Amazon cloud data base) to create a database table and test how data or traffic transfers over instances within the VPC.

## 7. REFERENCES

- [1] S. Oriyano, P. Cabera. "Introduction to Software Load Balancing with Amazon EC2," Amazon EC2, 2008.
- [2] T. C. Chieu, A. Mohindra, A. A. Karve, A. Segal. "Dynamic Scaling of Web Applications in a Virtualized Cloud Computing Environment," *e-Business Engineering, 2009. ICEBE '09. IEEE International Conference on*, pp.281-286, 21-23 Oct. 2009.
- [3] <http://www.nist.gov/itl/csd/cloud-102511.cfm>
- [4] <http://www.nist.gov/itl/cloud/upload/cloud-def-v15.pdf>
- [5] M. Mao; J. Li; M. Humphrey. "Cloud auto-scaling with deadline and budget constraints," *Grid Computing (GRID), 2010 11th IEEE/ACM International Conference*, pp.41-48, 25-28 Oct. 2010. [Online]. Available:
- [6] R. Cushing, S. Koulouzis, A. Belloum, M. Bubak, Prediction-based Auto-scaling of Scientific Workflows, MGC '11 Proceedings of the 9th International Workshop on Middleware for Grids, Clouds and e-Science, December 12, 2011, Lisbon, Portugal, DOI 10.1145/2089002.2089003
- [7] Dean, J. and Ghemawat, S., "Mapreduce: simplified data processing on large clusters," *Commun. ACM*, 51, 1, pp. 107-113, 2008.
- [8] S. Venugopal, L. Han; P. Ray. "Auto-scaling emergency call centres using cloud resources to handle disasters," *Quality of Service (IWQoS), 2011 IEEE 19th International* pp.1-9, 6-7 June 2011. [Online]. Available:
- [9] Sourya. Article "A History of Cloud Computing," 2011
- [10] Amazon.com, Article "Amazon Media Room: History & Timeline", September 2011.
- [11] [https://aws.amazon.com/amis?\\_encoding=UTF8&ami\\_provider\\_id=2&jiveRedirect=1&selection=ami\\_provider\\_id](https://aws.amazon.com/amis?_encoding=UTF8&ami_provider_id=2&jiveRedirect=1&selection=ami_provider_id)
- [12] <http://docs.amazonaws.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>
- [13] Elastic Load Balancing Service on Amazon  
<http://aws.amazon.com/elasticloadbalancing/>
- [14] <http://aws.amazon.com/ec2/instance-types/>
- [15] <http://docs.amazonaws.com/AmazonElastiCache/latest/UserGuide/Concepts.RegionsAndAvailabilityZones.html>
- [16] <http://awsdocs.s3.amazonaws.com/AutoScaling/latest/as-dg.pdf>

- [17] N. Roy, A. Dubey, A. Gokhale. “Efficient Autoscaling in the Cloud Using Predictive Models for Workload Forecasting,” *Cloud Computing (CLOUD), 2011 IEEE International Conference on*, vol., no., pp. 500-507, 4-9 July 2011. [Online]. Available:
- [18] Qi Zhang · Lu Cheng · Raouf Boutaba, Cloud computing: state-of-the- art and research challenges, Springer, 20 April 2010.
- [19] Abouzeid, a., bajda-Pawlikowski, K., abadi, d.J., silberschatz, a., and rasin, a. Hadoopdb: an architectural hybrid of Mapreduce and dbMs technologies for analytical workloads. In Proceedings of the Conference on Very Large Databases (lyon, France, 2009);
- [20] Zhou, Dynamic Random Access for Hadoop Distributed File System, International Conference on Distributed Computing Systems Workshops, 2012